IBM

IBM ILOG CPLEX Optimization Studio CP Optimizer User's Manual

Version 12 Release 8

Contents

Chapter 1. Welcome to the CP Optimizer User's Manual

This is the *CP Optimizer User's Manual*.

Overview

CP Optimizer is a software library which provides a constraint programming engine.

The CP Optimizer feature of the IBM® ILOG® CPLEX® Optimizers is a software library which provides a constraint programming engine targeting both constraint satisfaction problems and optimization problems, including problems involving scheduling. This engine, designed to be used in a "model & run" development process, contains powerful methods for finding feasible solutions and improving them. The strength of the optimizer removes the need for you to write and maintain a search strategy.

CP Optimizer is based on IBM ILOG Concert Technology. Concert Technology offers a library of classes and functions that enable you to define models for optimization problems. Likewise, CP Optimizer offers a library of classes and functions that enable you to find solutions to the models. Though the CP Optimizer defaults will prove sufficient to solve most problems, CP Optimizer offers a variety of tuning classes and parameters to control various algorithmic choices.

IBM ILOG Concert Technology and CP Optimizer provide application programming interfaces (APIs) for Microsoft .NET Framework languages, C++ and Java™. The CP Optimizer part of an application can be completely integrated with the rest of that application (for example, the graphical user interface, connections to databases and so on) because it can share the same objects.

About this manual

The *CP Optimizer User's Manual* provides conceptual information about the features of CP Optimizer.

This is the *CP Optimizer User's Manual*. It offers explanations of how to use CP Optimizer effectively. All of the CP Optimizer functions and classes used in this manual are documented in the *CP Optimizer Reference Manuals*. As you study this manual, you will probably consult the appropriate reference manual from time to time, as it contains precise details on classes and their members.

Prerequisites

The *CP Optimizer User's Manual* assumes that you have a working knowledge of one of the programming languages of the available APIs and have installed CP Optimizer.

CP Optimizer requires a working knowledge of the Microsoft .NET Framework, C++ or Java. However, it does not require you to learn a new language since it does not impose any syntactic extensions on your programming language of choice.

If you are experienced in constraint programming or operations research, you are probably already familiar with many concepts used in this manual. However, no experience in constraint programming or operations research is required to use this manual. The *Getting Started with CP Optimizer* manual provides a tutorial introduction to many of the topics covered in this manual.

You should have IBM ILOG Concert Technology and CP Optimizer installed in your development environment before starting to use this manual. Moreover, you should be able to compile, link and execute a sample program provided with CP Optimizer.

Related documentation

The *CP Optimizer User's Manual* is part of a collection of manuals. You will likely need to refer to the other manuals in the collection as you use this manual.

The following documentation ships with CP Optimizer and will be useful for you to refer to as you use this manual.

- The *Getting Started with CP Optimizer Manual* introduces CP Optimizer with tutorials that lead you through describing, modeling and solving problems.
- The *CP Optimizer Extensions User's Manual* explains how to use the advanced features, such as propagators, custom constraints and custom search, of CP Optimizer effectively.
- The *CP Optimizer Reference Manuals* document the IBM ILOG Concert Technology and CP Optimizer classes and functions used in the *CP Optimizer User's Manual*. The reference manuals also explain certain concepts more formally. There are three reference manuals; one for each of the available APIs.
- The *Release Notes for CP Optimizer* list new and improved features, changes in the library and documentation and issues addressed for each release.

Installing CP Optimizer

The installation directions for CP Optimizer are in the Electronic Product Delivery package.

In this manual, it is assumed that you have already successfully installed the IBM ILOG Concert Technology and CP Optimizer libraries on your platform (that is, the combination of hardware and software you are using). If this is not the case, you will find installation instructions in your IBM ILOG Electronic Product Delivery package. The instructions cover all the details you need to know to install IBM ILOG Concert Technology and CP Optimizer on your system.

Typographic and Naming Conventions

The naming conventions used in *CP Optimizer User's Manual* manual follow the standard practices of the programming language being used.

Important ideas are *italicized* the first time they appear.

In this manual, the examples are given in C++. In the **C++ API**, the names of types, classes and functions defined in the IBM ILOG Concert Technology and CP Optimizer libraries begin with Ilo.

The name of a class is written as concatenated words with the first letter of each word in upper case (that is, capital). For example,

```
IloIntVar
```

A lower case letter begins the first word in names of arguments, instances and member functions. Other words in the identifier begin with an uppercase letter. For example:

```
IloIntVar aVar;
IloIntVarArray::add;
```

Names of data members begin with an underscore, like this:

```
class Bin {
public:
  IloIntVar      _type;
  IloIntVar      _capacity;
  IloIntVarArray _contents;
  Bin (IloModel    model,
       IloIntArray capacity,
       IloInt      nTypes,
       IloInt      nComponents);
  void display(const IloCP cp);
};
```

Generally, accessors begin with the key word `get`. Accessors for Boolean members begin with `is`. Modifiers begin with `set`.

Names of classes, methods and symbolic constants in the **C#** and the **Java API**s correspond very closely to those in the **C++ API** with these systematic exceptions:

- In the **C# API** and the **Java API**, namespaces are used. For **Java**, the namespaces are `ilog.cp` and `ilog.concert`. For **C#**, the namespaces are `ILOG.CP` and `ILOG.Concert`.
- In the **C++ API** and **the Java API**, the names of classes begin with the prefix `Ilo` whereas in the **C# API** they do not.
- In the **C++ API** and the **Java API**, the names of methods conventionally begin with a lowercase letter, for example, `startNewSearch`, whereas in the **C# API**, the names of methods conventionally begin with an uppercase letter, for example, `StartNewSearch`, according to Microsoft practice.

To make porting easier from platform to platform, IBM ILOG Concert Technology and CP Optimizer isolate characteristics that vary from system to system.

For that reason, you are encouraged to use the following identifiers for basic types in C++:

- `IloInt` stands for signed long integers;
- `IloNum` stands for double precision floating-point values ;
- `IloBool` stands for Boolean values: `IloTrue` and `IloFalse`.

You are not obliged to use these identifiers, but it is highly recommended if you plan to port your application to other platforms.

Chapter 2. Using CP Optimizer

Solving constraint programming problems with CP Optimizer can be broken into three steps: describing the problem, modeling the problem and finding solutions to the model of the problem. A basic constraint programming problem model consists of decision variables and constraints on those variables. Finding a solution to a model involves constraint propagation and search.

Overview of CP Optimizer

CP Optimizer is a software library which provides constructs for modeling and solving constraint programming problems.

The CP Optimizer feature of the IBM ILOG CPLEX Optimizers is a software library which provides a constraint programming engine targeting both satisfiability problems and optimization problems. This engine, designed to be used in a "model & run" development process, contains powerful search methods for finding feasible solutions and improving them. The strength of the optimizer removes the need for you to write and maintain a search strategy.

CP Optimizer is based on IBM ILOG Concert Technology. Concert Technology offers a library of classes and functions that enable you to define models for optimization problems. Likewise, CP Optimizer offers a library of classes and functions that enable you to find solutions to the models. Though the CP Optimizer defaults will prove sufficient to solve most problems, CP Optimizer offers a variety of tuning classes and parameters to control various algorithmic choices.

The three-stage method

The three-stage method of constraint programming using CP Optimizer involves describing, modeling and solving.

To find a solution to a problem using CP Optimizer, you use a three-stage method: describe, model and solve.

The first stage is to *describe* the problem in natural language. For more information, see the section "Describe" on page 6.

The second stage is to use Concert Technology classes to *model* the problem. The model is composed of decision variables and constraints. *Decision variables* are the unknown information in a problem. Each decision variable has a *domain* of possible *values*. The *constraints* are limits or restrictions on combinations of values for these decision variables. The model may also contain an *objective*, an expression that can be maximized or minimized. For more information, see the section "Model" on page 6.

The third stage is to use CP Optimizer classes to *solve* the problem. Solving the problem consists of finding a value for each decision variable while simultaneously satisfying the constraints and maximizing or minimizing an objective, if one is included in the model. The CP Optimizer *engine* (also called "the optimizer") uses two techniques for solving optimization problems: *search strategies* and *constraint propagation*. For more information, see the section "Solve the problem" on page 7.

In this section, the three stages of describe, model and solve are executed on a simple problem to underscore the basic concepts in constraint programming.

The problem is to find values for x and y given the following information:

- $x + y = 17$
- $x - y = 5$
- x can be any integer from 5 through 12
- y can be any integer from 2 through 17

Describe

The first stage in solving a constraint programming problem with CP Optimizer is to describe the problem using natural language.

The first stage is to *describe* the problem in natural language.

What is the unknown information, represented by the decision variables, in this problem?

- The values of x and y, where x is an integer between 5 and 12 inclusive and y is as integer between 2 and 17 inclusive.

What are the limits or restrictions on combinations of these values, represented by the constraints, in this problem?

- $x + y = 17$
- $x - y = 5$

Though the describe stage of the process may seem trivial in a simple problem like this one, you will find that taking the time to fully describe a more complex problem is vital for creating a successful program. You will be able to code your program more quickly and effectively if you take the time to describe the model, isolating the decision variables and constraints.

Model

The second stage in solving a constraint programming problem with CP Optimizer is to model the problem. The model is composed of decision variables and constraints. The model may also contain an objective.

Decision variables

Decision variables represent the unknown information in a constraint programming problem.

Decision variables represent the unknown information in a problem. Decision variables differ from normal programming variables in that they have domains of possible values and may have constraints placed on the allowed combinations of theses values. For this reason, decision variables are also known as *constrained variables*. In this example, the decision variables are x and y.

Each decision variable has a domain of possible values. In this example, the domain of decision variable x is [5..12], or all integers from 5 to 12. The domain of decision variable y is [2..17], or all integers from 2 to 17.

Note:

In *CP Optimizer* and Concert Technology, square brackets denote the *domain* of decision variables. For example, [5 12] denotes a domain as a set consisting of precisely two integers, 5 and 12. In contrast, [5..12] denotes a domain as a range of integers, that is, the interval of integers from 5 to 12, so it consists of 5, 6, 7, 8, 9, 10, 11 and 12.

Constraints

Constraints in a model represent the limit on the combinations of values for decision variables.

Constraints are limits on the combinations of values for variables. There are two constraints on the decision variables in this example: *x + y = 17* and *x - y = 5*.

Solve the problem

The third stage in solving a constraint programming problem with CP Optimizer is to search for a solution and solve the problem.

Solution

A solution to a constraint programming problem is a set of value assignments to the constrained variables.

A *solution* is a set of value assignments to the constrained variables such that each variable is assigned exactly one value from its domain and such that together these values satisfy the constraints. If there is an objective in the model, then an *optimal solution* is a solution that optimizes the objective function. Solving the problem consists of finding a solution for the problem or an optimal solution, if an objective is included in the model. The CP Optimizer engine uses efficient algorithms for finding solutions to constraint satisfaction and optimization problems.

Search space

The search space of a constraint programming problem is all combinations of the values in the domains of the decision variables.

The CP Optimizer engine explores the *search space* to find a solution. The search space is all combinations of values. One way to find a solution would be to explicitly study each combination of values until a solution was found. Even for this simple problem, this approach is obviously time-consuming and inefficient. For a more complicated problem with many variables, the approach would be unrealistic.

The optimizer uses two techniques to find a solution: search heuristics and constraint propagation. Additionally, the optimizer performs two types of constraint propagation: *initial constraint propagation* and *constraint propagation during search*.

Initial constraint propagation

Constraint propagation is a powerful technique used by CP Optimizer in the search for solutions to constraint programming problems. The initial constraint propagation removes values from domains that will not take part in any solution.

First, the CP Optimizer engine performs an initial constraint propagation. The initial constraint propagation removes values from domains that will not take part in *any* solution. Before propagation, the domains are:

```
D(x) = [5 6 7 8 9 10 11 12]
D(y) = [2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17]
```

To get an idea of how initial constraint propagation works, consider the constraint *x + y = 17*. If you take the smallest number in the domain of *x*, which is 5, and add it to the largest number in the domain of *y*, which is 17, the answer is 22. This combination of values (*x = 5, y = 17*) violates the constraint *x + y = 17*. The only value of *x* that would work with *y = 17* is *x = 0*. However, there is no value of 0 in the domain of *x*, so *y* cannot be equal to 17. The value *y = 17* cannot take part in any solution. The domain reduction algorithm employed by the constraint propagation engine removes the value *y = 17* from the domain of *y*. Similarly, the propagation engine removes the following values from the domain of *y*: 13, 14, 15 and 16.

Likewise, if you take the largest number in the domain of *x*, which is 12, and add it to the smallest number in the domain of *y*, which is 2, the answer is 14. This combination of values (*x = 12, y = 2*) violates the constraint *x + y = 17*. The only value of *x* that would work with *y = 2* is *x = 15*. However, there is no value of 15 in the domain of *x*, so *y* cannot be equal to 2. The value of *y = 2* cannot take part in any solution. the propagation engine removes the value *y = 2* from the domain of *y*. For the same reason, the domain reduction algorithm employed by the propagation engine removes the following values from the domain of *y*: 2, 3 and 4.

After initial propagation for the constraint *x + y = 17*, the domains are:
```
D(x) = [5 6 7 8 9 10 11 12]
D(y) = [5 6 7 8 9 10 11 12]
```

Now, examine the constraint *x - y = 5*. If you take the value 5 in the domain of *x*, you can see that the only value of *y* that would work with *x = 5* is *y = 0*. However, there is no value of 0 in the domain of *y*, so *x* cannot equal 5. The value *x = 5* cannot take part in any solution. The propagation engine removes the value *x = 5* from the domain of *x*. Using similar logic, the propagation engine removes the following values from the domain of *x*: 6, 7, 8 and 9. Likewise, the domain reduction algorithm employed by the propagation engine removes the following values from the domain of *y*: 8, 9, 10, 11 and 12.

Returning to the other constraint, there are no further values that can be removed from the variables. After initial propagation, the search space has been reduced in size. The domains are now:
```
D(x) = [10 11 12]
D(y) = [5 6 7]
```

Constructive search

Along with constraint propagation, constructive search strategies are used by CP Optimizer in the search for solutions to constraint programming problems.

After initial constraint propagation, the search space is reduced. CP Optimizer uses a constructive search strategy to guide the search for a solution in the remaining part of the search space. It may help to think of the strategy as one that traverses a *search tree*. The *root* of the tree is the starting point in the search for a solution; each *branch* descending from the root represents an alternative in the search. Each combination of values in the search space can be seen as a *leaf node* of the search tree.

The CP Optimizer engine executes a search strategy that guides the search for a solution. The optimizer "tries" a value for a variable to see if this will lead to a

solution. To demonstrate how the optimizer uses search strategies to find a solution, consider a search strategy that specifies that the optimizer should select variable *x* and assign it the lowest value in the domain of *x*. For the first *search move* in this strategy, the optimizer assigns the value 10 to the variable *x*. This move, or search tree branch, is not permanent. If a solution is not found with *x = 10*, then the optimizer can undo this move and try a different value of *x*.

Constraint propagation during search

Constraint propagation is combined with a constructive search strategy is used by CP Optimizer in the search for solutions to constraint programming problems.

The CP Optimizer engine performs constraint propagation during search. This constraint propagation differs from the initial constraint propagation. The initial constraint propagation removes all values from domains that will not take part in *any* solution. Constraint propagation during search removes all values from the *current* domains that violate the constraints. You can think of constraint propagation during search in the following way. In order to "try" a value for a variable, the optimizer creates "test" or *current* domains. When constraint propagation removes values from domains during search, values are only removed from these "test" domains.

To continue the same example, suppose that, based on the search strategy, the optimizer has assigned the value 10 to the decision variable *x*. Working with the constraint *x + y = 17*, constraint propagation reduces the domain of *y* to [7]. However, this combination of values (*x = 10, y = 7*) violates the constraint *x - y = 5*. The optimizer removes the value *y = 7* from the current domain of *y*. At this point, the domain of *y* is empty, and the optimizer encounters a *failure*. The optimizer can then conclude that there is no possible solution with the value of 10 assigned to *x*.

When the optimizer decides to try a different value for the decision variable *x*, the domain of *y* is at first restored to the values [5 6 7]. It then reduces the domain of *y* based on the new value assigned to *x*.

This simple example demonstrates the basic concepts of constructive search and constraint propagation. To summarize, solving a problem consists of finding a value for each decision variable while simultaneously satisfying the constraints. The CP Optimizer engine uses two techniques to find a solution: constructive search with search strategies and constraint propagation. Additionally, the optimizer performs two types of constraint propagation: initial constraint propagation and constraint propagation during search.

The initial constraint propagation removes values from domains that will not take part in any solution. After initial constraint propagation, the search space is reduced. This remaining part of the search space, where the CP Optimizer engine will use constructive search with a search strategy to search for a solution, is called the search tree. Constructive search is a way to "try" a value for a variable to see if this will lead to a solution. The optimizer performs constraint propagation during search. Constraint propagation during search removes all values from the current or "test" domains that violate the constraints. If the optimizer cannot find a solution after a series of choices, these can be reversed and alternatives can be tried. The CP Optimizer engine continues to search using the constructive search and constraint propagation during search until a solution is found.

Scheduling in CP Optimizer

CP Optimizer offers classes and functions specially adapted to modeling and solving problems in scheduling.

Basic building blocks of scheduling models

The basic building blocks of scheduling include time intervals and the constraints amongst those intervals.

In addition to constrained integer variables, CP Optimizer provides a set of modeling features for applications dealing with scheduling over time. Although in CP Optimizer, time points are represented as integers, time is effectively continuous because the range of time points is potentially very wide.

A consequence of scheduling over effectively continuous time is that the evolution of some known quantities over time (for instance the instantaneous efficiency/speed of a resource or the earliness/tardiness cost for finishing an activity at a given date) needs to be compactly represented in the model.

Most of the scheduling applications consist of scheduling in time a set of activities, tasks or operations that have a start and an end time. In CP Optimizer, this type of decision variable is captured by the notion of *interval decision variable*.

Several types of constraints are expressed on and between interval decision variables:
- to limit the possible positions of an interval decision variable (forbidden start/end or "extent" values),
- to specify precedence relations between two interval decision variables and
- to relate the position of an interval variable with one of a set of interval decision variables (such as with spanning, synchronization, or alternative constraints).

An important characteristic of scheduling problems is that intervals may be optional and whether to execute an interval or not may be a decision variable. In CP Optimizer, this is captured by the notion of a Boolean presence status associated with each interval decision variable. Logical relations can be expressed between the presence of interval variables, for instance to state that whenever interval *a* is present then interval *b* must also be present.

Another aspect of scheduling is the allocation of limited resources to time intervals. The evolution of a resource over time can be modeled by three types of decision variables and expressions:
- The evolution of a disjunctive resource over time can be described by the sequence of intervals that represent the activities executing on the resource. CP Optimizer introduces the notion of an *interval sequence variable*. Constraints and expressions are available to control the sequencing of a set of interval variables.
- The evolution of a cumulative resource often needs a description of how the resource accumulated usage evolves over time. CP Optimizer provides *cumul function expressions* that can be used to constrain the evolution of resource usage over time.
- The evolution of a resource of infinite capacity, the state of which can vary over time is captured in CP Optimizer by *state functions*. The dynamic evolution of a

state function can be controlled using transition distances and constraints for specifying conditions on the state function that must be satisfied during fixed or variable intervals.

Some classical cost functions in scheduling are earliness/tardiness costs, makespan and activities execution/non-execution costs. CP Optimizer generalizes these classical cost functions and provides a set of basic expressions that can be combined together to express a large spectrum of scheduling cost functions that can be efficiently exploited by the CP Optimizer search.

Search space for scheduling models

The search space for scheduling models is usually independent from the size of the domain

For interval variables, the search does not enumerate the values in the domain, and the size of the search space is usually independent from the size of the domain. The search space is estimated as the number of possible permutations of the n interval variables of the problem whose log (base 2) is estimated as $n.log2(n)$.

This estimation is slightly adjusted. For an optional interval variable, there is an additional Boolean decision to be made on its presence status. In addition, the interval dependencies induced by an alternative constraints need to be taken into account; the selected interval bi in an alternative constraint $alternative(a,\{b1,...,bn\})$ has the same start and end values as the alternative master interval variable a.

Using search parameters

It is possible to set parameters on the CP Optimizer object to control the output, to control the constraint propagation, to limit the search and to control the search engine. It is important to observe that any parameter change from its default is displayed at the head of the search log

Setting parameters

The parameters in CP Optimizer can be set using a method of the class representing the optimizer.

In the **C++ API** of CP Optimizer, you set a parameter on the optimizer with a call to `IloCP::setParameter`. The first argument to this function is either `IloCP::IntParam` or `IloCP::NumParam`. The second argument is a value of type `IloInt`, `IloNum` or the enumerated type `IloCP::ParameterValues`. To set a parameter on the optimizer in the C++ API, you use the method `IloCP::setParameter`, for example:

```
IloCP cp(model);
cp.setParameter(IloCP::SearchType, IloCP::DepthFirst);
cp.solve();
```

In the **Java API** of CP Optimizer, you set a parameter on the optimizer with a call to `IloCP.setParameter`. The first argument to this function is either `IloCP.IntParam` or `IloCP.DoubleParam`. The second argument is a value of type `int` or `double` or an instance of a subclass of `IloCP.ParameterValues`. To set a parameter on the optimizer in the Java API, you use the method `IloCP.setParameter`, for example:

```
IloCP cp = new IloCP();
// add variables and constraints
cp.setParameter(IloCP.IntParam.SearchType,
                IloCP.ParameterValues.DepthFirst);
cp.solve();
```

Likewise, in the **C# API** of CP Optimizer, you set a parameter on the optimizer with a call to `CP.SetParameter`. The first argument to this function is either `CP.IntParam` or `CP.DoubleParam`. The second argument is a value of type `Int32` or `Double` or an instance of a subclass of `CP.ParameterValues`. To set a parameter on the optimizer in the C# API, you use the method `CP.SetParameter`, for example:

```
CP cp = new CP();
// add variables and constraints
cp.SetParameter(CP.IntParam.SearchType,
                CP.ParameterValues.DepthFirst);
cp.Solve();
```

Some parameters may not be changed while there is an active search, such as between calls to the optimizer methods `startNewSearch` and `endSearch`. You can change any limit, such as `ChoicePoinLimit`, `BranchLimit`, `TimeLimit`, `SolutionLimit` and `FailLimit` during search.

The appropriate values are detailed in the *CP Optimizer Reference Manuals*. Most of the search parameters available for use in CP Optimizer are discussed in more detail throughout this manual.

Time mode parameter

The time mode parameter defines the manner in which CP Optimizer measures time.

CP Optimizer uses time for both display purposes and for limiting the search. These timings can be measured either by CPU time or by elapsed time, and the time mode parameter defines how time is measured in CP Optimizer.

When multiple processors are available and the number of workers is greater than one, then the CPU time can be greater than the elapsed time by a factor up to the number of workers.

In the **C++ API** of CP Optimizer, the time mode is controlled with the parameter `IloCP::TimeMode`. A value of `IloCP::CPUTime` indicates that time should be measured as CPU time, `IloCP::ElapsedTime` indicates that time should be measured as elapsed time. The default is `IloCP::ElapsedTime`.

In the **Java API** of CP Optimizer, the time mode is controlled with the parameter `IloCP.IntParam.TimeMode`. A value of `IloCP.ParameterValues.CPUTime` indicates that time should be measured as CPU time, `IloCP.ParameterValues.ElapsedTime` indicates that time should be measured as elapsed time. The default is `IloCP.ParameterValues.ElapsedTime`.

Likewise, in the **C# API** of CP Optimizer, the time mode is controlled with the parameter `CP.IntParam.TimeMode`. A value of `CP.ParameterValues.CPUTime` indicates that time should be measured as CPU time, `CP.ParameterValues.ElapsedTime` indicates that time should be measured as elapsed time. The default is `CP.ParameterValues.ElapsedTime`.

Chapter 3. Modeling a problem with Concert Technology

After describing your constraint satisfaction or optimization problem, you use Concert Technology classes to model the problem. A Concert Technology model consists of a set of objects, known as modeling objects. Each decision variable, each constraint and the objective function in a model are all represented by objects of the appropriate class.

Using the environment

The environment manages internal modeling issues.

Creating the environment

The environment, which handles communication channels and memory management, must be created before the modeling objects are created.

The first step in an IBM ILOG Concert Technology application using the C++ API is to create the environment, an instance of the class `IloEnv`.

The environment manages internal modeling issues; it handles output, memory management for modeling objects and termination of search algorithms. In the Microsoft .NET Framework languages and Java APIs, issues regarding the environment are handled internally.

Normally an application needs only one environment, but you can create as many environments as you want. Typically, the environment is created early in the main part of an application, like this:

```
IloEnv env;
```

In the **C++ API**, every Concert Technology model and every optimizer object must belong to an environment. In programming terms, when you construct a model, you must pass one instance of `IloEnv` as an argument to that constructor.

In the **Java API**, the Concert Technology functions for creating modeling objects are defined in the interface `IloModeler` and implemented in the class `IloCP`.

Likewise, in the **C# API**, the functions for creating modeling objects are defined in the interface `IModeler`, and the class `CP` inherits them.

Environment and communication channels

The communication channels provided in the environment include output channels for different types of information.

In the C++ API, an instance of `IloEnv` in your application initializes its default output channels for general information, for warnings and for error messages.

Each environment maintains its own channels. The channels associated with an environment are `IloEnv::out`, `IloEnv::warning` and `IloEnv::error`. By default, these output streams are defined as `std::cout`. You can redirect these streams by calling `IloEnv::setOut(ostream &)` and other related member functions of `IloEnv`.

In the Microsoft .NET Framework languages and Java APIs, the native streams are used directly. To redirect output generated by the optimizer, you use the method `IloCP.setOut` in the **Java API** and the method `CP.SetOut` in the **C# API**. For example, in the C# API, the output from the optimizer can be redirected to the native error stream using the following code:

```
CP cp = new CP();
cp.SetOut(Console.Error);
```

Environment and memory management

The environment handles memory management of the model and algorithm objects.

When your C++ application deletes an instance of `IloEnv`, Concert Technology will automatically delete all models, algorithms (optimizers) and other objects depending on that environment as well.

To allocate on the environment memory pool in C++, you must pass the environment as an argument to the `new` operator:

```
  MyObject*  myobject = new (env) MyObject();
```

Memory allocated in the environment is reclaimed when the environment is terminated by the member function `IloEnv::end`. You must not use the delete operator for objects allocated on the environment memory pool. The destructor of these objects will be called when the memory is reclaimed.

To free memory used by a model in the **Java API**, you use the method `IloCP.end`. To free memory used by a model in the **C# API**, you use the method `CP.End`.

Note:

Environment

An instance of the class `IloEnv` manages the internal modeling issues, which include handling output, memory management for modeling objects and termination of search algorithms.

This instance is typically referred to as the *environment*. Normally an application needs only one environment, but you can create as many environments as you want.

In the C# and Java APIs, the environment object is not public. To free memory used by a model in the Java API, you use the method `IloCP.end`. To free memory used by a model in C# API, you use the method `CP.End`.

Managing data

Concert Technology provides classes that simplify the management of data.

Arrays

Arrays are used extensively in modeling with Concert Technology.

The data for an IBM ILOG Concert Technology model is often presented in terms of arrays.

In the Microsoft .NET Framework languages and the Java APIs, the native array classes are used to store data in arrays and are passed as arguments to many Concert Technology functions. In the C++ API, objects of the class `IloIntArray` can be used to store integer data in arrays.

Elements of the class `IloIntArray` can be accessed like elements of standard C++ arrays, but the class also offers a wealth of additional functions. For example, Concert Technology arrays are extensible; in other words, they transparently adapt to the required size when new elements are added using the method `add`. Conversely, elements can be removed from anywhere in the array with the method `remove`. Concert Technology arrays also provide debugging support when compiled in debug mode by using assertion statements to ensure that no element beyond the array bounds is accessed. Input and output operators (that is, `operator<<` and `operator>>`) are provided for arrays. For example, the code produces the following output:

`[1, 2, 3]`

This output format can be read back in with the `operator>>`, for example:

```
std::cin >> data;
```

When you have finished using an array and want to reclaim its memory, call the method `end`; for example, `data.end`. When the environment ends, all memory of arrays belonging to that same environment is returned to the system as well. Thus, in practice you do not need to call `end` on an array (or any other Concert Technology object) just before calling `IloEnv::end`.

Note:

Array of integer values

Arrays of integer values are represented by the class `IloIntArray` in the C++ API of Concert Technology. These arrays are extensible.

When you use an array, you can access a value in that array by its index, and the `operator[]` is overloaded for this purpose.

In the C# and Java APIs, the native arrays are used.

Finally, the **C++ API** of Concert Technology provides the template class `IloArray<X>` to create array classes for your own type X. This technique can be used to generate multidimensional arrays. All the functions mentioned here are supported for `IloArray` classes except for input/output, which depends on the input and output operator being defined for type X.

Note:

Extensible array

In the C++ API, Concert Technology provides the template class `IloArray` which makes it easy for you to create classes of arrays for elements of any given class. In other words, you can use this template to create arrays of Concert Technology objects; you can also use this template to create arrays of arrays (that is, multidimensional arrays).

When you use an array, you can access a value in that array by its index, and the `operator[]` is overloaded for this purpose.

The classes you create in this way consist of extensible arrays. That is, you can add elements to the array as needed.

In the **Java API**, arrays used by a CP Optimizer method might be created as such:

```
int[] fixedCost ={ 1, 2, 4, 3 };
IloIntVar[] open = cp.intVarArray(4,0,1);
IloIntExpr obj = cp.prod(open, fixedCost);
```

In the **C# API**, arrays used by a CP Optimizer method might be created as:

```
int[] fixedCost ={ 1, 2, 4, 3 };
IIntVar[] open = cp.IntVarArray(4,0,1);
IIntExpr obj = cp.Prod(open, fixedCost);
```

Sets of tuples

Sets of tuples can be used to model data that is related.

In many constraint applications, it is necessary to process a huge quantity of data. For instance, the features of some products can be described as a relation in a database or in text files. In this case, a useful data modeling object is a *tupleset*, or a set of tuples.

A *tuple* is an ordered set of values represented by an array. Tuples are useful for representing allowed combinations of data in a model. A *set of integer tuples* in a model is represented by an instance of a tupleset.

The elements of a tupleset are tuples of integer values, represented by arrays. The number of values in a tuple is known as the *arity* of the tuple, and the arity of the tuples in a set is called the arity of the set. (In contrast, the number of tuples in the set is known as the cardinality of the set.)

In the **C++ API** of CP Optimizer, the class `IloTupleSet` represents tuplesets.

In the **Java API** of CP Optimizer, the interface `IloTupleset` represents tuplesets.

In the **C# API** of CP Optimizer, the interface `ITupleSet` represents tuplesets.

Note:

Set of tuples

An *integer tuple* is an ordered set of values represented by an array. A set of integer tuples in a model is represented by a tupleset.

The number of values in a tuple is known as the arity of the tuple.

Consider as an example a bicycle factory that can produce thousands of different models. For each model of bicycle, a relation associates the features of that bicycle such as size, weight, color and price. This information can be used in a constraint programming application that allows a customer to find the bicycle that most closely fits a specification.

Then the tupleset `bicycleSet` defines the set of possible combinations of features. In the C++ API, the tupleset is created and built as follows:

```
IloIntTupleSet bicycleSet(env, 5);
bicycleSet.add(IloIntArray(env, 5, 1, 57, 12, 3, 1490));
bicycleSet.add(IloIntArray(env, 5, 2, 57, 13, 5, 1340));
```

```
        bicycleSet.add(IloIntArray(env, 5, 3, 60, 14, 3, 1790));
        bicycleSet.add(IloIntArray(env, 5, 4, 65, 14, 7, 1550));
        bicycleSet.add(IloIntArray(env, 5, 5, 67, 15, 2, 2070));
        bicycleSet.add(IloIntArray(env, 5, 6, 70, 15, 2, 1990));
```

In the Java API, the tupleset is created using the method `IloCP.intTable` and built as follows:

```
        IloIntTupleSet bicycleSet = cp.intTable(5);
        int[][] tuples = {{1, 57, 12, 3, 1490},
                          {2, 57, 13, 5, 1340},
                          {3, 60, 14, 3, 1790},
                          {4, 65, 14, 7, 1550},
                          {5, 67, 15, 2, 2070},
                          {6, 70, 15, 2, 1990}};
        cp.addTuple(bicycleSet, tuples[0]);
        cp.addTuple(bicycleSet, tuples[1]);
        cp.addTuple(bicycleSet, tuples[2]);
        cp.addTuple(bicycleSet, tuples[3]);
        cp.addTuple(bicycleSet, tuples[4]);
        cp.addTuple(bicycleSet, tuples[5]);
```

In the C# API, the tupleset is created using the method `CP.IntTable` and built as follows:

```
        IIntTupleSet bicycleSet = cp.IntTable(5);
        int[][] tuples = { new int [] {1, 57, 12, 3, 1490},
                           new int [] {2, 57, 13, 5, 1340},
                           new int [] {3, 60, 14, 3, 1790},
                           new int [] {4, 65, 14, 7, 1550},
                           new int [] {5, 67, 15, 2, 2070},
                           new int [] {6, 70, 15, 2, 1990}};
        cp.AddTuple(bicycleSet, tuples[0]);
        cp.AddTuple(bicycleSet, tuples[1]);
        cp.AddTuple(bicycleSet, tuples[2]);
        cp.AddTuple(bicycleSet, tuples[3]);
        cp.AddTuple(bicycleSet, tuples[4]);
        cp.AddTuple(bicycleSet, tuples[5]);
```

A tupleset can be used as an argument to a compatibility constraint in order to enforce the possible combinations allowed for a solution.

Piecewise linear functions

Piecewise linear functions are typically used in scheduling models to model a known function of time.

In CP Optimizer, piecewise linear functions are typically used in modeling a known function of time, for instance the cost that is incurred for completing an activity after a known date. A piecewise linear function is a function defined on an interval *[xmin, xmax)* which is partitioned into segments such that over each segment, the function is linear.

When two consecutive segments of the function are colinear, these segments are merged so that the function is always represented with the minimal number of segments.

In the **C++ API** of CP Optimizer, the interface `IloNumToNumSegmentFunction` represents piecewise linear functions.

In the **Java API** of CP Optimizer, the class `IloNumToNumSegmentFunction` represents piecewise linear functions.

In the **C# API** of CP Optimizer, the interface `INumToNumSegmentFunction` represents piecewise linear functions.

Note:

Piecewise linear function

A piecewise linear function is a function that is defined on an interval that is partitioned into segments such that over each segment, the function is linear.

Each interval *[x1, x2)* on which the function is linear is called a *segment*.

When two consecutive segments of the function are colinear, these segments are merged so that the function is always represented with the minimal number of segments.

Step functions

Step functions are typically used in scheduling models to model the efficiency of a resource over time.

In CP Optimizer, stepwise functions are typically used to model the efficiency of a resource over time. A stepwise function is a special case of piecewise linear function where all slopes are equal to 0 and the domain and image of the function are integer.

When two consecutive steps of the function have the same value, these steps are merged so that the function is always represented with the minimal number of steps.

In the **C++ API** of CP Optimizer, the class `IloNumToNumStepFunction` represents step functions.

In the **Java API** of CP Optimizer, the interface `IloNumToNumStepFunction` represents step functions.

In the **C# API** of CP Optimizer, the interface `INumToNumStepFunction` represents step functions.

Note:

Step function

Step functions are a special case of piecewise linear function where all slopes are equal to 0 and the domain and image of the function are integer.

Each interval *[x1, x2)* on which the function has the same value is called a *step*.

When two consecutive steps of the function have the same value, these steps are merged so that the function is always represented with the minimal number of steps.

Defining decision variables and expressions

Decision variables represent the unknown information in a problem. Expressions are built using decision variables, constants and other expressions.

Integer decision variables

Integer decision variables represent unknown information in a problem.

Integer decision variables represent unknown information in a problem and differ from normal programming variables in that they have domains of possible values and may have constraints placed on the allowed combinations of theses values.

These decision variables are defined by the lower bound and upper bound for the variable. In the **C++ API**, the following constructor creates a constrained integer variable with bounds -1 and 10:

```
IloIntVar myIntVar(env, -1, 10);
```

In the **Java API**, you create a constrained integer variable with bounds -1 and 10 with the method `IloCP.intVar`:

```
IloIntVar myIntVar = cp.intVar(-1, 10);
```

In the **C# API**, you create a constrained integer variable with bounds -1 and 10 with the method `CP.IntVar`:

```
IIntVar myIntVar = cp.IntVar(-1, 10);
```

In any solution, the value assigned by the optimizer to an integer decision variable must be between the lower and upper bounds, inclusive, of that variable. During the modeling stage, you may also modify the lower and upper bounds of a variable.

In the **C++ API** of CP Optimizer, the class `IloIntVar` represents integer decision variables.

In the **Java API** of CP Optimizer, the interface `IloIntVar` represents integer decision variables.

In the **C# API** of CP Optimizer, the interface `IIntVar` represents integer decision variables.

Note:

Integer decision variable

Integer decision variables represent unknown information in a problem.

Typically, the possible values are represented as a domain of integers with an upper bound and a lower bound. These variables are said to be constrained because constraints can be placed on them.

An interval variable can be specified as being optional, meaning that the interval can be absent from the solution schedule

Interval decision variables

Interval decision variables represent intervals of time whose positions in time are unknown.

An *interval decision variable* represents an unknown of a scheduling problem, in particular an interval of time during which something happens (an activity is carried out) whose position in time is unknown. An *interval* is characterized by a

start value, an end value and a size. The start and end of an interval variable must be in [IloIntervalMin..IloIntervalMax]. An important feature of interval decision variables is that they can be *optional*, that is, it is possible to model that an interval variable can be absent from the solution schedule.

Sometimes the intensity of "work" is not the same during the whole interval. For example, consider a worker who does not work during weekends (his work intensity during weekends is 0%) and on Friday he works only for half a day (his intensity during Friday is 50%). For this worker, 7 man-days work will span for longer than just 7 days. In this example 7 man-days represent what is called the *size* of the interval: that is, the length of the interval would be if the intensity function was always at 100%. To model such situations, a range for the size of an interval variable and an integer stepwise intensity function can be specified. The length of the interval will be at least long enough to cover the work requirements given by the interval size, taking into account the intensity function.

In the **C++ API** of CP Optimizer, the class `IloIntervalVar` represents interval decision variables.

In the **Java API** of CP Optimizer, the interface `IloIntervalVar` represents interval decision variables.

In the **C# API** of CP Optimizer, the interface `IIntervalVar` represents interval decision variables.

The domain of an interval variable is displayed as shown in this example:
`A1[0..1: 10..990 -- (5..10)5..990 --> 25..1000]`

After the name of the interval decision variable (here `A1`), the first range (here `0..1`) represents the domain of the Boolean presence status of the interval variable. The presence range `0..1` represents an optional interval variable whose status has still not been fixed, `0` an absent interval variable and `1` a present interval variable. An absent interval variable is displayed as:
`A1[0]`

When the interval variable is possibly present, the remaining fields describe the position of the interval variable:
- the first range in the remaining fields represents the domain of the interval start,
- the second range (between parenthesis) represents the domain of the interval size,
- the third range represents the domain of the interval length and
- the fourth and last range represents the domain of the interval end.

Note that the second range may be omitted in case the size and length of the interval variable are necessarily equal. When the values are fixed, the ranges `min..max` are replaced by a single value. For instance, the following display represents a fixed interval variable of size 5 that is present, starts at 10 and ends at 35:
`A1[1: 10 -- (5)25 --> 35]`

Note:

Interval decision variable

Interval decision variables represent unknowns of a scheduling problem

An interval has a start time, an end time, a size and a length. An interval variable allows for these values to be variable in the model.

Interval sequence decision variables

Interval decision sequence variables represent the orderings of sets of interval variables.

An *interval sequence decision variable* is defined on a set of interval variables The value of an interval sequence variable represents a total ordering of the interval variables of the set. Any absent interval variables are not considered in the ordering.

The sequence alone does not enforce any constraint on the relative position of interval end points. For instance, an interval variable *a* could be sequenced before an interval variable *b* in a sequence *p* without any impact on the relative position between the start/end points of *a* and *b* (*a* could still be fixed to start after the end of *b*). Different semantics can be used to define how a sequence constrains the positions of intervals. For instance, the no overlap constraint on a sequence variable specifies that the interval variables in the set do not overlap and that their start and end values are ordered according to the total ordering of the sequence.

In the **C++ API** of CP Optimizer, the class `IloIntervalSequenceVar` represents integer sequence decision variables.

In the **Java API** of CP Optimizer, the interface `IloIntervalSequenceVar` represents integer sequence decision variables.

In the **C# API** of CP Optimizer, the interface `IIntervalSequenceVar` represents integer sequence decision variables.

Note:

Interval sequence variable

Interval sequence decision variables represent a sequences of interval variables. A sequence can contain a subset of the variables or be empty. In a solution, the sequence will represent a total order over all the intervals in the set that are present in the solution. The assigned order of interval variables in the sequence does not necessarily determine their relative positions in time in the schedule.

Expressions

Expressions are constructed from decision variables and expressions.

Decision variables are most often used to build expressions, which in turn are used to create the objective and constraints of the model.

In the **C++ API** of IBM ILOG Concert Technology, expressions are generally represented by the class `IloExpr` and its subclasses `IloIntExpr`, representing integer expressions, and `IloNumExpr`, representing numeric (floating-point) expressions. Note that the parent classes `IloIntExprArg` and `IloNumExprArg` are used internally by Concert Technology to build expressions. You should not use `IloIntExprArg` or `IloNumExprArg` directly.

In the **Java API** of Concert Technology, expressions are generally represented by the interfaces IloIntExpr, representing integer expressions, and IloNumExpr, representing numeric expressions.

In the **C# API** of Concert Technology, expressions are generally represented by the interfaces IIntExpr, representing integer expressions, and INumExpr, representing numeric expressions.

An integer expression can be written explicitly; for example, using the **C++ API**:

```
1*x[1] + 2*x[2] + 3*x[3]
```

where x is assumed to be an array of decision variables. Using the **Java API**, the same expression could be written as:

```
cp.sum(cp.sum(cp.prod(1,x[1]),cp.prod(2,x[2])), cp.prod(3,x[3]))
```

Using the **C# API**, the same expression could be written as:

```
cp.Sum(cp.Sum(cp.Prod(1,x[1]),cp.Prod(2,x[2])), cp.Prod(3,x[3]))
```

Expressions can also be created incrementally, with a loop; for example using the **C++ API**:

```
IloIntExpr expr(env);
for (int i = 0; i < x.getSize(); ++i)
  expr += data[i] * x[i];
```

Using the **Java API**, the same expression could be written as:

```
IloIntExpr expr =  cp.constant(0);
for (int i = 0; i < x.length; ++i)
  expr = cp.sum(expr, cp.prod(data[i], x[i]));
```

Using the **C# API**, the same expression could be written as:

```
IIntExpr expr =  cp.Constant(0);
for (int i = 0; i < x.Length; ++i)
  expr = cp.Sum(expr, cp.Prod(data[i], x[i]));
```

In the **C++ API**, when you have finished using an expression (that is, you created a constraint or objective with it) you can delete it by calling its method end, for example:

```
expr.end();
```

Many Concert Technology functions return expressions. You can build even more complicated expressions by using the Concert Technology functions representing absolute value, integer division, maximum, minimum and others. For example, IloAbs(x) - IloMax(y,z) is a valid expression representing the difference between the absolute value of x and the maximum of y and z.

Note:

Expression

Values may be combined with decision variables and other expressions to form expressions.

To combine values, variables and other expressions to form an expression, you can use, among other functions, the operators:

- addition (+),

- subtraction (**-**),
- multiplication (__*__),
- division (**/**),
- self-assigned addition (**+=**) and
- self-assigned subtraction (**-=**).

Expressions are discussed in more detail in Chapter 4, "Constraints and expressions in CP Optimizer," on page 29.

Domains of variables and expressions

The set of possible values of a decision variable or expression is known as its domain.

An integer decision variable is typically created with a lower bound and an upper bound.

The initial *domain* of the variable contains precisely the values from the lower bound to the upper bound, inclusive. It is important to declare the initial domains to be as small as possible, in order to keep the search space as small as possible. As the optimizer works to find a solution for the problem, it will temporarily remove values from the domain of each decision variable. The search algorithms are discussed in detail in Chapter 6, "Search in CP Optimizer," on page 61.

During search, expressions have domains of possible values like decision variables. Unlike variables, these domains are not stored but instead calculated as needed from the basic elements of the expression.

Note:

Declaring a domain

Domains for integer decision variables are declared as part of the model. In order for the search to perform efficiently, it is important that the domains be as tight, or small, as possible.

Declaring the objective

An objective function is used to express the cost of each potential solution.

An *objective function* is used to express the cost of each potential solution. When an objective is added to a model, the problem is considered to be an optimization problem. The CP Optimizer search handles models with at most one objective function, though the modeling API provided by Concert Technology does not impose this restriction.

In the **C++ API** and the **Java API**, objective functions are represented by `IloObjective`. In the **C# API**, objectives are represented by `IObjective`. An objective function can be specified by creating an objective instance directly; however, it is common to use the functions which return an instance of `IloObjective`. In the **C++ API**, the functions are `IloMinimize` and `IloMaximize`. In the **Java API**, these functions are `IloModeler.minimize` and `IloModeler.maximize`. In the **C# API**, these functions are `IModeler.Minimize` and `IModeler.Maximize`. For example, in the **C++ API**, the following code defines an objective to minimize the expression `x[1] + 2*x[2] + 3*x[3]`:

```
IloObjective obj
  = IloMinimize(env,
                x[1] + 2*x[2] + 3*x[3]);
```

Using the **Java API**, the same objective could be written using the method
`IloCP.minimize`:

```
IloObjective obj
  = cp.minimize(cp.sum(cp.sum(x[1],cp.prod(2,x[2])),cp.prod(3,x[3])));
```

Using the **C# API**, the same objective could be written using the method
`CP.Minimize`:

```
IObjective obj
  = cp.Minimize(cp.Sum(cp.Sum(x[1],cp.Prod(2,x[2])),cp.Prod(3,x[3])));
```

Note:

Objective

An *objective* expresses the cost of possible solutions. The optimal solution to an
optimization problem is the feasible solution that, depending on the problem type,
minimizes or maximizes the cost.

After you create an objective, you must explicitly add it to the model in order for it
to be taken into account by the optimizer.

Adding constraints

Various types of constraints are available in Concert Technology for use in CP
Optimizer.

Constraint classes

Arithmetic and more complex constraints are available in Concert Technology.

IBM ILOG Concert Technology allows you to express constraints on decision
variables and expressions. In the **C++ API** and the **Java API**, constraints are
represented by `IloConstraint`. In the **C# API**, constraints are represented by the
`IConstraint` interface.

Arithmetic constraints

Arithmetic constraints are created by using arithmetic operators between
expressions and decision variables.

Arithmetic constraints can be created in a variety of ways.

In the C++ API, you can express constraints between expressions using the
following operators:
- equality (==),
- less than or equal to (<=),
- less than (<),
- greater than or equal to (>=),
- greater than (>) and
- not equal to (!=).

For example, you can write a constraint that one expression is not equal to another in the **C++ API**:

```
x[1] + 2*x[2] + 3*x[3] != 4*x[1]*x[2]
```

Using the **Java API**, the constraint could be written as:

```
cp.neq(cp.sum(x[1],
            cp.sum(cp.prod(2,x[2]), cp.prod(3,x[3]))) ,
                  cp.prod(4,cp.prod(x[1],x[2])))
```

Likewise, using the **C# API**, the constraint could be written as:

```
cp.Neq(cp.Sum(x[1],
            cp.Sum(cp.Prod(2,x[2]), cp.Prod(3,x[3]))) ,
                  cp.Prod(4,cp.Prod(x[1],x[2])))
```

Explicitly, this constraint enforces that the values the CP Optimizer search assigns to the decision variables x[1], x[2] and x[3] must be such that the expression x[1] + 2*x[2] + 3*x[3] does not equal the expression 4*x[1]*x[2].

After you create a constraint, you must explicitly add it to the model in order for it to be taken into account by the optimizer.

For more details on expressing constraints, refer to the section "Arithmetic constraints and expressions" on page 29.

Specialized constraints

Specialized constraints make it possible to express complicated relationships between variables.

One of the most powerful features of CP Optimizer is its provision for different kinds of constraints.

You have already seen arithmetic constraints that enable you to combine decision variables with the usual arithmetic operators (such as addition, subtraction, multiplication, division) to create linear and nonlinear expressions. In addition, CP Optimizer supports constraints on expressions that are not arithmetic: absolute value, minimum, maximum.

Specialized constraints make it possible to express complicated relations between variables, for example, relations that would require an exponential number of arithmetic constraints or relations that could not be expressed at all in arithmetic terms. Specialized constraints enter into such considerations as counting values, maintaining load weights and other critical activities.

Specialized constraints prove useful in a great many fields, in such mathematical puzzles as magic sequences and magic squares, but also in very practical real-world problems of allocation and scheduling. For example, the predefined "all different" constraint states that the value assigned to each variable in an array must be different from the values assigned to every other variable in that array.

Also, the single specialized all different constraint on n variables is logically equivalent to $(n-1)n/2$ instances of the "not equal" constraint (!=), but in terms of performance, that single constraint is much more efficient. In short, a specialized constraint such as is highly useful, both in terms of ease of expression and computational performance.

For detailed information on the specialized constraints that are available in CP Optimizer, refer to the section "Specialized constraints on integer decision variables" on page 33 and "Constraints and expressions on interval decision variables" on page 37 in Chapter 4, "Constraints and expressions in CP Optimizer," on page 29.

Combining constraints

Complex constraints may be created by combining simpler constraints using arithmetic or logical operators.

CP Optimizer provides a simple yet powerful way to define a constraint by combining other constraints.

This facility is based on the basic idea that constraints have value. The value of a constraint is true (`IloTrue` in C++) if the constraint is satisfied and false (`IloFalse`) if the constraint is not satisfied. Not only are they treated as if they have a Boolean value, such as true or false, but effectively as if the value is 0 or 1.

Note:

Logical constraint

A logical constraint is created by combining constraints or placing constraints on other constraints. Logical constraints are based on the idea that constraints have value. CP Optimizer handles constraints not only as if they have a Boolean value, such as true or false, but effectively as if the value is 0 or 1. This allows you to combine constraints into expressions or impose constraints on constraints.

You can use the following logical operators and constraints to combine constraints:
- not (!),
- and (&&),
- or (||),
- equivalence (==),
- exclusive or (!=) and
- implication (`IloIfThen`) (An instance of `IloIfThen` represents a condition constraint.).

Furthermore, constraints can be combined using arithmetic operators.

Some of the specialized constraints cannot be used in logical constraints.

For more details on expressing logical constraints, please refer to the section "Logical constraints" on page 32.

Note:

Constraint

Constraints specify the restrictions on the values that may be assigned to decision variables. To create a constraint for a model, you can:
- use an arithmetic operator between decision variables and expressions to return a constraint,
- use a function that returns a constraint,

- use a specialized constraint or
- use a logical operator between constraints which returns a constraint.

After you create a constraint, you must explicitly add it to the model in order for it to be taken into account by the CP Optimizer engine.

Constraints are discussed in more detail in Chapter 4, "Constraints and expressions in CP Optimizer," on page 29.

Formulating a problem

Modeling objects, such as objectives, variables and constraints, are contained in a model.

A model is a container of modeling objects, such as objectives, variables and constraints. You must explicitly add a constraint to the model or the CP Optimizer search engine will not use it in the search for a solution.

In the **C++ API** of CP Optimizer, a model is an instance of the class `IloModel`. Decision variables, objectives and constraints can be added to the model with the method `IloModel::add`.

In the **Java API** of CP Optimizer, the methods for adding object to the model are defined in the interfaces `IloModel` and `IloModeler` and implemented in the class `IloCP`. In particular, `IloCP.add` is one function for adding objects to the model.

Likewise, in the **C# API** of CP Optimizer, the methods for adding object to the model are defined in the interface `IModeler` and `IModelImpl` and implemented in the class `CP`. In particular, `CP.Add` is one function for adding objects to the model.

To create a model using the **C++ API**, the first step is to create an instance of the class `IloEnv`:

```
IloEnv env;
```

(Note that creating an environment is not necessary in the C# and the Java APIs.)

The initialization of the environment creates internal data structures to be used in the rest of the code. Once this is done, you can create model objects; here illustrated in the **C++ API**:

```
IloIntVar x(env, 0, 10);
IloConstraint ct = (x != 0);
```

In the **Java API**, you must create the `IloCP` object before creating the modeling objects:

```
IloCP cp = new IloCP();
IloIntVar x = cp.intVar(0, 10);
IloConstraint ct = cp.neq(x, 0);
```

Likewise in the **C# API**, you must create the `CP` object before creating the modeling objects:

```
CP cp = new CP();
IIntVar x = cp.IntVar(0, 10);
IConstraint ct = cp.Neq(x, 0);
```

Once this is done, you can create model and fill it with modeling objects; here illustrated in the **C++ API**:

```
IloModel model(env);
model.add(ct);
```

In the **Java API**, you have already created the `IloCP` object before creating the modeling objects, so adding the constraint to the model is simply:

```
cp.add(ct);
```

Likewise in the **C# API**, you have already created the `CP` object before creating the modeling objects, so adding the constraint to the model is simply:

```
cp.Add(ct);
```

As soon as the model is completed, you are ready to solve the problem. The processes for searching for solutions to a model are introduced in Chapter 6, "Search in CP Optimizer," on page 61.

When your problem solving is finished, you can reclaim memory for all modeling objects and clean up internal data structures by calling `IloEnv::end` for every environment you have created. This should be always be done before you exit your application.

To free memory used by a model in the **Java API**, you use the method `IloCP.end`. To free memory used by a model in the **C# API**, you use the method `CP.End`.

Note:

Model

A model is a container for modeling objects such as decision variables, objectives and constraints.

Chapter 4. Constraints and expressions in CP Optimizer

Constraints and expressions are building blocks that can be used to create models in CP Optimizer applications.

Arithmetic constraints and expressions

Arithmetic constraints and expressions may consist of linear and nonlinear arithmetic terms.

Arithmetic expressions

Arithmetic expressions consist of arithmetic terms that are combined by arithmetic operations.

An *arithmetic expression* consists of arithmetic terms, such as x, x^2, xy, or $3xy^2$, combined by arithmetic operations, such as addition, subtraction, multiplication, and division.

Arithmetic expressions can appear in arithmetic relations, such as equality and inequality. Arithmetic expressions can be linear (such as $x < z + 3y$) or nonlinear (such as $x < z^2 + y^*z + 3y$). CP Optimizer supports arithmetic constraints over integer decision variables. It also supports integer or numeric (floating-point) expressions.

Note:

CP Optimizer supports integer decision variables. Moreover, it is possible to constrain floating-point expressions. It is also possible to use a floating-point expression as a term in an objective function. For example, a floating-point cost function can be calculated from expressions by using integer decision variables, like this: `cost = x/1000 + y/3`.

There are no inherent restrictions on the magnitude of arithmetic operations that you can perform with CP Optimizer expressions and decision variables. The only limitation that you must bear in mind in your CP Optimizer application is any possibility of overflow due to the size and configuration of your platform (that is, limits on hardware and operating system). For example, if you multiply the largest possible integer available on your platform by the largest possible integer, you risk overflow because of limitations of your platform.

In the **C++ API** of CP Optimizer, there are overloaded operators for building arithmetic expressions and stating constraints over them. For many arithmetic operations, such as addition, exponentiation, or modular arithmetic, there are global functions that offer a variety of signatures that accommodate many combinations of decision variables and integer or numeric values in expressions.

In the **Java API** of CP Optimizer, there are methods to enable you to build arithmetic expressions to state constraints. Many such methods are defined in the interface `IloModeler`, implemented in the class `IloCP`.

Likewise, in the **C# API** of CP Optimizer, there are similar methods for building arithmetic expressions and stating constraints over them. The interface `IModeler` defines these methods, and the class `CP` inherits them.

An *arithmetic constraint* involves one or more decision variables in an arithmetic expression.

Table 1 summarizes the methods that support arithmetic operations to create expressions to use in constraints. Table 2 summarizes the methods and functions that return expressions for use in constraints. Table 3 on page 31 summarizes the methods that implement arithmetic constraints. The reference manuals of the application programming interfaces (APIs) document each method, operator or global function more fully.

In those tables, the names of applications where the expressions or constraints are used appear in parentheses, like this (`sports.cpp`). The applications are part of the product, available in the appropriate subdirectory of the `examples` directory of your installation.

Table 1. Arithmetic operations for use in constraints

Arithmetic operation	C++ API	Java API	C# API

addition	`operator+`
	
	`IloSum`
	
	`(sports.cpp)`
		`IloModeler.sum`
		
		`(Facility.java)`
			`IModeler.Sum`
			
			`(Steelmill.cs)`

subtraction	`operator-`
	
	`IloDiff`
		`IloModeler.diff`
		
		`(Sports.java)`
			`IModeler.Diff`
			
			`(Sports.cs)`

multiplication	`operator*`	`IloModeler.prod`
		
		`(Facility.java)`
			`IModeler.Prod`
			
			`(Facility.cs)`

scalar product	`IloScalProd`
	
	`(facility.cpp)`
		`IloModeler.scalProd`	`IModeler.ScalProd`

integer division	`IloDiv`
	
	`(sports.cpp)`
		`IloCP.div`
		
		`(Sports.java)`
			`CP.Div`
			
			`(Sports.cs)`

floating-point division	`operator/`	`IloCP.quot`	`CP.Quot`

modular arithmetic	`operator%`
	
	`(sports.cpp)`
		`IloCP.modulo`
		
		`(Sports.java)`
			`CP.Modulo`
			
			`(Sports.cs)`

Table 2. Arithmetic expressions for use in constraints

Expression	C++ API	Java API	C# API
standard deviation	`IloStandardDeviation`	`IloCP.standardDeviation`	`CP.StandardDeviation`

minimum	`IloMin`
	
	`(talent.cpp)`
		`IloModeler.min`
		
		`(Sports.java)`
			`IModeler.Min`
			
			`(Sports.cs)`

maximum	`IloMax`	`IloModeler.max`	`IModeler.Max`

Table 2. Arithmetic expressions for use in constraints (continued)

Expression	C++ API	Java API	C# API
counting	IloCount (teambuilding.cpp) IloCountDifferent (alloc.cpp)	IloCP.count (Facility.java) IloCP.countDifferent (Alloc.java)	CP.Count (Sports.cs) CP.CountDifferent (Alloc.cs)
absolute value	IloAbs (sports.cpp)	IloModeler.abs (Sports.java)	IModeler.Abs (Sports.cs)
element or index	IloElement	IloCP.element (Facility.java)	CP.Element (Steelmill.cs)

Table 3. Arithmetic constraints

Arithmetic constraint	C++ API	Java API	C# API
equal to	operator==	IloModeler.eq	IModeler.Eq
not equal to	operator!=	IloCP.neq	CP.Neq
strictly less than	operator<	IloCP.lt	CP.Lt
strictly greater than	operator>	IloCP.gt	CP.Gt
less than or equal to	operator<=	IloModeler.le	IModeler.Le
greater than or equal to	operator>=	IloModeler.ge	IModeler.Ge

Element expressions

Element expressions index arrays by decision variables.

An element expression indexes an array of values by a decision variable. In other words, it allows you to treat the index into an array as a decision variable, as something that can be constrained.

This kind of constrained indexing is conventionally expressed in other disciplines by a very large number of binary (0-1) or Boolean variables. However, in CP Optimizer this indirect relation can be stated straight forwardly and efficiently as an expression that can appear in constraints.

For example, in a logistics model you might use an element expression in a constraint to index products available in warehouses. Or in a sports scheduling application, your model might include two arrays, both listing all teams. An element expression might appear in constraints between the two arrays to enforce which teams play each other.

In the **C++ API** of CP Optimizer, the global function IloElement has many signatures that enable you to build element expressions. The index operator, operator[], allows you to index an array of integers or constrained integer variables using a constrained integer variable.

In the **Java API** of CP Optimizer, the methods IloCP.element allow you to build element expressions. See the example Steelmill.java for a sample of this expression in an application.

Likewise, in the **C# API** of CP Optimizer, the methods `CP.Element` represent element expressions. For samples of this expression, see `Teambuilding.cs`, `Facility.cs`, or `Steelmill.cs`.

Division expression examples

Both integer division and floating-point division expressions can be modeled in arithmetic expressions.

Both integer division and floating-point division can be modeled in arithmetic expressions. To clarify the distinction between integer division and floating-point division of decision variables in constraints, consider these brief examples.

Let x be a decision variable that can assume the value 9:
- In the **C++ API**, `IloIntVar x(env, 9, 9, "x");`
- In the **Java API**, `IloIntVar x = cp.intVar(9, 9);`
- In the **C# API**, `IIntVar x = cp.IntVar(9, 9);`

Then **integer division** with `IloDiv`, or `IloCP.div` or `CP.Div` of x divided by 10 results in the value 0 (zero).
- Consequently, a constraint stating "`IloDiv(x, 10) == 0`" is true.
- In contrast, a constraint stating "`IloDiv(x, 10) >= 0.5`" is false.

However, **floating-point division** with `operator/`, or `IloCP.quot` or `CP.Quot` of x divided by 10 results in the value 0.9.
- Consequently, a constraint stating "`x/10 == 0`" is false.
- A constraint stating "`x/10 >= 0.5`" is true.

Logical constraints

Logical constraints make it possible to model complicated logical relations.

Logical constraints enforce logical relations. Logical relations include:
- logical-and (conjunction); for example, a warehouse in a logistics application must be in the same region **and** stock the product;
- logical-or (disjunction); for example, a crew in a rostering application must include a nurse **or** an emergency medical technician;
- logical-not (negation); for example, total cost must **not** exceed budget;
- if-then (implication); for example, **if** inventory is below threshold, **then** re-order.

A logical constraint makes it possible to express complicated relations between decision variables, relations that would require an impractical number of arithmetic constraints or relations that could not be expressed at all in arithmetic terms. Table 4 summarizes the logical constraints available in CP Optimizer.

Table 4. Logical constraints in CP Optimizer

Logical constraint	C++ API	Java API	C# API

logical-and (conjunction)	`IloAnd`
	 `operator&&`
	 `(facility.cpp)`
		`IloModeler.and`	`IModeler.And`

Table 4. Logical constraints in CP Optimizer (continued)

Logical constraint	C++ API	Java API	C# API

logical-or (disjunction)	`IloOr`
	
	`operator||`
	
	`(teambuilding.cpp)`
		`IloModeler.or`
		
		`(Teambuilding.java)`
			`IModeler.Or`
			
			`(Teambuilding.cs)`

logical-not (negation)	`operator!`
	
	`IloNot`
	
	`(talent.cpp)`
		`IloModeler.not`	`IModeler.Not`

logical if-then (implication)	`IloIfThen`	`IloModeler.ifThen`
		
		`IloCP.ifThenElse`
			`IModeler.IfThen`

Compatibility constraints

Compatibility constraints enforce allowed or forbidden combinations of values among decision variables.

CP Optimizer offers two predefined constraints that enforce compatibility among decision variables. These compatibility constraints make it easier for you to declare allowed or forbidden values that a decision variable or a group of decision variables may assume. These constraints can apply to any number of variables simultaneously.

You use a tuple (that is, an ordered set, such as an array) to declare allowed or forbidden values. This tuple has arity equal to the number of decision variables under consideration. That is, the number of fields in the tuple is the same as the number of decision variables that the constraint applies to. Each tuple declares an allowed or forbidden combination.

In the **C++ API** of CP Optimizer, the global functions `IloAllowedAssignments` and `IloForbiddenAssignments` return compatibility constraints. For samples of these constraints, see `teambuilding.cpp` or `sports.cpp`.

In the **Java API** of CP Optimizer, the methods `IloCP.allowedAssignments` and `IloCP.forbiddenAssignments` represent compatibility constraints. For samples of these constraints, see `Teambuilding.java` or `Sports.java`.

In the **C# API** of CP Optimizer, the methods `CP.AllowedAssignments` and `CP.ForbiddenAssignments` represent compatibility constraints. For samples of these constraints, see `Teambuilding.cs` or `Sports.cs`.

Specialized constraints on integer decision variables

Specialized constraints on integer decision variables are designed and implemented to reduce domains of variables efficiently during a search.

Overview

Specialized constraints on integer decision variables are designed and implemented to reduce domains of variables efficiently during a search.

Theoretically, these constraints can be written from elementary arithmetic constraints and expressions, but CP Optimizer offers you these specialized constraints, ready to use in your application. They have been designed and implemented to reduce domains of variables efficiently during a search.

All different constraint

The all different constraint is a specialized constraint which forces every decision variable in a given group to assume a value different from the value of every other decision variable in that group.

An all different constraint forces every decision variable in a given group to assume a value different from the value of every other decision variable in that group. In other words, no two of those decision variables will have the same integer value when this constraint is satisfied.

For example, if the decision variables x, y and z share the domain of integers from 1 through 10, inclusive, then an all different constraint applied to x, y and z could return such solutions as 1, 2, 3 or 4, 6, 8 or 9, 5, 2. However, 1,1, 2 would not be a solution to an all different constraint over those three variables.

In the **C++ API** of CP Optimizer, the class `IloAllDiff` represents all different constraints. For samples of this constraint, see `teambuilding.cpp` or `sports.cpp`.

In the **Java API** of CP Optimizer, the method `IloCP.allDiff` represents all different constraints. For samples of this constraint, see `Teambuilding.java` or `Sports.java`.

In the **C# API** of CP Optimizer, the method `CP.AllDiff` represents all different constraints. For samples of this constraint, see `Teambuilding.cs` or `Sports.cs`.

Minimum distance constraint

The minimum distance constraint is a specialized constraint which forces every decision variable in a given group to assume values which differ by at least the specified amount.

A minimum distance constraint resembles an all different constraint in that it accepts a group of decision variables (such as an array of integer decision variables) and ensures that they differ from each other. In fact, it makes sure that they differ from one another by a given amount, so that they are at least a certain distance apart from each other.

For example, the decision variables x, y and z share the domain of integers from 1 through 10, inclusive, and the minimum distance between them should be 3, then a minimum distance constraint applied to x, y and z with distance 3 could return such solutions as 1, 4, 8 or 7, 10, 4. However, 2, 4, 6 would not be an acceptable solution because y would be closer to x or z than the stipulated distance.

In the **C++ API** of CP Optimizer, the class `IloAllMinDistance` represents a minimum distance constraint.

In the **Java API** of CP Optimizer, the method `IloCP.allMinDistance` represents a minimum distance constraint.

In the **C# API** of CP Optimizer, the method `CP.AllMinDistance` represents a minimum distance constraint.

Packing constraint

The packing constraint is a specialized constraint which maintains the load of a group of containers or bins, given a group of weighted items and an assignment of items to containers.

A packing constraint maintains the load of a group of containers or bins, given a group of weighted items and an assignment of items to containers.

To understand what a packing constraint can do, consider n items and m containers. Suppose the containers are numbered consecutively from zero. Each item i has an integer weight w_i. Each item can go in only one container, and each item has associated with it a constrained integer variable p_i to specify in which container item i is placed. Associated with each container j is an integer variable l_j representing the load in that container. The load is the sum of the weights of the items which have been assigned to that container. A container is used if at least one item is placed in that container.

Given this scenario, you can use a packing constraint to:
* You can set the capacity of each container. In other words, you can place an upper bound on the load variable.
* With a packing constraint, you can also make sure that the total sum of the loads of the containers is equal to the sum of the weights of the items being placed.
* You can specify the number of containers used. In fact, you can specify a particular group of containers to use.

In the **C++ API** of CP Optimizer, the class `IloPack` returns a packing constraint. For samples of this constraint, see `ppp.cpp` or `steelmill.cpp`.

In the **Java API** of CP Optimizer, the method `IloCP.pack` represents a packing constraint. For samples of this constraint, see `Ppp.java` or `Steelmill.java`.

In the **C# API** of CP Optimizer, the method `CP.Pack` represents a packing constraint. For samples of this constraints, see `Ppp.cs` or `Steelmill.cs`.

Inverse constraint

The inverse constraint is a specialized constraint which forces two groups of decision variables to be in strict correspondence with each other.

An inverse constraint forces two groups of decision variables to be in strict correspondence with each other. The two groups of decision variables are conventionally represented by two arrays. An inverse constraint makes sure that the i-th item of one array is the inverse of the i-th item of the other array, and vice versa.

In formal terms, if the length of the arrays `f` and `invf` is `n`, then the inverse constraint guarantees that:
* for all `i` in the interval `[0,n-1]`, then `invf[f[i]]==i`;
* for all `j` in the interval `[0,n-1]`, then `f[invf[j]]==j`.

In the **C++ API** of CP Optimizer, the class `IloInverse` represents inverse constraints. For samples of this constraint, see `sports.cpp` or `talent.cpp`.

In the **Java API** of CP Optimizer, the method `IloCP.inverse` represents inverse constraints. For a sample of this constraint, see `Sports.java`.

Likewise, in the **C# API** of CP Optimizer, the method `CP.Inverse` represents inverse constraints. For a sample of this constraints, see `Sports.cs`.

Lexicographic constraint

The lexicographic constraint is a specialized constraint which maintains order between two groups of expressions.

A lexicographic constraint maintains order between two groups of expressions. The two groups of expressions are conventionally specified in two arrays that must be the same length. A lexicographic constraint compares the expressions, index by index, to see which is greater in lexicographic terms.

Informally, you might think of a lexicographic constraint as putting two arrays in alphabetic order relative to each other.

In formal terms, a lexicographic constraint over two arrays x and y maintains that x is less than or equal to y in the lexicographical sense of the term. That is, either both arrays are equal, index by index, or there exists an index i strictly less than the length of the array x such that for all $j < i$, $x[j] = y[j]$ and $x[i] < y[i]$.

In the **C++ API** of CP Optimizer, the global function `IloLexicographic` returns a lexicographic constraint.

In the **Java API** of CP Optimizer, the method `IloCP.lexicographic` returns a lexicographic constraint.

Likewise, in the **C# API** of CP Optimizer, the method `CP.Lexicographic` returns a lexicographic constraint.

Distribution constraint

The distribution constraint is a specialized constraint which counts the number of occurrences of several values among the decision variables in an array of decision variables.

A distribution constraint is a counting constraint used to count the number of occurrences of several values among the constrained variables in an array of constrained variables. You can also use an instance of this class to force the constrained variables of an array to assume values in such a way that only a limited number of the constrained variables assume each value.

To understand what a distribution constraint can do, consider an array of decision variables, called `vars`. Given an array of integers `values` of length n called `values` and an array of the same length of decision variables called `cards`. The i-th element of the array `cards` represents the number of times the i-th element of `values` appears in the array of decision variables `vars`.

When an instance of this class is created by a constructor with no "values" argument), then the array of values that are being counted must be an array of consecutive integers starting with 0 (zero). In that case, for each i, the i-th element of `cards` is equal to the number of occurrences of i in the array `vars`.

In the **C++ API** of CP Optimizer, the class `IloDistribute` represents distribution constraints.

In the **Java API** of CP Optimizer, the method `IloCP.distribute` represents distribution constraints.

Likewise, in the **C# API** of CP Optimizer, the method `CP.Distribute` represents distribution constraints.

Constraints and expressions on interval decision variables

CP Optimizer provides specialized expressions and constraints on interval decision variables.

Expressions on interval decision variables

CP Optimizer provides specialized expressions on interval decision variables.

Numerical expressions can be created from interval decision variables. These expressions can be used to:
- define a term for the cost function and
- connect interval variables to integer and floating point expressions.

CP Optimizer provides access to the integer expressions that represent the start of, end of, length of and size of an interval variable. Special care must be taken in the case of optional intervals: in this case an integer value *absVal* must be specified which represents the value of the expression when the interval is absent. If this value is omitted, it is assumed to be 0.

Numerical expressions are also provided that allow you to evaluate a piecewise linear function on these four attributes of an interval variable. A numerical value *absVal* can be specified that represents the value of the expression when the interval is absent. If this value is omitted, it is assumed to be 0.

In the **C++ API** of CP Optimizer, the global functions `IloStartOf`, `IloEndOf`, `IloLengthOf` and `IloSizeOf` return integer expressions representing the various attributes of an interval decision variable. The global functions `IloStartEval`, `IloEndEval`, `IloLengthEval` and `IloSizeEval` return numerical expressions representing the value of the given function evaluated on the appropriate attribute the given of interval variable whenever the interval variable is present.

In the **Java API** of CP Optimizer, the methods `IloCP.startOf`, `IloCP.endOf`, `IloCP.lengthOf` and `IloCP.sizeOf` return integer expressions representing the various attributes of an interval decision variable. The global functions `IloCP.startEval`, `IloCP.endEval`, `IloCP.lengthEval` and `IloCP.sizeEval` return numerical expressions representing the value of the given function evaluated on the appropriate attribute the given of interval variable whenever the interval variable is present.

In the **C# API** of CP Optimizer, the methods `CP.StartOf`, `CP.EndOf`, `CP.LengthOf` and `CP.SizeOf` return integer expressions representing the various attributes of an interval decision variable. The global functions `CP.StartEval`, `CP.EndEval`, `CP.LengthEval` and `CP.SizeEval` return numerical expressions representing the value of the given function evaluated on the appropriate attribute the given of interval variable whenever the interval variable is present.

Forbidden values constraints

Interval variables can be constrained to not start, end or overlap a set of fixed dates

It may be necessary to state that an interval cannot start, cannot end or cannot overlap a set of fixed dates. CP Optimizer provides constraints for modeling these requirements. Given an interval variable and an integer stepwise function, the forbidden start constraint states that whenever the interval is present, it cannot start at a point where the function evaluates to zero. Given an interval variable a and an integer stepwise function F, the forbidden end constraint states that whenever the interval is present, it cannot end at a value t where $F(t-1) = 0$. Given an interval variable and an integer stepwise function, the forbidden extent constraint states that whenever the interval is present, it cannot overlap a point where the function evaluates to zero.

In the **C++ API** of CP Optimizer, the global functions `IloForbidStart`, `IloForbidEnd` and `IloForbidExtent` return constraints representing restrictions on when an interval can start, end or be placed.

In the **Java API** of CP Optimizer, the methods `IloCP.forbidStart`, `IloCP.forbidEnd` and `IloCP.forbidExtent` return constraints representing restrictions on when an interval can start, end or be placed.

In the **C# API** of CP Optimizer, the methods `CP.ForbidStart`, `CP.ForbidEnd` and `CP.ForbidExtent` return constraints representing restrictions on when an interval can start, end or be placed.

Precedence constraints on interval variables

Interval variables can be constrained to have relative temporal position.

This section describes common constraints in scheduling, namely, precedence constraints. These constraints restrict the relative temporal position of interval variables in a solution. For instance a precedence constraint can model the fact that a given activity must end before another given activity starts (optionally with some minimum delay). If one or both of the interval variables of the precedence constraint is absent, then the precedence is systematically considered to be true and thus, it does not impact the schedule.

Note:

Precedence constraint

Precedence constraints are used to specify when one interval variable must start or end with respect to the start or end time of another interval. The following types of precedence constraints are available; if a and b denote interval variables, both interval variables are present and `delay` is a number or integer expression (0 by default), then:

- `endBeforeEnd(a, b, delay)` constrains that at least the given delay should elapse between the end of a and the end of b. It imposes the inequality `end(a) + delay <= end(b)`
- `endBeforeStart(a, b, delay)` constrains that at least the given delay should elapse between the end of a and the start of b. It imposes the inequality `end(a) + delay <= start(b)`.

- endAtEnd(a, b, delay) constrains the given delay to separate the end of a and the end of b. It imposes the equality end(a) + delay == end(b).
- endAtStart(a, b, delay) constrains the given delay to separate the end of a and the start of b. It imposes the equality end(a) + delay == start(b).
- startBeforeEnd(a, b, delay) constrains that at least the given delay should elapse between the start of a and the end of b. It imposes the inequality start(a) + delay <= end(b).
- startBeforeStart(a, b, delay) constrains that at least the given delay should elapse between the start of a and the start of b. It imposes the inequality start(a) + delay <= start(b).
- startAtEnd(a, b, delay) constrains the given delay to separate the start of a and the end of b. It imposes the equality start(a) + delay == end(b).
- startAtStart(a, b, delay) constrains the given delay to separate the start of a and the start of b. It imposes the equality start(a) + delay == start(b).

If either interval a or b is not present in the solution, the constraint is automatically satisfied, and it is as if the constraint was never imposed.

In the **C++ API** of CP Optimizer, the global functions IloEndBeforeEnd, IloEndBeforeStart, IloEndAtEnd, IloEndAtStart, IloStartBeforeEnd, IloStartBeforeStart, IloStartAtEnd and IloStartAtStart return precedence constraints.

In the **Java API** of CP Optimizer, the methods IloCP.endBeforeEnd, IloCP.endBeforeStart, IloCP.endAtEnd, IloCP.endAtStart, IloCP.startBeforeEnd, IloCP.startBeforeStart, IloCP.startAtEnd and IloCP.startAtStart return precedence constraints.

In the **C# API** of CP Optimizer, the methods CP.EndBeforeEnd, CP.EndBeforeStart, CP.EndAtEnd, CP.EndAtStart, IloCP.startBeforeEnd, IloCP.startBeforeStart, IloCP.startAtEnd and IloCP.startAtStart return precedence constraints.

Logical constraints on interval variables

Logical constraints on interval variables may be modeled using presence constraints.

The presence status of interval variables can be restricted by logical constraints. The presence constraint states that a given interval variable must be present. Of course, this constraint may be used in logical constraints. For example, given two optional intervals *a* and *b*, a logical constraint could be that if interval *a* is present then b must be present too. This can be modeled using the presence constraint.

In the **C++ API** of CP Optimizer, the global function IloPresenceOf returns a constraint that states that a given interval must be present.

In the **Java API** of CP Optimizer, the method IloCP.presenceOf returns a constraint that states that a given interval must be present.

In the **C# API** of CP Optimizer, the method CP.PresenceOf returns a constraint that states that a given interval must be present.

Isomorphism constraints on interval variables

The isomorphism constraint constrains two arrays of interval variables.

The isomorphism constraint "synchronizes" the presence of each index pair of interval variables in the two arrays. The isomorphism constraint is used when a set of operations is subject to two independent scheduling models.

Consider two sets of interval variables, *A* and *B*. This constraint states that, in a solution, there is a 1-to-1 correspondence between the present intervals of *A* and the present intervals of *B*; two intervals in correspondence have the same start and end values. The isomorphism constraint is useful to enforce some patterns of constraints on a set of interval variables

In the **C++ API** of CP Optimizer, the class `IloIsomorphism` represents an isomorphism constraint between two interval variable arrays.

In the **Java API** of CP Optimizer, the method `IloCP.isomorphism` returns an isomorphism constraint between two interval variable arrays.

In the **C# API** of CP Optimizer, the method `CP.Isomorphism` returns an isomorphism constraint between two interval variable array.

Constraints on groups of interval variables

CP Optimizer provides specialized constraints on groups of interval variables.

Constraints over groups of intervals allow hierarchical creation of the model by "encapsulating" a group of interval variables by one "high level" interval.

The spanning constraint states that a given interval spans over all present intervals from a given set. That is, the given interval starts together with the first present interval from the set and ends together with the last one.

In the **C++ API** of CP Optimizer, the class `IloSpan` represents a spanning constraint.

In the **Java API** of CP Optimizer, the method `IloCP.span` returns a spanning constraint.

In the **C# API** of CP Optimizer, the method `CP.Span` returns a spanning constraint.

The alternative constraint models an exclusive alternative among a set of intervals. If a given interval *a* is present then exactly one of the intervals of the set is present and interval *a* starts and ends together with this chosen one.

In the **C++ API** of CP Optimizer, the class `IloAlternative` represents an alternative constraint.

In the **Java API** of CP Optimizer, the method `IloCP.alternative` returns an alternative constraint.

In the **C# API** of CP Optimizer, the method `CP.Alternative` returns an alternative constraint.

The synchronization constraint makes the intervals in the given set start and end together with an interval *a*, if *a* is present.

In the **C++ API** of CP Optimizer, the class `IloSynchronize` represents a synchronization constraint.

In the **Java API** of CP Optimizer, the method `IloCP.synchronize` returns a synchronization constraint.

In the **C# API** of CP Optimizer, the method `CP.Synchronize` returns a synchronization constraint.

Sequence constraints on interval variables and interval sequence variables

Groups of interval variables may be constrained to have relative positions of the start and end points of the interval variables.

An interval sequence variable alone does not enforce any constraint on the relative position of the end points of the interval variables. For instance, an interval variable *a* could be sequenced before an interval variable *b* in a sequence *p* without any impact on the relative position between the start/end points of *a* and *b* (*a* could still be fixed to start after the end of *b*). Different semantics can be used to define how a sequence constrains the positions of intervals.

There are four sequencing constraints available on sequence variables. The first in sequence constraint on an interval variable and an interval sequence variable states that if the interval variable is present, then it will be the first interval of the sequence. The last in sequence constraint on an interval variable and an interval sequence variable states that if an interval variable is present, it will be the last interval of the sequence. The before constraint on an interval sequence variable and two interval variable states that if both interval variables are present, then the first interval variable will appear in the sequence before the other interval. The previous constraint on an interval sequence variable and two interval variable states that if both interval variables are present then the first interval variable will be just before the other in the sequence, that is, it will appear before the second and no other interval will be sequenced between the two intervals.

In the **C++ API** of CP Optimizer, the functions `IloFirst, IloLast, IloBefore` and `IloPrev` represent the sequencing constraints.

In the **Java API** of CP Optimizer, the methods `IloCP.first, IloCP.last, IloCP.before` and `IloCP.prev` represent the sequencing constraints.

In the **C# API** of CP Optimizer, the methods `CP.First, CP.Last, CP.Before` and `CP.Prev` represent the sequencing constraints.

The no overlap constraint on an interval sequence variable *p* states that the sequence defines a chain of non-overlapping intervals, any interval in the chain being constrained to end before the start of the next interval in the chain. This constraint is typically useful for modeling disjunctive resources.

A transition distance matrix can be specified, which defines the minimal distance that must separate two consecutive intervals in the sequence.

In the **C++ API** of CP Optimizer, the class `IloNoOverlap` represents a no overlap constraint.

In the **Java API** of CP Optimizer, the method `IloCP.noOverlap` returns a no overlap constraint.

In the **C# API** of CP Optimizer, the method `CP.NoOverlap` returns a no overlap constraint.

Constraints and expressions on cumulative (cumul) function expressions

Cumulative usage of renewable resources can be represented by cumulative function expressions.

Cumul function expressions

Cumulative function expressions represent the cumulative usage of renewable resources. .

In scheduling problems involving cumulative resources (also known as renewable resources), the cumulated usage of the resource by the activities is usually represented by a function of time. An activity usually increases the cumulated resource usage function at its start time and decreases it when it releases the resource at its end time (pulse function). For resources that can be produced and consumed by activities (for instance the content of an inventory or a tank), the resource level can also be described as a function of time, production activities will increase the resource level whereas consuming activities will decrease it. In these type of problems, the cumulated contribution of activities on the resource can be represented by a function of time and constraints can be posted on this function, for instance a maximal or a safety level.

CP Optimizer introduces the notion of cumulative (cumul) function expression which is a function that represents the sum of individual contributions of intervals. A family of elementary cumul function expressions describes the individual contribution of an interval variable (or a fixed interval of time). These expressions cover the main use-cases mentioned previously: pulse for usage of a cumulative resource, step for resource production/consumption. When the elementary cumul functions that define a cumul function are fixed (and thus, so are their related intervals), the expression is fixed. CP Optimizer provides several constraints over cumul functions. These constraints allow restriction of the possible values of the function over the complete horizon or over some fixed or variable interval. For applications where the actual quantity of resource that is used, produced or consumed by intervals is an unknown of the problem, expressions are available for constraining these quantities.

In the **C++ API** of CP Optimizer, the class `IloCumulFunctionExpr` represents a cumul function expression.

In the **Java API** of CP Optimizer, the method `IloCP.cumulFunctionExpr` returns a cumul function expression.

In the **C# API** of CP Optimizer, the method `CP.CumulFunctionExpr` returns a cumul function expression.

Elementary cumul function expressions

Elementary cumul functions represent the effects of an interval variable on a cumul function.

The effects of an interval variable on a cumul function are modeled using elementary cumul functions. An elementary pulse function increases or decreases the level of the function by a given amount for the length of a given interval variable or fixed interval. An elementary step function increases or decreases the

level of the function by a given amount at a given time. An elementary step at start function increases or decreases the level of the function by a given amount at the start of a given interval variable. An elementary step at end function increases or decreases the level of the function by a given amount at the end of a given interval variable.

In the **C++ API** of CP Optimizer, the functions `IloPulse`, `IloStep`, `IloStepAtStart` and `IloStepAtEnd` return elementary cumul functions.

In the **Java API** of CP Optimizer, the methods `IloCP.pulse`, `IloCP.step`, `IloCP.stepAtStart` and `IloCP.stepAtEnd` return elementary cumul functions.

In the **C# API** of CP Optimizer, the methods `CP.Pulse`, `CP.Step`, `CP.StepAtStart` and `CP.StepAtEnd` return elementary cumul functions.

Expressions on cumul function expressions

The height of cumul functions is represented by a family of integer expressions.

The actual height of a cumul function may be an unknown of the problem. CP Optimizer provides a family of integer expressions to control this height. Some of these expressions define a range of possible values for the actual height of the function when a given interval variable is present. The actual height is an unknown of the problem, and some integer expressions to control this height are provided.

The height at start expression returns the value of the contribution of the given interval *a* to the given cumul function evaluated at the start of *a* that is, it measures the contribution of interval *a* to the cumul function at its start point. Similarly, the height at end expression returns the value of the contribution of the given interval *a* to the given cumul function evaluated at the end of *a*. An additional integer value *absVal* can be specified at the construction of the expression in which case that will be the value returned by the expression when the interval is absent otherwise, if no value is specified, the expression will be equal to 0 when the interval is absent.

In the **C++ API** of CP Optimizer, the functions `IloHeightAtStart` and `IloHeightAtEnd` return integer expressions representing the total contribution of the start or and of an interval variable to a cumul function.

In the **Java API** of CP Optimizer, the methods `IloCP.heightAtStart` and `IloCP.heightAtEnd` return integer expressions representing the total contribution of the start or and of an interval variable to a cumul function.

In the **C# API** of CP Optimizer, the methods `CP.HeightAtStart` and `CP.HeightAtEnd` return integer expressions representing the total contribution of the start or and of an interval variable to a cumul function.

Constraints on cumul function expressions

Cumul functions can be constrained to evaluate within a certain range of values for a fixed interval or an interval variable.

Cumul functions can be constrained to evaluate within a certain range of values for a fixed interval or an interval variable. A simple less than or equal constraint on inequality constraint on a cumul function states that the domain of values of the cumul function cannot exceed the given value anywhere along the temporal axis. An "always equal" range constraint on a cumul function states that the cumul

expression function must evaluate to a given value everywhere on a fixed interval or from the start to end of a given interval variable. An "always in" range constraint on a cumul function states that the domain of the cumul expression function must evaluate to within a certain range at each point within a fixed interval or from the start to end of a given interval variable.

Range constraints on cumul functions cannot be used in logical constraints of CP Optimizer, because any logical constraint involving interval variables must be captured via the presence Boolean value on the interval handled by the presence constraint.

In the **C++ API** of CP Optimizer, the functions `IloAlwaysIn`, `IloAlwaysEqual`, `operator<=` and `operator>=` return range constraints on cumul functions.

In the **Java API** of CP Optimizer, the methods `IloCP.alwaysIn`, `IloCP.alwaysEqual`, `IloCP.le` and `IloCP.ge` return range constraints on cumul functions.

In the **C# API** of CP Optimizer, the methods `CP.AlwaysIn`, `CP.AlwaysEqual`, `CP.Le` and `CP.Ge` return range constraints on cumul functions.

Constraints on state functions

The evolution of a state function may be constrained using specialized constraints.

A state function is a decision variable whose value is a set of non-overlapping intervals over which the function maintains a particular non-negative integer state. In between those intervals, the state of the function is not defined, typically because of an ongoing transition between two states.

A set of constraints is available to restrict the evolution of a state function. These constraints allow you to specify that:
* the state of the function must be defined and should remain equal to a given state everywhere over a given fixed or variable interval.
* the state of the function must be defined and should remain constant (no matter its value) everywhere over a given fixed or variable interval.
* intervals requiring the state of the function to be defined cannot overlap a given fixed or variable interval.
* everywhere over a given fixed or variable interval, the state of the function, if defined, must remain within a given range of states.

In the **C++ API** of CP Optimizer, the functions `IloAlwaysIn`, `IloAlwaysEqual`, `IloAlwaysConstant` and `IloAlwaysNoState` return state constraints.

In the **Java API** of CP Optimizer, the methods `IloCP.alwaysIn`, `IloCP.alwaysEqual`, `IloCP.alwaysConstant` and `IloCP.alwaysNoState` return state constraints.

In the **C# API** of CP Optimizer, the methods `CP.AlwaysIn`, `CP.AlwaysEqual`, `CP.AlwaysConstant` and `CP.AlwaysNoState` return state constraints.

Chapter 5. Constraint propagation in CP Optimizer

CP Optimizer solves a model using constraint propagation and constructive search with search strategies.

Overview

CP Optimizer solves a model using constraint propagation and constructive search with search strategies.

CP Optimizer solves a model using constraint propagation and constructive search with search strategies. There may be multiple algorithms available for reducing domains for decision variables incident on a specific constraint type.

The examples in this section are for illustrating the effect of different propagation levels and have been written using the C++ API.

The term *constraint propagation* refers to two distinct processes, domain reduction and communicating these reductions amongst the constraints. These two processes are described in a general sense first, then the details of domain reduction for each type of constraint follow the general description.

Domain reduction

Constraint propagation reduces the domains of decision variables.

Each decision variable in an IBM ILOG Concert Technology model has a *domain* that is the set of its possible values.

For instance, the domain of the decision variable

```
IloIntVar x(env, -1, 2);
```

is the set of values {-1, 0, 1, 2} , represented as [-1..2].

Note:

In IBM ILOG Concert Technology and CP Optimizer, square brackets denote the *domain* of decision variables. For example, [5 12] denotes a domain as a set consisting of precisely two integers, 5 and 12. In contrast, [5..12] denotes a domain as a range of integers, that is, the interval of integers from 5 to 12, so it consists of 5, 6, 7, 8, 9, 10, 11 and 12.

A constraint is stated over one or more decision variables and restricts the possible assignments of values to these variables. The possible assignments are the solutions of the constraint. For instance, the constraint $x <= y$ allows the assignments $x = 0, y = 1$ or $x = 2, y = 2$ but not $x = 3, y = 2$.

A constraint can perform domain reduction on its decision variables to eliminate from their domains values that do not belong to a solution of the constraint. When the domain reduction algorithm is such that it removes all the values that do not belong to a solution, the process is called full domain reduction. Full domain

reduction sometimes can be very costly in terms of computation time, so, in practice, the domain reduction performed by a constraint does not necessarily eliminate all inconsistent values.

As an example, consider the variables x and y, each of which has the initial domain [1..10] and the constraint $x + y <= 5$.

The domain of each variable is the set of integer values from 1 to 10. A solution to the constraint x + y <= 5 is the assignment to x and y any combination of values from the set {1, 2, 3, 4}. The assignment $x = 5$ does not lead to a solution for this constraint as there is no value in the domain of y that satisfies the constraint when $x = 5$. A full domain reduction for this constraint eliminates all values greater than 4 from the domains of x and y, and thus the domains of x and y are reduced to the interval [1..4].

Constraint propagation

Constraint propagation is the process of communicating the domain reduction of a decision variable to all of the constraints that are stated over this variable.

Constraint propagation is the process of communicating the domain reduction of a decision variable to all of the constraints that are stated over this variable. This process can result in more domain reductions. These domain reductions, in turn, are communicated to the appropriate constraints. This process continues until no more variable domains can be reduced or when a domain becomes empty and a failure occurs. An empty domain during the initial constraint propagation means that the model has no solution.

For example, consider the decision variables y with an initial domain [0..10], z with an initial domain [0..10] and t with an initial domain [0..1], and the constraints

```
y + 5*z <= 4
t != z
t != y
```

over these three variables.

The domain reduction of the constraint $y + 5*z <= 4$ reduces the domain of y to [0..4] and z to [0]. The variable z is thus fixed to a single value. Constraint propagation invokes domain reduction of every constraint involving z. Domain reduction is invoked again for the constraint $y + 5*z <= 4$, but the variable domains cannot be reduced further. Domain reduction of the constraint $t != z$ is invoked again, and because z is fixed to 0, the constraint removes the value 0 from the domain of t. The variable t is now fixed to the value 1, and constraint propagation invokes domain reduction of every constraint involving t, namely $t != z$ and $t != y$. The constraint that can reduce domains further is $t != y$. Domain reduction removes the value 1 from the domain of y.

Constraint propagation is performed on constraints involving y; however, no more domain reduction can be achieved and the final domains are:

- $y = [0\ 2..4]$,
- $z = [0]$ and
- $t = [1]$.

To invoke the constraint propagation process in CP Optimizer, the `propagate` function of the optimizer object is called. In the **C++ API**, this function is

`IloCP::propagate`; in the **Java API**, this function is `IloCP.propagate`; and in the **C#
API**, this function is `CP.Propagate`. This function invokes domain reduction on
every constraint of the model and propagates the domain reductions. It returns
true (`IloTrue` in the C++ API) if propagation succeeds; in other words, if no empty
domains result. It returns false (`IloFalse` in the C++ API) otherwise.

As an example using the C++ API, a code that invokes propagation on the model
above is;

```
IloIntVar y(env, 0, 10);
IloIntVar z(env, 0, 10);
IloIntVar t(env, 0, 1);
IloModel model(env);
model.add(y + 5*z <= 4);
model.add(t != z);
model.add(t != y);
IloCP cp(model);
if (cp.propagate()){
  cp.out() << " Domains reduced: " << std::endl;
  cp.out() << " Domain of y = " << cp.domain(y) << std::endl;
  cp.out() << " Domain of z = " << cp.domain(z) << std::endl;
  cp.out() << " Domain of t = " << cp.domain(t) << std::endl;
}else{
  cp.out() << " Model has no solution." << std::endl;
}
```

The call to the method `IloCP::domain(IloIntVar x)` is directed to an output
stream to display the current domain of the decision variable x. Running this code
produces the output:

```
 Domains reduced:
 Domain of y = [0 2..4]
 Domain of z = [0]
 Domain of t = [1]
```

Propagation of arithmetic constraints

Propagation of arithmetic constraints works to reduce the bounds of the incident
decision variables.

In the C++ API, CP Optimizer can handle arithmetic expressions created with the
operators +, -, * and /. Arithmetic constraints are created with arithmetic
expressions and operators like ==, <=, >=, < and >.

For example, the constraint $17*p*q + x/y <= 100$, where p, q, x and y are decision
variables, can be handled by CP Optimizer. These operators are discussed in more
detail in the section "Arithmetic constraints and expressions" on page 29. The Java
API and C# API equivalents of the operators are listed there as well.

Apart from a few cases that are described below, arithmetic constraints do not
achieve full domain reduction because there does not always exist an efficient
algorithm for full domain reduction. Thus, domain reduction is mostly applied to
bounds of decision variables and is called bound reduction. Bound reduction is
considered as a good trade-off between the number of values removed and the
efficiency of the domain reduction algorithm.

For example, consider the model:

```
IloIntVar x(env, -7, 7);
IloIntVar y(env, -7, 7);
IloModel model(env);
model.add(0.5*x + 3*y == 5);
```

The two solutions of this model are $x = -2$, $y = 2$ and $x = 4$, $y = 1$. Full constraint propagation would give the domain of x as [-2 4] and the domain of y as [1 2]. However, CP Optimizer does not perform this domain reduction. The constraint propagation engine does not create "holes" in the domain of x. The reduced domain of x is [-2..4].

There are exceptions to this behavior. These are binary constraints of the form $y == a*x + b$, where y and x are variables and a and b are numerical values. In this case, full domain reduction is achieved.

For instance consider the constraint $y == 2*x$ over the variables x with domain [1..3] and y with domain [0..10]. This constraint forces y to be even. Full domain reduction is performed and reduces the domain of y to [2 4 6]. The main reason for achieving full domain reduction in this case is that it does not hurt the efficiency of constraint propagation, and it can be effective to propagate holes in domains from a constraint to another when there are linking constraints like $x == y$ in a model.

Another case is that of linear inequalities such as $x + 3y - 4z <= 10$. Achieving bound reduction for these constraints is sufficient to achieve full domain reduction.

CP Optimizer provides bound reduction for expressions such as absolute value, minimum, maximum and piecewise linear functions. On these expressions, achieving bound reduction is sufficient to maintain full domain reduction.

In the C++ API, an absolute value expression is created with the function `IloAbs`. Consider the following code:

```
IloIntVar x(env, -10, 20);
IloIntVar y(env, -3, 4);
IloModel model(env);
model.add(y == IloAbs(x));
IloCP cp(model);
if (cp.propagate()){
  cp.out() << " Domains reduced: " << std::endl;
  cp.out() << " Domain of x = " << cp.domain(x) << std::endl;
  cp.out() << " Domain of y = " << cp.domain(y) << std::endl;
}else{
  cp.out() << " Model has no solution." << std::endl;
}
```

Running this code, the domains of both x and y are reduced. The domain of y is reduced so that it is positive and the domain of x is reduced to take into account the maximum value of y:

```
Domains reduced:
Domain of x = [-4..4]
Domain of y = [0..4]
```

In the C++ API, minimum and maximum values expressions over a set of variables are created with the `IloMin` and `IloMax` expressions. For example consider the model:

```
IloIntVar x(env, 0, 10);
IloIntVar y(env, 4, 6);
IloIntVar u(env, 2, 10);
IloModel model(env);
model.add(u == IloMin(x, y));
IloCP cp(model);
if ( cp.propagate() ) {
  cp.out() << " Domains reduced: " << std::endl;
  cp.out() << " Domain of x = " << cp.domain(x) << std::endl;
  cp.out() << " Domain of y = " << cp.domain(y) << std::endl;
```

```
      cp.out() << " Domain of u = " << cp.domain(u) << std::endl;
    }else{
      cp.out() << " Model has no solution." << std::endl;
    }
```

The value of *u* cannot exceed the smallest upper bound of *x* and *y*, that is 6. Moreover, *x* nor *y* cannot have a value smaller than the lower bound of *u*, which is 2. The domains of the variables after running this code are:

```
Domains reduced:
 Domain of x = [2..10]
 Domain of y = [4..6]
 Domain of u = [2..6]
```

Propagation of logical constraints

Domain reduction of decision variables incident on logical constraints are propagated using the truth values of the subconstraints.

CP Optimizer can process logical constraints. Logical constraints are stated over arithmetic constraints and also over some specialized constraints.

The connectors for creating logical constraints are described in detail in "Logical constraints" on page 32.

- negation (!),
- conjunction (&&),
- disjunction (||),
- implication (IloIfThen),
- equivalence (==) and
- exclusive or (!=).

The semantics of these connectors are the usual logical semantics. The arguments of these connectors are constraints that, in turn, possibly could be logical constraints. For instance, consider two tasks whose durations are represented by the integer values *dx* and *dy* and whose starting dates are represented by the decision variables *x* and *y*. To model the constraint that the two tasks do not overlap, you can state the disjunction:

(x >= y + dy) || (y >= x + dx)

To impose this constraint only when another condition is true (for example when *z == 3*), you can write the logical constraint (shown here using the C++ API):

```
    IloIfThen(env, z == 3, (x >= y + dy) || (y >= x + dx))
```

To understand the domain reduction achieved by these constraints, it is important to observe that a constraint has a truth value. Depending on the domains of the variables, the truth value of a constraint is true when the constraint is definitely satisfied. The truth value is false when the constraint is definitely violated. Otherwise, both values are possible.

Consider, for instance, the constraint *z == 3*. It has a truth value of false when the domain of *z* is [0..2], that is, when it does not contain 3. It has a truth value of true when *z* is fixed to 3. The truth value of this constraint is not yet determined when the domain of *z* contains the value 3 and at least one other value. The truth value is undetermined if the domain of *z* is [2..3]; it is not yet known if the constraint *z == 3* will be satisfied (*z = 3*) or violated (*z = 2*).

In practice, each constraint appearing in a logical expression is associated with a Boolean decision variable whose value is the truth value of the constraint. These Boolean variables are mentioned in the search log as additional variables.

Here is an example of a logical model using the C++ API:

```
    IloIntVar x(env, 0, 5);
    IloIntVar y(env, 7, 20);
    IloIntVar z(env, -10, 20);
    IloModel model(env);
    model.add(x >= y || z < 7);
    model.add(IloIfThen(env, z != 10, y == 10));
    IloCP cp(model);
 if (cp.propagate()) {
    cp.out() << " Domains reduced: " << std::endl;
    cp.out() << " Domain of x = " << cp.domain(x) << std::endl;
    cp.out() << " Domain of y = " << cp.domain(y) << std::endl;
    cp.out() << " Domain of z = " << cp.domain(z) << std::endl;
 }else{
    cp.out() << " Model has no solution." << std::endl;
 }
```

In the first disjunction, the constraint $x >= y$ is obviously violated. To satisfy this disjunction, the constraint $z < 7$ is imposed. This forces the maximum of the domain of z to be 6. The domain reduction of the implication (`IloIfThen`) constraint is invoked since the domain of z has changed. The left member of the implication ($z\ !=\ 10$) is now satisfied and the constraint imposes the constraint $y\ ==\ 10$ which reduces the domain of y. No more domain reduction can be done at this point and this code produces the output:

```
 Domains reduced:
 Domain of y = [10]
 Domain of z = [-10..6]
```

This manner of reducing domains by imposing constraints is very efficient, but it has some limitations. It reacts only to violation or satisfaction of constraints. For instance, the domain reduction of the constraint

$(x == 2)\ ||\ (x == 3)$

will not reduce the domain of x to [2..3]. It waits for one of the constraints to be violated before imposing the other one.

Since any arithmetic or logical constraint can have a truth value, these constraints can appear where any other expression can appear. In particular, such a constraint can be combined in an arithmetic expression. An expression such as

$1 + (x < 7)$

is a valid expression and can be handled by CP Optimizer. Constraints can be created over these expressions. For instance, recall the disjunction

$(x >= y + dy)\ ||\ (y >= x + dx)$

It can be expressed by the constraint

$(x >= y + dy) + (y >= x + dx) >= 1$

Since only one of the disjunction members can be true, it can also be stated as

$(x >= y + dy) + (y >= x + dx) == 1$

Similarly, the implication (here in the C++ API)

```
 IloIfThen(env, z == 3, (x >= y + dy) || (y >= x + dx))
```

can be expressed by

$(z == 3) <= ((x >= y + dy) || (y >= x + dx))$

or by

$(z == 3) <= (x >= y + dy) + (y >= x + dx)$

A typical application of such constraints are cardinality constraints. When among a set of constraints at least two of them must be true, one can state a constraint like

$(x >= y[0]) + (x >= y[1]) + (x >= y[2]) + (x >= y[3]) + (x >= y[4]) >= 2$

The statement of this constraint with only logical connectors is possible but would need a large number of disjunctions.

Propagation of specialized constraints and expressions

CP Optimizer provides tailored constraint propagation algorithms for specialized constraints and expressions.

Overview

CP Optimizer provides tailored constraint propagation algorithms from specialized constraints and expressions.

A fundamental and powerful feature of CP Optimizer are the predefined specialized constraints. In theory, basic constraints like arithmetic and logical ones can model any kind of constraint. A specialized constraint is equivalent to a set of arithmetic or logical constraints. In most of the cases, the specialized constraint achieves more domain reduction than the equivalent set of basic constraints and performs domain reduction more efficiently.

Specialized constraints prove useful in the practical real-world problems of allocation and scheduling. These constraints make it possible to express complicated relations between decision variables, for example, relations that would require a huge number of arithmetic constraints. Specialized constraints enter into such considerations as counting values, maintaining load weights and other such critical activities.

A typical example of the gain specialized constraints can provide is illustrated with the following model implemented using the C++ API:

```
IloIntVar x(env, 1, 2);
IloIntVar y(env, 1, 2);
IloIntVar z(env, 1, 2);
IloModel model(env);
model.add(x != y);
model.add(x != z);
model.add(y != z);
```

Each decision variable has the values 1 and 2 in its domain, and the constraints state that the variables must be all assigned different values. This model clearly has no solution. However, no constraint can perform domain reduction. Examining the constraint *x != y*, both *x = 1, y = 2* and *x = 2, y = 1* are solutions. Therefore, no value can be removed from domains of *x* or *y* if the constraints are considered independently.

To achieve more domain reduction by having a more global view on this model, the three constraints need to be replaced by the all different constraint:

```
model.add(IloAllDiff(env, IloIntVarArray(env, 3, x, y, z)));
```

In order to achieve full domain reduction, the inference level of this constraint, that is the strength of its domain reduction, must be set to the extended level by changing the parameter `IloCP::AllDiffInferenceLevel`:

```
cp.setParameter(IloCP::AllDiffInferenceLevel, IloCP::Extended);
```

Invoking propagation by running the code:

```
IloCP cp(model);
cp.setParameter(IloCP::AllDiffInferenceLevel, IloCP::Extended);
if (cp.propagate()){
  cp.out() << " Domains reduced: " << std::endl;
  cp.out() << " Domain of x is " << cp.domain(x) << std::endl;
  cp.out() << " Domain of y is " << cp.domain(y) << std::endl;
  cp.out() << " Domain of z is " << cp.domain(z) << std::endl;
}else{
  cp.out() << " Model has no solution." << std::endl;
}
```

produces the output:

```
 Model has no solution.
```

because the domain reduction has now created an empty domain for one the variables *x*, *y*, or *z*.

Inference levels

CP Optimizer provides tuning parameters which are used to adjust the constraint propagation inference levels.

The inference level of a constraint, that is, the strength of the domain reduction is achieves, is controlled by tuning parameters. There is a parameter for each specialized constraint whose inference level can be changed.

In the **C++ API**, these parameters are:
- `IloCP::AllDiffInferenceLevel`,
- `IloCP::DistributeInferenceLevel`,
- `IloCP::CountInferenceLevel`,
- `IloCP::CountDifferentInferenceLevel`,
- `IloCP::SequenceInferenceLevel`,
- `IloCP::AllMinDistanceInferenceLevel`,
- `IloCP::ElementInferenceLevel`,
- `IloCP::PrecedenceInferenceLevel`,
- `IloCP::IntervalSequenceInferenceLevel`,
- `IloCP::NoOverlapInferenceLevel`,
- `IloCP::CumulFunctionInferenceLevel` and

- `IloCP::StateFunctionInferenceLevel`.

For the **Java API**, the inference level parameter for the all different constraint is `IloCP.IntParam.AllDiffInferenceLevel`. For the **C# API**, the inference level for the all different constraint is `CP.IntParam.AllDiffInferenceLevel`. The parameters in the **C# API** and the **Java API** for the other specialized constraints are similarly formed.

The effects of changing the values of these inference level parameters will be described in the following sections.

In the **C++ API**, the possible values of these parameters are:
- `IloCP::Default`,
- `IloCP::Low`,
- `IloCP::Basic`,
- `IloCP::Medium` and
- `IloCP::Extended`.

For the **Java API**, the default value is `IloCP.ParameterValues.Default`. For the **C# API**, the default inference level is `CP.ParameterValues.Default`.

The strength of the domain reduction increases as the parameter value moves from the low value to the extended value. Consider the following model written in the **C++ API** where the variables of an array must be all different:

```
IloIntVarArray x(env);
x.add(IloIntVar(env,  1, 2));
x.add(IloIntVar(env,  1, 2));
x.add(IloIntVar(env,  0, 2));
x.add(IloIntVar(env,  IloIntArray(env, 4, 1, 2, 4, 6, 8)));
x.add(IloIntVar(env,  IloIntArray(env, 4, 1, 2, 4, 6, 8)));
x.add(IloIntVar(env,  1, 9));
x.add(IloIntVar(env,  8, 8));
IloModel model(env);
model.add(IloAllDiff(env, x));
```

Setting the inference level of the all different constraints to the basic level and propagating:

```
IloCP cp(model);
cp.setParameter(IloCP::AllDiffInferenceLevel, IloCP::Basic);
if ( cp.propagate() )
  cp.out() << " Domains of x are " << cp.domain(x) << std::endl;
else
  cp.out() << " Model has no solution." << std::endl;
```

produces the output:

```
 Domains of x are [[1..2] [1..2] [0..2] [1..2 4 6] [1..2 4 6] [1..7 9] [8]]
```

The basic level of the all different constraint reduces the domain of variables by eliminating the value of fixed variables (here $x[6]$) from the domain of other variables (here $x[3]$, $x[4]$ and $x[5]$).

When the inference level is set to the medium level, the bounds of the decision variables are reduced further:

```
 Domains of x are [[1..2] [1..2] [0] [4 6] [4 6] [3..7 9] [8]]
```

The values 1 and 2 are shared by the two decision variables x[0] and x[2] and thus cannot appear in the domains of any other variables. This reasoning is applied to reduce bounds of the decision variables. Finally, the extended inference level achieves full domain reduction:

```
Domains of x are [[1..2] [1..2] [0] [4 6] [4 6] [3 5 7 9] [8]]
```

More domain reduction involves more computation time, thus there is a trade-off when solving a problem.

By default, the value of each of inference level is the default level. This level forces the inference level of the constraints to be the value of the parameter (in the **C++ API**) IloCP::DefaultInferenceLevel. Changing the value of this parameter is a way to change the inference level of several constraint types at the same time. The possible values of this parameter are all of the inference levels except for the default level. Its default value is the basic level.

In the **C++ API**, the code:

```
cp.setParameter(IloCP::DefaultInferenceLevel, IloCP::Extended);
```

specifies that the inference level of any constraint type whose inference level is unchanged (or set to the default level) is the extended level.

In the **Java API**, the code for setting the default inference level to the extended level is:

```
cp.setParameter(IloCP.IntParam.DefaultInferenceLevel,
                IloCP.ParameterValues.Extended);
```

In the **C# API**, the code for setting the default inference level to the extended level is:

```
cp.SetParameter(CP.IntParam.DefaultInferenceLevel,
                CP.ParameterValues.Extended);
```

In the following sections, information about the domain reduction achieved by specialized constraints with respect to their inference levels is provided.

The element expression

The constraint propagation algorithm for element expressions reduces domains efficiently.

The element expression indexes an array of values with a decision variable. This expression is used to associate a cost or a distance to the value of a variable, for example.

Suppose there is a decision variable x that chooses a delivery customer and suppose that the distances to the customers are:
- 7 for customer 1,
- 12 for customer 2,
- 5 for customer 3 and
- 21 for customer 4.

If the variable y is to be equal to the distance of the chosen customer, you can write:

```
IloIntArray distance(env, 4, 7, 12, 5, 21);
model.add(y == IloElement(distance, x));
```

This constraint states that y is equal to the x-th element of the distance array. In the **C++ API**, another way to state the same constraint is to write:

```
model.add(y == distance[x]);
```

This constraint achieves full domain reduction. In other words, after domain reduction, for each value of x there is a value for y that satisfies the constraint and vice-versa.

Here is an example of a model with one element expression:

```
IloIntVar x(env, 0, 3);
IloIntVar y(env, 0, 20);
IloModel model(env);
model.add(y == IloElement(IloIntArray(env, 4, 7, 12, 5, 21), x));
IloCP cp(model);
if ( cp.propagate() ) {
  cp.out() << " Domains reduced: " << std::endl;
  cp.out() << " Domain of x is " << cp.domain(x) << std::endl;
  cp.out() << " Domain of y is " << cp.domain(y) << std::endl;
} else
  cp.out() << " Model has no solution." << std::endl;
```

Running this code produces the output:

```
Domains reduced:
Domain of x is [0..2]
Domain of y is [5 7 12]
```

The value 3 is removed from the domain of x because the fourth element in the array is 21, which does not belong to the domain of y. The domain of y is reduced to the set of values that are indexed by values of x, that is 5, 7 and 12.

The counting expression

Inference levels can be used to adjust the propagation of the counting expression.

The specialized counting expression counts the number of times a value appears in the domain of a set of variables. This is useful to count or to constrain the number of times an object or a feature is selected or used.

For instance, assume there are 5 customers, and for each customer, a supplier needs to be chosen. For each customer, there is a list of compatible suppliers. Supplier 1 must not supply more than 2 customers, and the expensive Supplier 2 must not supply more than one customer.

To model the compatibility between suppliers and customers, a decision variable is introduced for each customer that will be fixed to the value of the supplier that supplies it:

```
IloIntVarArray cust(env);
cust.add(IloIntVar(env,  1, 2));
cust.add(IloIntVar(env,  1, 2));
cust.add(IloIntVar(env,  1, 2));
cust.add(IloIntVar(env,  0, 5));
cust.add(IloIntVar(env,  1, 3));
```

The constraints can then be expressed the following way:

```
IloModel model(env);
model.add(IloCount(cust, 1) <= 2);
model.add(IloCount(cust, 2) <= 1);
```

To propagate the constraints of the model add:

```
      IloCP cp(model);
      if (cp.propagate())
        cp.out() << " Domains of cust are " << cp.domain(cust) << std::endl;
      else
        cp.out() << " Model has no solution." << std::endl;
```

Running this program produces the output:

```
 Domains of cust are [[1..2] [1..2] [1..2] [0..5] [1..3]]
```

Using the default inference level of the constraint does not result in any domain
reduction. To get the highest level of domain reduction on this example, the
inference level can be set to the extended level by adding:

```
      cp.setParameter(IloCP::CountInferenceLevel, IloCP::Extended);
```

The output becomes:

```
 Domains of cust are [[1..2] [1..2] [1..2] [0 3..5] [3]]
```

Since values 1 and 2 are the only values in the domain of cust[0], cust[1] and
cust[3], the counting constraints impose that one of these variables will have the
value 2 and two of them will have the value 1. Therefore these values are removed
from the domains of cust[3] and cust[4].

Setting the inference level of the counting expression to the basic or the low level
provides basic but efficient domain reduction. Setting the inference level to the
medium level provides a stronger reduction but on bounds only. Finally, setting
the inference level to the extended level provides full domain reduction.

The distribution constraint

Inference levels can be used to adjust the propagation of the distribution
constraint.

The specialized distribution constraint is an aggregation of counting expressions.

This constraint operates on an array of decision variables varArray, an array of
values valueArray and an array of cardinality variables cardArray.

In the **C++ API**, the constraint is written:

```
  IloDistribute(env, cardArray, valueArray, varArray);
```

It constrains the number of occurrences of a value valueArray[i] in the array
varArray to be equal to cardArray[i].

For instance, reconsidering the example for the counting expression, the model:

```
      IloModel model(env);
      model.add(IloCount(cust, 1) >= 3);
      model.add(IloCount(cust, 2) >= 1);
```

is equivalent to the model:

```
  IloIntVarArray cardArray(env);
  cardArray.add(IloIntVar(env, 3, IloIntMax));
  cardArray.add(IloIntVar(env, 1, IloIntMax));
  IloIntArray valueArray(env, 2, 1, 2);
  IloModel model(env);
  model.add(IloDistribute(env, cardArray, valueArray, cust));
```

Then, setting the inference level to the default level gives the same domain reduction as counting expressions:

```
 Domains of cust are [[1..2] [1..2] [1..2] [0..5] [1..3]]
```

Constraints using the counting expression are aggregated to distribution constraints and thus provide a more global view than otherwise.

The compatibility and incompatibility constraints

Compatibility and incompatibility constraints can be used to model difficult constraints and to improve solving efficiency.

The compatibility constraint is a specialized constraint that is defined by explicitly specifying the set of assignments that are solutions to the constraint.

Compatibility constraints are used frequently in constraint programming applications. There are two broad categories of use:

1. when external data defines a constraint, and such a constraint is difficult to state with arithmetic or logical constraints, and
2. when you want to improve the efficiency of the solving process by modeling a subproblem by assignments.

External data as constraints

In many constraint applications, it is necessary to process a huge quantity of data. For instance, the features of some products can be described as a relation in a database or in text files.

Consider as an example a bicycle factory that can produce thousands of different models. For each model of bicycle, a relation associates the features of that bicycle such as size, weight, color, price. This information can be used in a constraint programming application that allows a customer to find the bicycle that most closely fits a specification.

In the bicycle example, illustrated here using the **C++ API**, an array of decision variables x is defined, where x[0] represents the identifier of the bicycle, x[1] its size, x[2] its weight, x[3] its color and x[4] its price:

```
    IloIntVarArray x(env, 5);
```

A compatibility constraint on x forces the values of x to be one of the combinations defined in the tupleset:

```
    model.add(IloAllowedAssignments(env, x, bicycleSet));
```

where bicycleSet defines the set of solutions to the constraint as an IloIntTupleSet:

```
    IloIntTupleSet bicycleSet(env, 5);
    bicycleSet.add(IloIntArray(env, 5, 1, 57, 12, 3, 1490));
    bicycleSet.add(IloIntArray(env, 5, 2, 57, 13, 5, 1340));
    bicycleSet.add(IloIntArray(env, 5, 3, 60, 14, 3, 1790));
    bicycleSet.add(IloIntArray(env, 5, 4, 65, 14, 7, 1550));
    bicycleSet.add(IloIntArray(env, 5, 5, 67, 15, 2, 2070));
    bicycleSet.add(IloIntArray(env, 5, 6, 70, 15, 2, 1990));
```

Another bicycle variable can be created by construction another array of variables:

```
    IloIntVarArray y(env, 5);
```

The same compatibility constraint can be placed on y:

```
model.add(IloAllowedAssignments(env, y, bicycleSet));
```

It is important to note that tuplesets can be large, and thus they are shared over compatibility constraints that use the same tupleset.

Here is a full example for illustrating domain reduction achieved by this constraint:

```
IloIntVarArray x(env, 5);
x[0] = IloIntVar(env, 0, 10);
x[1] = IloIntVar(env, 40, 60);
x[2] = IloIntVar(env, 10, 20);
x[3] = IloIntVar(env, 5, 6);
x[4] = IloIntVar(env, 1000, 5000);

IloIntTupleSet bicycleSet(env, 5);
bicycleSet.add(IloIntArray(env, 5, 1, 57, 12, 5, 1490)); // tuple 0
bicycleSet.add(IloIntArray(env, 5, 2, 57, 13, 1, 1340)); // tuple 1
bicycleSet.add(IloIntArray(env, 5, 3, 60, 14, 5, 1790)); // tuple 2
bicycleSet.add(IloIntArray(env, 5, 4, 65, 14, 4, 1550)); // tuple 3
bicycleSet.add(IloIntArray(env, 5, 5, 67, 15, 2, 2070)); // tuple 4
bicycleSet.add(IloIntArray(env, 5, 6, 70, 15, 5, 1990)); // tuple 5

IloModel model(env);
model.add(IloAllowedAssignments(env, x, bicycleSet));

IloCP cp(model);
if (cp.propagate())
  cp.out() << " Domains of vars = " << cp.domain(x) << std::endl;
else
  cp.out() << " Model has no solution." << std::endl;
```

Running this code produces the output:

```
 Domains of vars = [[1 3] [57 60] [12 14] [5] [1490 1790]]
```

There are only two possible solutions to the constraint due to the domains of the variables in the array x. With the domain of x[3], only tuples 0, 2 and 5 are possible. The domain of x[1] eliminates the tuple 5. Thus the final domains are the union of the possible values in the two remaining solutions.

The forbidden assignments constraint plays a symmetrical role: the tupleset represents the set of non-solutions of the constraint. This is useful when there are fewer non-solutions than solutions.

Both constraints achieve full domain reduction and do not support inference levels.

Improving efficiency: a compatibility constraint for a subproblem

A modeling trick that may dramatically reduce the computation time needed to solve a problem consists in identifying a difficult subproblem, computing all the solutions of the subproblem and storing them in a tuple set and then creating a compatibility constraint.

This approach is not restricted to constraint programming but is a general approach: facing a difficult problem, it can be easier to solve it by:
• decomposing the problem into subproblems,
• solving the different subproblems and

- connecting the solutions of the subproblems to produce a solution to the whole problem.

An allowed assignments constraint forces the values of the variables of the problem to be one of the solutions of the subproblem. Thus the connection of the solution of the subproblem with the remainder of the problem is automatically handled. The advantage of this approach is that when searching for a solution of the whole problem, instead of always retrieving the solutions of the subproblem in the search, the allowed assignments constraint forces values to one of these solutions. In other terms, the work of solving the subproblem is factored and done only once, before the search, and not several times during the search (which is potentially a huge number of times).

This approach also has a drawback: the solutions of the subproblem must be found first so this must be practical (the solutions should not be too numerous). Allowed assignment constraints with several hundreds of thousands of tuples can be handled, but subproblems with billions of solutions cannot be handled in this way. Nevertheless, it is possible to set a bound on the number of solutions of the subproblem to precompute and store in the tupleset. In this case, the problem solved is a restriction of the initial problem, but this may be useful in practice.

An example of such an approach is available in the file `teambuilding.cpp` in the `examples/src/cpp` directory.

Constraint aggregation

CP Optimizer may aggregate constraints to improve the efficiency of propagation.

The CP Optimizer engine may sometimes preprocess the model in order to improve the efficiency of the propagation engine.

The optimizer may aggregate constraints by reducing some groups of basic constraints into less basic ones. By doing this, the propagation algorithm may be able to achieve more domain reduction. For example, if the model has a set of inequality constraints, the engine may work to combine some of these into an all different constraint.

The constraint aggregator is on by default. It can be turned off by setting the value of the aggregator parameter to "Off".

In the **C++ API**, the parameter is `IloCP::ConstraintAggregation` and the value `IloCP::Off`.

For the **Java API**, the aggregator parameter is `IloCP.IntParam.ConstraintAggregation` and the value is `IloCP.ParameterValues.Off`.

For the **C# API**, the aggregator parameter is `CP.IntParam.ConstraintAggregation` and the value is `CP.ParameterValues.Off`.

Chapter 6. Search in CP Optimizer

CP Optimizer uses constructive search strategies to find a solution to a constraint programming problem.

Overview

CP Optimizer uses constructive search strategies to find a solution to a constraint programming problem.

A methodical approach to developing an application using CP Optimizer begins with a verbal description of the problem and moves directly to the model or representation. Once the model is created, you are ready to solve it using the optimizer. For most problems, using the basic solve function is all that is needed for solving the model. Nonetheless, CP Optimizer offers a variety of controls that allow you to tailor the solution process for your specific needs.

The CP Optimizer engine searches for an optimal solution or, if the problem is a satisfiability one, a solution of a model. A solution is an assignment of values to variables such that every constraint in the model is satisfied.

A naive way to find a solution would be to explicitly study each possible combination of values for decision variables until a solution is found. The CP Optimizer search implicitly generates such combinations but in a very efficient manner using constraint propagation. It produces an optimal solution for optimization problems and at least one solution for satisfiability problems.

The implicit generation of combinations uses *constructive search strategies*. A constructive strategy attempts to build a solution by choosing a non-fixed decision variable and a value for that variable. The chosen variable is then fixed to the chosen value, and constraint propagation is triggered. This operation is called *branching*, and the fixing is also called a "branch". Constraint propagation reduces the domains of variables and, consequently, the currently possible combinations. After propagation terminates, another non-fixed variable, if one exists, is chosen, and the process repeats until all decision variables are fixed. However, if a fixing fails because it cannot lead to a solution, the constructive strategy *backtracks* and chooses another value for the variable.

The CP Optimizer constructive search strategies are guided towards optimal solutions in order to converge rapidly. The CP Optimizer search uses a variety of guides and uses the most appropriate one depending on the model structure and on constraint propagation. This section explains the standard search algorithms used.

Searching for solutions

CP Optimizer uses constructive search strategies to find a solution to a constraint programming problem.

Overview

CP Optimizer uses constructive search strategies to find a solution to a constraint programming problem. While the built-in algorithm will generally be sufficient, it may be helpful for you to know about the other algorithms in the case that you need to tune the optimizer.

CP Optimizer provides a number of search strategies for you to use. While the built-in algorithm will generally be sufficient, it may be helpful for you to know about the other algorithms in the case that you need to tune the optimizer.

In the **C++ API** of CP Optimizer, you use the class `IloCP`, which is a subclass of `IloAlgorithm`, to execute and control the search. The constructor for `IloCP` takes an `IloModel` as its argument. As with the environment, once you are finished with the optimizer, you call the method `IloCP::end` to reclaim the memory used by the optimizer.

In the **Java API** of CP Optimizer, you use the class `IloCP` to execute and control the search.

Likewise, in the **C# API** of CP Optimizer, you use the class `CP` to execute and control the search to find a solution to a problem expressed in a model.

Note:

CP optimizer

The class `IloCP` in the C++ API and the Java API and the class `CP` in the C# API can be used to employ different algorithms for solving problems modeled with Concert Technology modeling classes.

An object of this class is sometimes referred to as the optimizer.

Solving an optimization problem

The basic search strategy in CP Optimizer can be used to solve optimization problems.

The basic algorithm for solving a model is invoked by calling the method `solve`, a member function of the optimizer object. (For example, in the **C++ API**, this is `IloCP::solve`. In the **Java API**, this method is `IloCP.solve`. In the **C# API**, this method is `CP.Solve`.) This function returns true (`IloTrue` in the C++ API) when the engine has found an optimal solution. If the model has no objective function, then this call returns true when a solution is found. The function returns false (`IloFalse` in the C++ API) when the problem has no solution.

Consider the following model with 3 integer variables, written in the C++ API:

```
IloIntVar x(env, 0, 7, "x");
IloIntVar y(env, 0, 7, "y");
IloIntVar z(env, 0, 7, "z");
IloIntVarArray all(env, 3, x, y, z);
model.add(IloMinimize(env, IloSum(all)));
model.add(IloAllDiff(env, all));
model.add(y == IloElement(IloIntArray(env, 7, 3, 7, 8, 8, 0, 1, 4), x));
```

To solve the problem represented by this model, you can call the method `IloCP::solve`, like this:

```
IloCP cp(model);
cp.setParameter(IloCP::LogVerbosity, IloCP::Quiet);
if (cp.solve()){
  cp.out() << " An optimal solution has been found"
          << ", objective value = " << cp.getObjValue()
          << ", x = " << cp.getValue(x)
          << ", y = " << cp.getValue(y)
          << ", z = " << cp.getValue(z) << std::endl;
} else {
  cp.out() << " The problem has no solution " << std::endl;
}
```

In this example, the search log output is deactivated for the sake of brevity. Information regarding the search log is presented in the section "The search log" on page 65. The method getValue, a member function of the optimizer object, takes a decision variable as an argument and returns the value of that variable in the solution that was found. The method getObjValue, a member function of the optimizer object, returns the value of the objective for this solution.

Running this code produces the output:

```
An optimal solution has been found, objective value = 4, x = 0, y = 3, z = 1
```

In the **C++ API** of CP Optimizer, you use the class IloCP and the methods IloCP::solve, IloCP::getValue and IloCP::getObjValue.

In the **Java API** of CP Optimizer, you use the class IloCP and the methods IloCP.solve, IloCP.getValue and IloCP.getObjValue.

In the **C# API** of CP Optimizer, you use the class CP and the methods CP.Solve and CP.GetValue and the member CP.ObjValue.

Accessing intermediate solutions

As CP Optimizer searched for an optimal solution to an optimization problem, the search will generally encounter a sequence of solutions that improve the objective function.

As the CP Optimizer engine solves an optimization problem, a sequence of solutions that improve the objective function are produced until an optimal solution is found. In some cases, you may want to have access to this sequence of intermediate solutions. CP Optimizer provides a simple interface that provides this access via member functions of the optimizer object.

To prepare the optimizer for search, you call the member function startNewSearch, a member function of the optimizer object. To instruct the optimizer to find a solution, you call the method next. This method returns true (IloTrue in the C++ API) if the optimizer finds a solution (not necessarily an optimal one) and false (IloFalse in the C++ API) if the problem is infeasible. To find the next solution, you again call the method next. If the method returns true, then the optimizer has found a new solution with a strictly better objective value than the previous one. If the optimizer does not find another solution, the value returned by this method is false. When you have finished searching, you can call the method end to free the memory and reset the state of the optimizer. A typical code using these methods of the optimizer object in the C++ API would be similar to:

```
IloCP cp(model);
cp.setParameter(IloCP::LogVerbosity, IloCP::Quiet);
cp.startNewSearch();
while(cp.next()){
  cp.out() << "objective value = " << cp.getObjValue()
```

```
            << ", x = " << cp.getValue(x)
            << ", y = " << cp.getValue(y)
            << ", z = " << cp.getValue(z) << std::endl;
    }
    cp.end();
```

Running this code produces the output:

```
objective value = 5, x = 4, y = 0, z = 1
objective value = 4, x = 0, y = 3, z = 1
```

This run finds two solutions, one with an objective value of 5 and then one with an objective value of 4, which is the optimum.

In the **C++ API** of CP Optimizer, you use the class `IloCP` and the methods `IloCP::startNewSearch` and `IloCP::next`.

In the **Java API** of CP Optimizer, you use the class `IloCP` and the methods `IloCP.startNewSearch` and `IloCP.next`.

In the **C# API** of CP Optimizer, you use the class `CP` and the methods `CP.StartNewSearch` and `CP.Next`.

Solving a satisfiability problem

The basic search strategy in CP Optimizer can be used to solve satisfiability problems.

The process for finding solutions to constraint satisfaction problems is similar to the process for finding solutions to optimization problems.

Consider this satisfiability problem, written in the C++ API:

```
    IloIntVar x(env, 0, 7, "x");
    IloIntVar y(env, 0, 7, "y");
    IloIntVar z(env, 0, 7, "z");
    IloIntVarArray all(env, 3, x, y, z);
    model.add(IloAllDiff(env, all));
    model.add(y == IloElement(IloIntArray(env, 7, 3, 7, 8, 8, 0, 1, 4), x));
```

A typical code for finding a solution to this model would look like:

```
    IloCP cp(model);
    cp.setParameter(IloCP::LogVerbosity, IloCP::Quiet);
    if (cp.solve()){
      cp.out() << " A solution has been found"
              << ", x = " << cp.getValue(x)
              << ", y = " << cp.getValue(y)
              << ", z = " << cp.getValue(z) << std::endl;
    } else {
      cp.out() << " The problem has no solution " << std::endl;
    }
```

Running this code produces the output:

```
A solution has been found, x = 4, y = 0, z = 1
```

For constraint satisfaction problems, the search stops at the first solution encountered, thus there are no intermediate solutions found. However, it is possible to produce all solutions to a constraint satisfaction problem. For the model at the start of this section, running the code:

```
            IloCP cp(model);
            cp.setParameter(IloCP::LogVerbosity, IloCP::Quiet);
            cp.startNewSearch();
            while(cp.next()){
              cp.out() << "x = " << cp.getValue(x)
                       << ", y = " << cp.getValue(y)
                       << ", z = " << cp.getValue(z) << std::endl;
            }
            cp.end();
```

produces the 30 distinct solutions of the model:

```
x = 4, y = 0, z = 1
x = 4, y = 0, z = 2
x = 4, y = 0, z = 3
x = 4, y = 0, z = 5
x = 4, y = 0, z = 6
x = 4, y = 0, z = 7
x = 0, y = 3, z = 1
x = 0, y = 3, z = 2
x = 0, y = 3, z = 4
x = 0, y = 3, z = 5
x = 0, y = 3, z = 6
x = 0, y = 3, z = 7
x = 5, y = 1, z = 0
x = 5, y = 1, z = 2
x = 5, y = 1, z = 3
x = 5, y = 1, z = 4
x = 5, y = 1, z = 6
x = 5, y = 1, z = 7
x = 1, y = 7, z = 0
x = 1, y = 7, z = 2
x = 1, y = 7, z = 3
x = 1, y = 7, z = 4
x = 1, y = 7, z = 5
x = 1, y = 7, z = 6
x = 6, y = 4, z = 0
x = 6, y = 4, z = 1
x = 6, y = 4, z = 2
x = 6, y = 4, z = 3
x = 6, y = 4, z = 5
x = 6, y = 4, z = 7
```

In the **C++ API** of CP Optimizer, you use the class IloCP and the methods IloCP::getValue, IloCP::getObjValue, IloCP::startNewSearch, IloCP::next and IloCP::end.

In the **Java API** of CP Optimizer, you use the class IloCP and the methods IloCP.solve, IloCP.getValue, IloCP.getObjValue, IloCP.startNewSearch, IloCP.next and IloCP.end.

In the **C# API** of CP Optimizer, you use the class CP and the methods CP.Solve, CP.GetValue, CP.GetObjValue, CP.StartNewSearch, CP.Next and CP.End.

The search log

During search, information regarding the progress of the optimizer is displayed to the output channel of the optimizer; this information is called the *search log* or simply the *log*.

Reading a search log

The search log provides detailed information of the model and the search.

A sample of the search log looks like:

```
! --------------------------------------------------------------------------
! Minimization problem - 13 variables, 13 constraints
! Presolve      : 5 extractables eliminated, 1 constraint generated
! Workers            = 2
! LogPeriod          = 100
! Initial process time : 0.00s (0.00s extraction + 0.00s propagation)
!  . Log search space  : 23.6 (before), 23.6 (after)
!  . Memory usage      : 290.4 kB (before), 290.4 kB (after)
! Using parallel search with 2 workers.
! --------------------------------------------------------------------------
!          Best Branches  Non-fixed   W       Branch decision
                      0         13                    -
+ New bound is 256
                      0         13    1   F          -
+ New bound is 633
                      0         13    1   F          -
+ New bound is 634
*           1704        15  0.02s     1       (gap is 62.79%)
*           1644        53  0.02s     1       (gap is 61.44%)
*           1525        88  0.02s     1       (gap is 58.43%)
            1525       100         1   1          2  = S_6
            1525       133         1   1   F      1  = O_4
+ New bound is 1080 (gap is 29.18%)
            1525       141         1   1   F      1  = O_0
+ New bound is 1304 (gap is 14.49%)
*           1414       159  0.02s     1       (gap is 7.78%)
            1414       169         1   1   F      1  = O_1
+ New bound is 1360 (gap is 3.82%)
*           1383       191  0.02s     1       (gap is 1.66%)
            1383       200         1   1   F      0  = O_1
            1383       201        10   1   F      1  = O_1
+ New bound is 1372 (gap is 0.8%)
            1383       215         5   1   F      1  = O_4
+ New bound is 1379 (gap is 0.29%)
            1383       231         5   1   F      1  = O_1
+ New bound is 1382 (gap is 0.072%)
            1383       300         4   1   F      2  = S_4
            1383       400         5   1   F      0 != S_5
            1383       500         5   1   F      0  = S_6
            1383       600         5   1   F      3  = S_7
! Time = 0.02s, Average fail depth = 7, Memory usage = 778.6 kB
! Current bound is 1382 (gap is 0.072%)
!          Best Branches  Non-fixed   W       Branch decision
            1383       700         5   1   F      1  = S_1
            1383       790         4   1   F      1  = O_1
+ New bound is 1383 (gap is 0%)
! --------------------------------------------------------------------------
! Search completed, 5 solutions found.
! Best objective        : 1383 (optimal - effective tol. is 0)
! Best bound             : 1383
! Number of branches     : 790
! Number of fails        : 379
! Total memory usage     : 812.0 kB (759.8 kB CP Optimizer + 52.2 kB Concert)
! Time spent in solve    : 0.02s (0.02s engine + 0.00s extraction)
! Search speed (br. / s) : 37619.0
! --------------------------------------------------------------------------
```

The first line of the log indicates the type of problem, along with the number of decision variables and constraints in the model. In this case, there is an objective included in the model, so the problem is reported to be a "Minimization problem".

When the model does not include an objective, the problem type is reported as a "Satisfiability problem". The number of search phases, if any, is also displayed in the first line.

The second line of the log shows the result of presolve. The number of model objects (extractables) eliminated is displayed along with the number of constraints generated to improve the formulation of the eliminated extractables if any.

Any parameter change from its default is displayed after the preprocessing information.

The next three lines of the log provide information regarding the initial constraint propagation. The "Initial process time" is the time in seconds spent at the root node of the search tree where the initial propagation occurs. This time encompasses the time used by the optimizer to load the model, called extraction, and the time spent in initial propagation. The value for "Log search space" provides an estimate on the size of the depth-first search tree; this value is the log (base 2) of the products of the domains sizes of all the decision variables of the problem. Typically, the estimate of the size of the search tree should be smaller after the initial propagation, as choices will have been eliminated. However, this value is always an overestimate of the log of the number of remaining leaf nodes of the tree because it does not take into account the action of propagation of constraints at each node. The memory used by the optimizer during the initial propagation is reported.

The log then display the type of search used, sequential or parallel; in the latter case, the number of workers used is also displayed.

In order to interpret the remainder of the log file, you may want to think about the search as a binary tree. The root of the tree is the starting point in the search for a solution; each branch descending from the root represents an alternative choice or decision in the search. Each of these branches leads to a node where constraint propagation during search will occur. If the branch does not lead to a failure and a solution is not found at a node, the node is called a choice point. The optimizer can make an additional decision and create two new alternative branches from the current node, or it can jump in the tree and search from another node.

The lines in the next section of the progress log are displayed periodically during the search and describe the state of the search. The display frequency of the progress log can be controlled with parameters of the optimizer.

The progress information given in a progress log update includes:
- Best: the value of the best solution found so far, in the case of an optimization problem;
- Branches: the number of branches explored in the binary search tree;
- Non-fixed: the number of uninstantiated (not fixed) model variables, or the elapsed time;
- W: the id of the worker at the branch currently under consideration by the optimizer;
- Branch decision: the decision made at the branch currently under consideration by the optimizer.

For optimization problems, whenever a solution is found, the log displays a specific line starting with "*". In this line, the time is displayed in place of the non-fixed value (which is always zero in this case), and the relative gap is

displayed when it is meaningful in place of the branch decision. For minimization (resp. maximization) problems, the best known lower (resp. upper) bound is displayed each time it is improved, in a line starting with "+".

The final lines of the log provide information about the entire search, after the search has terminated. This information about the search includes:

- Termination status line: the conditions under which the search terminated and the number of solutions found during search;
- Best objective: the value of the best solution found during search along with the effective optimality tolerance;
- Best bound: in case of optimality, this is equal to the best objective;
- Number of branches: the number of branches explored in the binary search tree;
- Number of fails: the number of branches that did not lead to a solution;
- Total memory usage: the memory used by IBM ILOG Concert Technology and the CP Optimizer engine;
- Time spent in solve: the elapsed time from start to the end of the search displayed in hh:mm:ss.ff format;
- Search speed: average time spent per branch.

If the search terminated due to a limit, this final part of the log could be quite different and the objective not optimal, as in the following sample:

```
! ----------------------------------------------------------------------------
! Minimization problem - 5100 variables, 11134 constraints, 1 phase
! Presolve      : 378 extractables eliminated
! TimeLimit          = 20
! LogPeriod          = 10000
! Initial process time : 0.63s (0.62s extraction + 0.00s propagation)
!  . Log search space  : 11936.1 (before), 11936.1 (after)
!  . Memory usage      : 7.3 MB (before), 7.3 MB (after)
! Using parallel search with 8 workers.
! ----------------------------------------------------------------------------
!        Best Branches  Non-fixed    W      Branch decision
                   0       2679                   -
+ New bound is 28
                   0       2679    1   F          -
+ New bound is 32
                   0       2679    1   F          -
+ New bound is 33
*            183    5481  6.30s       4       (gap is 81.97%)
             183   10000      2549    5   F    0  = _int4922
             183   10000      2520    6   F    0  = _int4920
             183   10000      2446    7        1  = _int3209
             183   10000      2528    8   F    0  = _int4922
             183   10000      2446    1        0  = _int3844
             183   10000      2469    2        0  = _int4941
             183   10000      2540    3   F    0  = _int4921
             183   10000         1    4        0  = _int3698
*            182   17106 13.81s       7       (gap is 81.87%)
             182   20000      2534    8   F    1  = _int4920
             182   20000      2553    5        0  = _int3447
             182   20000         1    7        2  = _int2240
*            181   20032 15.69s       7       (gap is 81.77%)
             181   20000      2446    1        1  = _int4943
             181   20000      2446    2        1  = _int4935
             181   20000      2515    6   F    1  = _int3466
! Time = 16.23s, Average fail depth = 61, Memory usage = 118.8 MB
! Current bound is 33 (gap is 81.77%)
!        Best Branches  Non-fixed    W      Branch decision
             181   20000      2529    3        0  = _int3255
             181   20000         1    4   F    0  = _int4515
```

```
*             180     22955  17.34s          7       (gap is 81.67%)
! --------------------------------------------------------------------------------
! Search terminated by limit, 4 solutions found.
! Best objective          : 180 (gap is 81.67%)
! Best bound               : 33
! Number of branches       : 167651
! Number of fails          : 67872
! Total memory usage        : 129.1 MB (122.4 MB CP Optimizer + 6.7 MB Concert)
! Time spent in solve       : 20.23s (19.60s engine + 0.62s extraction)
! Search speed (br. / s) : 8551.4
! --------------------------------------------------------------------------------
```

In this case, the objective and the best bound don't fit and the gap is displayed.

For multicriteria optimization, the log displays the objective of the best solution found, the best known bound and the gap as in the following example.

```
! --------------------------------------------------------------------------------
! Minimization problem - 102 variables, 114 constraints
! TimeLimit            = 4
! Workers              = 2
! LogPeriod            = 25000
! Initial process time : 0.00s (0.00s extraction + 0.00s propagation)
!  . Log search space  : 280.6 (before), 280.6 (after)
!  . Memory usage       : 585.4 kB (before), 585.4 kB (after)
! Using parallel search with 2 workers.
! --------------------------------------------------------------------------------
!        Best Branches  Non-fixed    W        Branch decision
                   0        102                    -
+ New bound is 2; 9
                   0        102    1   F           -
+ New bound is 3; 15
                   0        102    1   F           -
+ New bound is 3; 16
*         86       421  0.03s         2       (gap is 96.51% @ crit. 1 of 2)
  New objective is 86; 12
*         42      1430  0.04s         1       (gap is 92.86% @ crit. 1 of 2)
  New objective is 42; 15
*         30      1226  0.05s         2       (gap is 90.00% @ crit. 1 of 2)
  New objective is 30; 15
*         28      5646  0.10s         2       (gap is 89.29% @ crit. 1 of 2)
  New objective is 28; 14
*         26     11964  0.20s         1       (gap is 88.46% @ crit. 1 of 2)
  New objective is 26; 13
          26     25000         1    2        10  = _int55
          26     25000         1    1        14 != _int51
          26     50000         1    1         4  = _int53
          26     50000         1    2        10  = _int59
          26     75000         1    2        12 != _int58
          26     75000         1    1   F     4 != _int54
          26      100k         1    2         0  = _int18
          26      100k         1    1   F     4  = _int48
          26      125k         1    2   F    12  = _int46
          26      125k        10    1   F     0  = _int88
          26      150k        13    2   F    12  = _int60
          26      150k         1    1         7 != _int51
! Time = 2.42s, Average fail depth = 19, Memory usage = 3.1 MB
! Current objective is 26; 13
! Current bound is 3; 16 (gap is 88.46% @ crit. 1 of 2)
!        Best Branches  Non-fixed    W        Branch decision
          26      175k         1    2         9  = _int60
          26      175k         1    1         5  = _int76
          26      200k         1    2         2  = _int25
          26      200k         1    1   F    10  = _int12
          26      225k         1    2         0 != _int5
          26      225k         1    1        13 != _int60
          26      250k        14    2         0  = _int18
```

```
! --------------------------------------------------------------------------
! Search terminated by limit, 5 solutions found.
! Best objective        : 26; 13 (gap is 88.46% @ crit. 1 of 2)
! Best bound            : 3; 16
! Number of branches    : 486763
! Number of fails       : 229441
! Total memory usage    : 3.4 MB (3.2 MB CP Optimizer + 0.1 MB Concert)
! Time spent in solve   : 4.00s (4.00s engine + 0.00s extraction)
! Search speed (br. / s) : 121599.6
! --------------------------------------------------------------------------
```

In the progress update, you can see that the column "Best" gives only the value of the first criteria of the best solution found so far. The values of all the criteria of the best solution found are displayed when a new solution is found. Likewise, The values of all the criteria of the best bound found are displayed when a new bound is found.

Finally, notice that the gap is computed on the first (i.e., most important in lexicographic optimization) non optimal criterion. The log displays the gap and the rank of of the criterion on which it is computed. For instance, if the first two criteria are within the optimality tolerance, then the gap on the third criterion is displayed as follows:

```
( gap is 25.50% @ crit. 3 of 4)
```

Note:

The CP Optimizer search log is meant for visual inspection only, not for mechanized parsing. In particular, the log may change from version to version of CP Optimizer in order to improve the quality of information displayed in the log. Any code based on the log output may have to be updated when a new version of CP Optimizer is released.

Search log parameters

Search log parameters control what and how much information is displayed.

The amount of information displayed by the log can be controlled with the log verbosity parameter. The display frequency of the progress information is controlled with the log period parameter.

In the **C++ API** of CP Optimizer, you set a parameter with a call to IloCP::setParameter. Log verbosity is controlled with the parameter IloCP::LogVerbosity. A value of IloCP::Quiet will suppress the log altogether. Suppressing the log is done with a line of code (here in the C++ API):

```
cp.setParameter(IloCP::LogVerbosity,IloCP::Quiet);
```

The display frequency of the progress information is controlled with the parameter IloCP::LogPeriod. By setting this parameter to a value of k, the log is displayed after every k search decisions.

In the **Java API** of CP Optimizer, you set a parameter with a call to IloCP.setParameter. Log verbosity is controlled with the parameter IloCP.IntParam.LogVerbosity. A value of IloCP.ParameterValues.Quiet will suppress the log altogether. The display frequency of the progress information is controlled with the parameter IloCP.IntParam.LogPeriod. By setting this parameter to a value of k, the log is displayed after every k search decisions.

Likewise, in the **C# API** of CP Optimizer, you set a parameter with a call to `CP.SetParameter`. Log verbosity is controlled with the parameter `CP.IntParam.LogVerbosity`. A value of `CP.ParameterValues.Quiet` will suppress the log altogether. The display frequency of the progress information is controlled with the parameter `CP.IntParam.LogPeriod`. By setting this parameter to a value of k, the log is displayed after every k search decisions.

Retrieving a solution

CP Optimizer provides objects to represent the solution of a problem. These objects allow for easy retrievability of the solution.

When the CP Optimizer engine has found a solution, you can examine the values that have been assigned to the model variables by using the method `getValue`, a member function of the optimizer object, with the argument being the model variable.

In some cases, you may want to save all or part of a solution for use later. IBM ILOG Concert Technology provides a solution class that is useful for transferring stored values from or to the active model objects associated with a solution.

In the **C++ API** of CP Optimizer, you use the class `IloSolution`, which is created on the environment.

In the **Java API** of CP Optimizer, you use the interface `IloSolution` to store and transfer values.

Likewise, in the **C# API** of CP Optimizer, you use the interface `ISolution` to store and transfer values.

You must use the methods `add`, member functions of the solution class to inform the solution that it should store the decision variable or those in the array of decision variables that is passed as an argument to the method. When the optimizer has found a solution that you want to store, you use the method `store` to store the solution.

Note:

Solution class

The class `IloSolution` in the C++ API, the interface `IloSolution` in the Java API and the interface `ISolution` in the C# API makes it possible to store the values from decision variables and also to fix those variables with stored values.

To illustrate using the C++ API, the optimal solution found for the optimization model from the section "Solving an optimization problem" on page 62 can be stored like this:

```
IloCP cp(model);
IloSolution solution(env);
solution.add(all);
if (cp.solve()) {
  solution.store(cp);
  cp.out() << "An optimal solution is " << solution << std::endl;
}
```

Running this code produces the output:
```
An optimal solution is IloSolution[ x[0] y[3] z[1] ]
```

In the **C++ API** of CP Optimizer, you use the class `IloSolution` and the methods `IloSolution::add` and `IloSolution::store`.

In the **Java API** of CP Optimizer, you use the interface `IloSolution` and the methods `IloSolution.add` and `IloSolution.store`.

In the **C# API** of CP Optimizer, you use the interface `ISolution` and the methods `ISolution.Add` and `ISolution.Store`.

Retrieving search information

Information about the search can be retrieved from the search log as well as by using functions provided by CP Optimizer.

There are several ways to obtain information regarding the search. One way is by examining the search log as discussed in the section "The search log" on page 65. Information about the search can also be obtained using the method `printInformation`, a member function of the optimizer object.

For the optimization example in the section "Solving an optimization problem" on page 62, adding a call to the `printInformation` method produces the output:

```
Number of variables    : 5
Number of constraints  : 5
Engine memory usage     : 707.9 kB
Time in last solve      : 0.01s (0.01s engine + 0.00s extraction)
Total time spent in CP  : 0.01s
```

This output contains the number of variables (additional variables are those added to the optimizer to help with the search) and the number of constraints in the problem. These values can be greater than the number of variables and constraints in the model because additional variables and constraints may have been added internally when the optimizer loaded the model. Finally, the memory consumption and some timings are reported.

A third way of retrieving information about the search is to use the method `getInfo`, a member function of the optimizer object. You can have access to search information such as the solving time, the number of branches, the number of failures and the memory usage.

In the **C++ API** of CP Optimizer, you use the method `IloCP::getInfo`. This function takes one argument, either of type `IloCP::NumInfo` or `IloCP::IntInfo`. For example, to have access to the solving time, use the argument `IloCP::SolveTime`.

In the **Java API** of CP Optimizer, you use the method `IloCP.getInfo`. This function takes one argument, either of type `IloCP.DoubleInfo` or `IloCP.IntInfo`. For example, to have access to the solving time, use the argument `IloCP.DoubleInfo.SolveTime`.

Likewise, in the **C# API** of CP Optimizer, you use the method `CP.GetInfo`. This function takes one argument, either of type `CP.DoubleInfo` or `CP.IntInfo`. For example, to have access to the solving time, use the argument `CP.DoubleInfo.SolveTime`.

For an exhaustive list of information that can be accessed, see the method `IloCP::getInfo` in the C++ API reference manual, `IloCP.getInfo` in the Java API reference manual, or `CP.GetInfo` in the .NET languages API reference manual.

Setting parameters on search

Parameters can be set on the search to limit the search as well as adjust the tolerance on optimality.

Search parameterization is an important feature of CP Optimizer. One use of parameters is to limit the search. The parameter `TimeLimit` sets a time limit on the time spent in search. The parameter `BranchLimit` limits the total number of branches (decisions) that are performed by the optimizer.

When a limit is set on the search process of an optimization problem, the call to the optimizer object member function `solve` terminates when the limit is reached. The function returns true when a solution is available and false otherwise. Note that the number of branches and the time limit can go slightly beyond the specified limit because the best solution found gets "replayed" (regenerated), and this can produce some extra time or branches.

In general, to obtain information on the reason the search ended, you can query the engine using the `getInfo` member function of the optimizer object (`IloCP::getInfo` in the **C++ API**, `IloCP.getInfo` in the **Java API** and `CP.GetInfo` in the **C# API**). with the argument `SearchStatus` (`IloCP::SearchStatus` in the **C++ API**, `IloCP.IntInfo.SearchStatus` in the **Java API**, and `CP.IntInfo.SearchStatus`in the **C# API**) or `SearchStopCause` (`IloCP::SearchStopCause` in the **C++ API**, `IloCP.IntInfo.SearchStopCause` in the **Java API**, and `CP.IntInfo.SearchStopCause` in the **C# API**). The meanings of the return values of this function are listed in the CP Optimizer reference manuals.

Another important search parameterization is the one that defines optimality. A solution is considered optimal if there does not exist a solution with a better objective function with respect to an *optimality tolerance*. This tolerance can be absolute and is controlled with the search parameter `OptimalityTolerance`. The relative optimality tolerance is controlled with the search parameter `RelativeOptimalityTolerance`.

For instance, if you consider that an improvement of 10 on your objective function is negligible, you can set this tolerance using the C++ API with the call:

```
cp.setParameter(IloCP::OptimalityTolerance, 10);
```

With this tolerance set, if an optimal solution of a minimization problem is found with an objective value of 900, then there does not exists a solution with an objective value of 890. There may exist solutions with objective values of 891 and higher, but these have been missed due to the tolerance. The default value for this tolerance is 1e-9.

Another example: if you want to find a solution within 3% of the optimal, you set the relative optimality tolerance using the C++ API with the call:

```
cp.setParameter(IloCP::RelativeOptimalityTolerance, 0.03);
```

With this tolerance set, if an optimal solution of a minimization problem is found with an objective value of 900, then there does not exists a solution with an objective value of 873 (= 900 - 900 *0.03). There may exist solutions with objective values of 874 and higher. The default value for this tolerance is 0.0001.

It is important to note that when both a relative and an absolute optimality tolerance are set, they act similarly to constraints, that is only the strongest applies.

Chapter 7. Tuning the CP Optimizer search

CP Optimizer provides a variety of search algorithms for solving constraint programming problems.

Overview

CP Optimizer provides a variety of search algorithms for solving constraint programming problems.

The performance of search is crucial for rapid convergence towards good solutions. CP Optimizer guides the search towards good solutions by examining model structure and observing constraint propagation and then choosing an adapted strategy.

In some cases, tuning the search may provide better performance than the built-in strategy. Tuning the search is possible by setting parameters for selecting other search types and by using tuning objects to provide additional structural information to the search. The latter is done by specifying search phases. Search phases provide a simple way to specify important variables, hierarchy between groups of variables and also the strategy for selecting decision variables to fix and the values to which the variables should be fixed. This section discusses how to set parameters on the optimizer, define evaluators for use in the search phase tuning objects and use multi-point search algorithms.

Using presolve

Presolve consists in modifying the model to improve it so that it can be solved more efficiently.

CP Optimizer presolves the input model before search. Presolve consists in modifying the model to improve it so that it can be solved more efficiently.

Presolve works by

- simplifying the model by reducing linear constraints by removing redundancies and fixed expressions and
- strengthening the model to achieve more domain reduction by aggregating constraints or factorizing common subexpressions.

Presolve can be helpful on models that are not well formulated due to the model itself or to the model data. The activation of presolve is controlled by a parameter. In the **C++ API**, the parameter is `IloCP::Presolve` which can be set to `IloCP::On` (the default) or `IloCP::Off`. In the **Java API**, the parameter is `IloCP.IntParam.Presolve`. In the **C# API**, the parameter is `CP.IntParam.Presolve`.

As a consequence, the overall engine memory consumption can increase because an internal model is created to perform the improvement operations. The presolve transformations are displayed in the output log of CP Optimizer. For example the line

```
! Presolve      : 6942 extractables eliminated, 158 constraints generated
```

at the beginning of the search shows that presolve as eliminated 6942 constraints or expressions and that a subset of constraints were aggregated to 158 constraints.

Using alternative search types

The alternative search types include depth-first search, restart search and multi-point search.

Overview

CP Optimizer using constructive search to build a solution along with other heuristics to improve search.

The CP Optimizer search is based on *constructive search*, which is a search technique that attempts to build a solution by fixing decision variables to values. While the built-in CP Optimizer search is based on this technique, the optimizer also uses other heuristics to improve search.

Three types of search are available in CP Optimizer: "restart", "depth-first" and "multipoint". The default search is "restart"; however, the type of the search can be changed either to improve performance or to debug a model or a strategy. The search type parameter controls the type of search applied to a problem.

In the **C++ API** of CP Optimizer, you use the method `IloCP::setParameter` to set the parameter `IloCP::SearchType`. This parameter has a default value of `IloCP::Restart`. Other values of this parameter are `IloCP::DepthFirst` and `IloCP::MultiPoint`.

In the **Java API** of CP Optimizer, you use the method `IloCP.setParameter` to set the parameter `IloCP.IntParam.SearchType`. This parameter has a default value of `IloCP.ParameterValues.Restart`. Other values of this parameter are `IloCP.ParameterValues.DepthFirst` and `IloCP.ParameterValues.MultiPoint`.

Likewise, in the **C# API** of CP Optimizer, you use the method `CP.SetParameter` to set the parameter `CP.IntParam.SearchType`. This parameter has a default value of `CP.ParameterValues.Restart`. Other values of this parameter are `CP.ParameterValues.DepthFirst` and `CP.ParameterValues.MultiPoint`.

Depth-first search

Depth-first search is a tree search algorithm such that each fixing, or instantiation, of a decision variable can be thought of as a branch in a search tree.

The depth-first search type applies constructive search directly. Depth-first search is a tree search algorithm such that each fixing, or instantiation, of a decision variable can be thought of as a branch in a search tree. The optimizer works on the subtree of one branch until it has found a solution or has proven that there is no solution in that subtree. The optimizer will not move to work on another section of the tree until the current one has been fully explored.

This type of search is quite useful while debugging your model and tuning the search, but generally will be less efficient than restart search because it does not easily recover from poor branching decisions.

Consider the optimization model implemented with the C++ API of Concert Technology:

```
    IloIntVar x(env, 0, 7, "x");
    IloIntVar y(env, 0, 7, "y");
    IloIntVar z(env, 0, 7, "z");
    IloIntVarArray all(env, 3, x, y, z);
    model.add(IloMinimize(env, IloSum(all)));
    model.add(IloAllDiff(env, all));
    model.add(y == IloElement(IloIntArray(env, 7, 3, 7, 8, 8, 0, 1, 4), x));
```

Depth-first search can be applied to this problem with the following code:

```
    IloCP cp(model);
    cp.setParameter(IloCP::SearchType, IloCP::DepthFirst);
    if (cp.solve()){
      cp.out() << " An optimal solution has been found"
               << ", objective value = " << cp.getObjValue()
               << ", x = " << cp.getValue(x)
               << ", y = " << cp.getValue(y)
               << ", z = " << cp.getValue(z) << std::endl;
    }
```

Running this code produces the output:

```
 ! ----------------------------------------------------------------------------
 ! Minimization problem - 3 variables, 2 constraints
 ! Workers             = 2
 ! SearchType          = DepthFirst
 ! Initial process time : 0.00s (0.00s extraction + 0.00s propagation)
 !  . Log search space  : 7.6 (before), 7.6 (after)
 !  . Memory usage       : 288.9 kB (before), 288.9 kB (after)
 ! Using parallel search with 2 workers.
 ! ----------------------------------------------------------------------------
 !        Best Branches  Non-fixed   W      Branch decision
                     0          3                  -
 + New bound is 0
 *          4          2  0.00s       1      (gap is 100.0%)
 ! ----------------------------------------------------------------------------
 ! Search completed, 1 solution found.
 ! Best objective        : 4 (optimal - effective tol. is 0)
 ! Best bound            : 0
 ! Number of branches    : 4
 ! Number of fails       : 2
 ! Total memory usage    : 718.3 kB (676.3 kB CP Optimizer + 42.0 kB Concert)
 ! Time spent in solve   : 0.00s (0.00s engine + 0.00s extraction)
 ! Search speed (br. / s) : 400.0
 ! ----------------------------------------------------------------------------
 An optimal solution has been found, objective value = 4, x = 0, y = 3, z = 1
```

Restart search

Restart search is the default search. The constructive search is restarted from time to time and guided towards an optimal solution.

The default search is called "restart"; in this type of search, the constructive search is restarted from time to time and guided towards an optimal solution.

In restart search, a depth-first search is restarted after a certain number of failures. The parameter `RestartGrowthFactor` controls the increase of this number between restarts. If the last fail limit was *f* after a restart, then, for next run, the new fail limit will be *f* times the value of this parameter. The initial fail limit can be controlled with the parameter `RestartFailLimit`.

In the **C++ API** of CP Optimizer, you use the method `IloCP::setParameter` and the values `IloCP::RestartGrowthFactor` and `IloCP::RestartFailLimit`.

In the **Java API** of CP Optimizer, you use the method `IloCP.setParameter` and the values `IloCP.DoubleParam.RestartGrowthFactor` and `IloCP.IntParam.RestartFailLimit`.

In the **C# API** of CP Optimizer, you use the method `CP.SetParameter` and the values `CP.DoubleParam.RestartGrowthFactor` and `CP.IntParam.RestartFailLimit`.

Multi-point search

Multi-point search creates a set of solutions and combines the solutions in the set in order to produce better solutions.

The third value for the search type parameter is for "multi-point" search. This search creates a set of solutions and combines the solutions in the set in order to produce better solutions.

Multi-point search is more diversified than depth-first or restart search, but it does not necessarily prove optimality or the inexistence of a solution. The search runs until the optimizer considers that the best solution found cannot be improved. Therefore it is recommended to set up a limit when using multi-point search.

For instance, to use multi-point search on this model, you can write:

```
IloCP cp(model);
cp.setParameter(IloCP::SearchType, IloCP::MultiPoint);
cp.setParameter(IloCP::BranchLimit, 10000);
if (cp.solve()){
  cp.out() << " A solution has been found"
          << ", objective value = " << cp.getObjValue()
          << ", x = " << cp.getValue(x)
          << ", y = " << cp.getValue(y)
          << ", z = " << cp.getValue(z) << std::endl;
}
```

Running this code produces the output:

```
! --------------------------------------------------------------------------
! Minimization problem - 3 variables, 2 constraints
! BranchLimit         = 10000
! Workers             = 2
! SearchType          = MultiPoint
! Initial process time : 0.00s (0.00s extraction + 0.00s propagation)
!  . Log search space  : 7.6 (before), 7.6 (after)
!  . Memory usage      : 288.9 kB (before), 288.9 kB (after)
! Using parallel search with 2 workers.
! --------------------------------------------------------------------------
!        Best Branches  Non-fixed   W      Branch decision
                    0          3            -
+ New bound is 0
*          10      22  0.00s      1      (gap is 100.0%)
*           9      65  0.00s      1      (gap is 100.0%)
*           5      86  0.00s      1      (gap is 100.0%)
*           4     248  0.00s      1      (gap is 100.0%)
! --------------------------------------------------------------------------
! Search completed, 4 solutions found.
! Best objective         : 4 (optimal - effective tol. is 0)
! Best bound             : 0
! Number of branches     : 248
! Number of fails        : 123
! Total memory usage     : 786.5 kB (744.5 kB CP Optimizer + 42.0 kB Concert)
! Time spent in solve    : 0.00s (0.00s engine + 0.00s extraction)
! Search speed (br. / s) : 24800.0
! --------------------------------------------------------------------------
 A solution has been found, objective value = 4, x = 0, y = 3, z = 1
```

For more information on multi-point search, refer to "Using multi-point search algorithms" on page 85.

Setting parameters for directing the search

CP Optimizer provides parameters that can be set in order to direct the search.

You can set parameters to choose the search strategy and to tune the optimizer.

For example, to switch the search strategy from the default search to depth-first search, you set the `SearchType` parameter to `DepthFirst`.

The search uses randomization in some strategies. The parameter `RandomSeed` sets the seed of the random generator used by these strategies.

Parameters may not be changed while there is an active search.

Ordering variables and values

Decision variables and their possible values can be ordered so that the optimizer can fix the key decision variables early in the process.

Grouping variables

Decision variables and their possible values can be ordered so that the optimizer can fix the key decision variables early in the process.

In many applications, there exists a group of key decision variables, such that once these variables are fixed, it is easy to extend the partial solution to the remaining variables.

Information about key variables can be given to the search by way of tuning objects called search phases. An instance of `IloSearchPhase` (`ISearchPhase` in the **C# API**) is created with an array of decision variables and passed to the search as argument to the method `setSearchPhases` before calling `solve` and `startNewSearch`.

Assume that in a model the decision variables in the array x are key variables. In the C++ API, you can pass this information to the search the following way:

```
cp.setSearchPhases(IloSearchPhase(env,x));
cp.solve();
```

This search phase forces the search to fix (instantiate) the decision variables from the array x before instantiating any other variable in the model.

Instantiation of groups of decision variables can be ordered by using several search phases. The search phases are passed to the search using a search phase array. The decision variables in the first phase are instantiated before the variables in the second one and so on.

Consider for instance, the two phases:

```
IloSearchPhase xPhase(env, x);
IloSearchPhase yPhase(env, y);
```

If we solve a model by calling:

```
        IloSearchPhaseArray phaseArray(env);
        phaseArray.add(xPhase);
        phaseArray.add(yPhase);
        cp.setSearchPhases(phaseArray);
        cp.solve();
```

Decision variables in x will be instantiated before variables in y that in turn will be instantiated before any variables that are not in x and y.

It is important to observe that when using search phases, the phases do not need to cover all variables of the model. The CP Optimizer search will instantiate all variables, and those that do not appear in any search phase will always be instantiated last.

Defining a constructive strategy

A constructive strategy can be defined using search phases. A search phase specifies the criteria for the order in which variables and values are chosen in the search.

In addition to variables, a search phase can be used to tune a search strategy by specifying the criteria for the order in which decision variables are chosen to be fixed and to which values these variables should be fixed. This strategy is then used as a constructive strategy to instantiate the decision variables of the phase.

In the **C++ API**, the constructor of a complete phase is:

```
  IloSearchPhase(IloIntVarArray x,
                 IloIntVarChooser varChooser,
                 IloIntValueChooser valueChooser);
```

The variable chooser (`IloIntVarChooser` in the **C++ API** and the **Java API**, `IIntVarChooser` in the **C# API**) defines how the next decision variable to fix in the search is chosen.

The value chooser (`IloIntValueChooser` in the **C++ API** and the **Java API**, `IIntValueChooser` in the **C# API**) defines how values are chosen for instantiating decision variables.

Note:

Search phase

A search phase is an object that is used to define instantiation strategies to help the CP Optimizer search. A search phase is composed of
• an array of decision variables to instantiate (or fix),
• a variable chooser that defines how the next variable to instantiate is chosen
• a value chooser that defines how values are chosen when instantiating variables.

Since V12.8, the `IloCPEngine` class is used in evaluators used by choosers. It provides access to the extracted model objects. See details about `IloCPEngine` in the CP Optimizer reference manuals.

Simple variable selection

To chose decision variables, you evaluate the variables with an evaluator.

To chose decision variables, you evaluate the variables with an evaluator.

In the C++ API, a decision variable evaluator is an instance of `IloIntVarEval`. This class implements the function:

```
IloNum IloIntVarEvalI::eval(IloCPEngine cp, IloIntVar x);
```

that returns the evaluation of the variable x.

Several predefined evaluators exist in CP Optimizer. For instance, the evaluator returned by the C++ API function call `IloDomainSize(env)` returns the current size of the domain of the variable that is being evaluated. Another example is the evaluator returned by `IloVarIndex(env, vars)` that returns the index in the array vars of the evaluated variable.

In order to select a decision variable with an evaluator, an instance of a variable selector needs to be created with the evaluator as an argument. Here are two functions in the C++ API that can create such selectors:

```
IloVarSelector IloSelectSmallest(IloIntVarEval e);
IloVarSelector IloSelectLargest(IloIntVarEval e);
```

The selector created by `IloSelectSmallest(IloDomainSize(env))` will select the variable in the search phase that has the smallest domain. The selector returned by `IloSelectRandomVar(env)` chooses a variable randomly. This selector is useful for breaking ties.

In the **C++ API** of CP Optimizer, you use the classes `IloIntVarEval` and `IloVarSelector` and the functions `IloDomainSize`, `IloVarIndex`, `IloSelectSmallest` and `IloSelectRandomVar`.

In the **Java API** of CP Optimizer, you use the interfaces `IloIntVarEval` and `IloVarSelector` and the methods `IloCP.domainSize`, `IloCP.varIndex`, `IloCP.selectSmallest` and `IloCP.selectRandomVar`.

In the **C# API** of CP Optimizer, you use the interfaces `IIntVarEval` and `IVarSelector` and the methods `CP.DomainSize`, `CP.VarIndex`, `CP.SelectSmallest` and `CP.SelectRandomVar`.

Simple value selection

To chose the value at which you fix a decision variable, you evaluate the values with an evaluator.

As with decision variables, values are evaluated with an evaluator.

In the C++ API, a value evaluator is an instance of `IloIntValueEval`. This class implements the function:

```
IloNum IloIntValueEvalI::eval(IloCPEngine cp, IloIntVar x, IloInt v)
```

that returns the evaluation of assigning the value v to the variable x.

Several predefined value evaluators exist in CP Optimizer. For instance, the evaluator returned by the C++ API function call `IloValue(env)` returns the value itself. Another example is the evaluator returned by `IloExplicitValueEval(env, valueArray, evalArray)` that returns the evaluation `evalArray[i]` when evaluating `valueArray[i]`.

In order to select a value with an evaluator, an instance of a value selector needs to be created with the evaluator as argument. Here are two functions in the C++ API that can create such selectors:

```
IloValueSelector IloSelectSmallest(IloIntValueEval e);
IloValueSelector IloSelectLargest(IloIntValueEval e);
```

The selector created by `IloSelectLargest(IloValue(env))` will select the largest value in the domain of the selected variable. The selector returned by `IloSelectRandomValue(env)` chooses a value randomly. This selector is useful for breaking ties.

In the **C++ API** of CP Optimizer, you use the classes `IloIntValueEval` and `IloValueSelector` and the functions `IloValue`, `IloExplicitValueEval`, `IloSelectLargest` and `IloSelectRandomValue`.

In the **Java API** of CP Optimizer, you use the interfaces `IloIntValueEval` and `IloValueSelector` and the methods `IloCP.value`, `IloCP.explicitValueEval`, `IloCP.selectLargest` and `IloCP.selectRandomValue`.

In the **C# API** of CP Optimizer, you use the interfaces `IIntValueEval` and `IValueSelector` and the methods `CP.Value`, `CP.ExplicitValueEval`, `CP.SelectLargest` and `CP.SelectRandomValue`.

Multi-criteria selection

To break ties in selecting values or variables, a multi-criteria selector is needed.

If it appears that several decision variables have the same smallest or largest evaluation, the one with the smallest index in the array given to the search phase is selected. Similarly for the selection of domain values, ties are broken by choosing the smallest value. To break ties using a different criteria, you need to use multi-criteria selection.

Assume you want to select the decision variable having the smallest domain and break ties with another selector by selecting the variable having the smallest minimum value in its domain.

For this purpose, you can select variables with an array of selectors. To implement the selection above in the C++ API, you need to create the array:

```
IloVarSelectorArray varSelArray(env);
varSelArray.add(IloSelectSmallest(IloDomainSize(env)));
varSelArray.add(IloSelectSmallest(IloDomainMin(env)));
```

Then a search phase whose variable chooser is the array of selectors is:

```
IloSearchPhase phase(env, varSelArray, IloSelectSmallest(IloValue(env)));
```

Furthermore, you can force the creation of ties by selecting several decision variables at the selection stage. For instance, the following selector array selects at least five variables among those having the smallest domain and break ties randomly:

```
IloVarSelectorArray varSelArray(env);
varSelArray.add(IloSelectSmallest(5, IloDomainSize(env)));
varSelArray.add(IloSelectRandomVar(env));
```

Multi-criteria selection can also be applied to value selection. For instance, the following array of value selectors selects the 5 smallest values in the domain of the decision variable and then chooses randomly among those values:

```
                IloValueSelectorArray valueSelArray(env);
                valueSelArray.add(IloSelectSmallest(5, IloValue(env)));
                valueSelArray.add(IloSelectRandomValue(env));
```

Search phases with selectors

A selector can be used with a search phase to indicate which variable should be instantiated to which value.

A search phase that instantiates the decision variables in the array x by choosing the variable having the smallest domain size and by assigning it to the largest value of its domain can be defined like this in the C++ API:

```
        IloSearchPhase phase(env, x,
                        IloSelectSmallest(IloDomainSize(env)),
                        IloSelectLargest(IloValue(env)));
```

A search phase or an array of search phases defined with selectors can be given to an instance of the optimizer by invoking the method setSearchPhases before calling solve or startNewSearch:

```
        cp.setSearchPhases(phase);
        cp.solve();
```

Defining your own evaluator

CP Optimizer provides an API to define custom evaluators.

CP Optimizer lets you control the order in which the values in the domain of a decision variable are tried during the search for a solution through the use of search phases.

In that way, you can exploit strategic information that you have about the problem to guide the search for a solution. There are a few predefined evaluators for the variable chooser and value chooser. However, you may encounter problems for which you would like to define your own evaluators.

An evaluator of integer variables is an object that is used by selectors of variables to define instantiation strategies. In the C++ API, an evaluator of integer variables is defined by implementing the pure virtual member function IloNum IloIntVarEvalI::eval(IloCPEngine cp, IloIntVar x) that returns a floating-point evaluation of the integer variable x.

An evaluator of integer values is an object that is used by selectors of values to define instantiation strategies. In the C++ API, an evaluator of integer value assignments is defined by implementing the pure virtual member function IloNum IloIntValueEvalI::eval(IloCPEngine cp, IloIntVar x, IloInt value) that returns an evaluation of fixing the decision variable x to value.

For example, assume that you want to control the order in which values are assigned to constrained integer variables, and in particular, you want the order of values chosen to be 0, -1, +1, -2, +2 and so forth. The following code uses an evaluator and value chooser to implement that approach.

```
class AbsValI : public IloIntValueEvalI {
public:
  AbsValI(IloEnv env) : IloIntValueEvalI(env) { }
  IloNum eval(IloCPEngine cp, IloIntVar x, IloInt value) {
    return IloAbs(value);
  }
};
```

```
IloIntValueEval AbsVal(IloEnv env) {
  return new (env) AbsValI(env);
}
```

It is a good idea then to test that code in a program like this:

```
IloEnv env;
try {
  IloModel model(env);
  IloIntVarArray x(env, 2, -2, 4);
  model.add(x);

  IloCP cp(model);
  cp.setParameter(IloCP::LogVerbosity,IloCP::Quiet);
  IloVarSelectorArray varSel(env);
  varSel.add(IloSelectRandomVar(env));
  IloValueSelectorArray valSel(env);
  valSel.add(IloSelectSmallest(AbsVal(env)));
  IloSearchPhase phase(env, x, varSel, valSel);
  cp.setSearchPhases(phase);
  cp.startNewSearch();
  while (cp.next()) {
    cp.out() << cp.getValue(x[0]) << " " << cp.getValue(x[1]) << std::endl;
  }
  cp.end();
} catch (IloException & ex) {
  env.out() << "Caught " << ex << std::endl;
}
env.end();
```

In the **C++ API** of CP Optimizer, you use the classes `IloIntVarEvalI` and
`IloIntValueEvalI` and the functions `IloIntVarEvalI::eval` and
`IloIntValueEvalI::eval`.

In the **Java API** of CP Optimizer, you use you use the classes `IloCustomIntVarEval`
and `IloCustomIntValueEval` and the functions `IloCustomIntVarEval.eval` and
`IloCustomIntValueEval.eval`.

At this time, custom evaluators are not available in the **C# API** of CP Optimizer.

Search phases for scheduling

Search phases can be used on interval variables as well as on sequence variables.

Two types of search phases are available for scheduling: search phases on interval
variables and search phases on sequence variables.

A search phase on interval variables works on a unique interval variable or on an
array of interval variables. During this phase CP Optimizer fixes the value of the
specified interval variable(s): each interval variable will be assigned a presence
status and for each present interval, a start and an end value. This search phase
fixes the start and end values of interval variables in an unidirectional manner,
starting to fix first the intervals that will be assigned a small start or end value.

A search phase on sequence variables works on a unique sequence variable or on
an array of sequence variables. During this phase CP Optimizer fixes the value of
the specified sequence variable(s): each sequence variable will be assigned a totally
ordered sequence of present interval variables. Note that this search phase also
fixes the presence statuses of the intervals involved in the sequence variables. This
phase does not fix the start and end values of interval variables. It is recommended

to use this search phase only if the possible range for start and end values of all interval variables is limited (for example by some known horizon that limits their maximal values).

Using multi-point search algorithms

Multi-point search is based on a pool of points. This pool can be managed via parameters.

One search technique available in CP Optimizer is a multi-point search algorithm. This algorithm is based on a pool of *search points*. A search point is a collection of decision variable assignments that may lead to either feasible or partial solutions (a partial solution has some variables which are still not fixed). The multi-point method starts with an initial pool of search points whose candidate assignments are generated with constructive search. It then produces new search points by learning new variable assignments from search points maintained in the pool. On an optimization problem, the multi-point search method is able to learn from partial solutions in order to produce feasible solutions.

Multi-point search produces solutions by first performing variable assignments proposed by each of the generated search points. It then attempts to complete the solution by invoking a tree-search based completion procedure. If no feasible solution can be produced, a solution with a maximal number of instantiated decision variable is retained.

The completion procedure used is basically the same as the one used by the restart method. If you have specified search phases, the multi-point search completion procedure will use the search phases as well. In addition, phase priorities will be respected when interpreting search point assignments. That is, if you have specified phases A, B and C, then, for each search point, assignments involving variables of A will be performed at first, followed by those of phase B, ending with those of phase C.

The search parameter `MultiPointNumberOfSearchPoints` controls the number of (possibly partial) solutions manipulated by the multi-point search algorithm. The default value is 30. A larger value will diversify the search, with possible improvement in solution quality at the expense of a longer run time. A smaller value will intensify the search, resulting in faster convergence at the expense of solution quality. Note that memory consumption increases proportionally to this parameter, for each search point must store each decision variable domain.

To use multipoint search in the **C++ API** of CP Optimizer, you use the method `IloCP::setParameter` and the values `IloCP::SearchType`, `IloCP::MultiPoint` and `IloCP::MultiPointNumberOfSearchPoints`.

To use multipoint search in the **Java API** of CP Optimizer, you use the method `IloCP.setParameter` and the values `IloCP.IntParam.SearchType`, `IloCP.ParamterValues.MultiPoint` and `IloCP.IntParam.MultiPointNumberOfSearchPoints`.

To use multipoint search in the **C# API** of CP Optimizer, you use the method `CP.SetParameter` and the values `CP.IntParam.SearchType`, `CP.ParamterValues.MultiPoint` and `CP.IntParam.MultiPointNumberOfSearchPoints`.

Solving lexicographic multi-objective problems

Providing a starting point can sometimes help the optimizer produce solutions more quickly.

There are cases where better solutions can be produced more quickly by providing a starting point, an instance of `IloSolution`, to the optimizer. In goal programming, the multi-objective optimization problem may involve a lexically ordered set of objective functions *(f1,f2,...,fn)*. It could be, for example, a detailed scheduling problem for which the main objective, *f1*, is to minimize resource allocation costs whereas a secondary objective, *f2*, is to minimize the makespan of the schedule given an optimal or good resource allocation cost. In this case, the problem can be solved in n successive steps: first, minimize objective *f1* to produce a solution *sol1*, then, add a constraint to avoid deteriorating *f1* and solve the problem with objective function *f2* using *sol1* as a starting point to produce a solution *sol2*, etc. The solution to a given step is a feasible solution for the next step. Setting a starting point may improve the performance of the optimizer engine. A typical code in the C++ API for implementing a multi-objective problem using starting points would be similar to:

```
IloEnv env;
IloModel model(env);
IloInt n = ...;
IloIntervalVarArray activities(env, n);
// ...

IloIntExpr f1 = ...;
IloIntExpr f2 = ...;
IloObjective obj1 = IloMinimize(env, f1);
IloObjective obj2 = IloMinimize(env, f2);
IloCP cp(model);

// Minimize f1
model.add(obj1);
cp.solve();

// Store solution
IloSolution sol1(env);
for (IloInt i=0; i<n; ++i) {
  // For illustration purpose, we only save start values
  sol1.setStart(activities[i], cp.getStart(activities[i]));
}

// Change objective
model.remove(obj1);
model.add(f1 <= cp.getValue(f1)); // f1 should not worsen
model.add(obj2);

// Minimize f2 using sol1 as starting point
cp.setStartingPoint(sol1);
cp.solve();
```

The starting point provided to the engine does not have to specify a value for each decision variable (it can specify a range of values or no information at all) and does not have to be a feasible solution for the problem being solved. If the starting point provides a fixed value for each decision variable of the problem and if it is feasible, the CP Optimizer search will first visit this solution when traversing the search space. In all other cases, the information contained in the starting point is used as a guideline for the search but there is no guarantee that the solutions traversed by the search will be "close" to the starting point solution.

Note: the starting point information is used by the restart and multi-point search types only. It is not used by the depth-first search.

To set a starting point in the **C++ API** of CP Optimizer, you use the method `IloCP::setStartingPoint`.

To set a starting point in the **Java API** of CP Optimizer, you use the method `IloCP.setStartingPoint`.

To set a starting point in the **C# API** of CP Optimizer, you use the method `CP.SetStartingPoint`.

Chapter 8. Designing models

While developing models for CP Optimizer can be straightforward, there are some principles that you should consider while working on a model.

Overview

Considering some design principles while creating a model for CP Optimizer can help the optimizer run more efficiently.

The principles described are meant to help you avoid the errors often made by new users of CP Optimizer when they design a model for a problem. Of course, not every problem benefits from a mechanical application of every principle mentioned here, but in general these principles should help you develop robust and efficient CP Optimizer programs.

Decrease the number of variables

Decreasing the number of variables, and thus reducing the size of the search space, is one model design principle to consider.

The unknowns of a given problem are typically represented in the model by decision variables. There are practical ways of decreasing the number of variables and thus reducing the size of the model and its search space.

Problems best solved with constraint-based programming are generally subject to intrinsic combinatorial growth. Even if reducing the domains of variables by propagating constraints makes it possible to reduce the search space, the initial size of this search space still remains a weighty factor in the execution time.

Principle

Consequently, good practice in designing a model should attempt to minimize the size of the search space in the first place. This size increases exponentially with the number of variables. Thus, limiting the number of such variables (even at the expense of enlarging their domains) can reduce the combinatorial complexity.

Example

This principle of reducing the number of decision variables can often be applied advantageously to resource allocation problems. For example, assume that C consumers must choose among R resources where:
- all the resources are available to every consumer;
- if consumer i chooses resource j, he or she incurs $cost[i,j]$;
- each consumer uses at most one resource.

First model

This problem is often represented in the following way:
- Create $C*R$ constrained integer variables $supplier_{i,j}$ with domain [0, 1] such that $supplier_{i,j} = 1$ if consumer i uses resource j.

- The constraints stating that each consumer uses at most one resource are represented this way: for each i, $\sum_{j=0}^{j=R-1} supplier_{i,j} \leq 1$

- The goal is to maximize $\sum_{i=0}^{i=C-1} \sum_{j=0}^{j=R-1} cost[i,j] \cdot supplier_{i,j}$

To evaluate the combinatorial complexity of the problem, consider the maximum number of configurations, called the *apparent complexity* of the problem. This figure is the size of the search space, that is, the worst case complexity of a generate and test algorithm.

In this model, the apparent complexity is 2^{R*C}, which is around 10^{30} if $R=C=10$.

Second model: using fewer variables

With CP Optimizer, a more efficient model can be represented. The alternate model can be written this way:

- Create a fake resource numbered 0.
- Create C constrained integer variables $supplier_i$ with domain [0..R] so that

$supplier_i = 0$ if consumer i uses no resource,

$supplier_i = j$ if consumer i uses resource j.

The constraints stating that each consumer uses at most one resource are necessarily satisfied, since a constrained integer variable can be fixed with only one value.

- The goal is to maximize $\sum_{i=0}^{C-1} cost[i,supplier_i]$.

The maximum number of solutions of this representation is $(R+1)^C$, which is 11^{10} if $R=C=10$. This value is much smaller than 10^{30} from the first model.

Use dual representation

Using dual representation to consider the problem for a different view is a model design principle to consider.

Object programming and straightforward object manipulation in CP Optimizer often make it possible to design a direct and very intuitive model of a problem. Nevertheless, it is important not to get stuck in this apparent simplicity if that model gives rise to a disproportionate amount of propagation or outright inefficiency.

Principle

Dual representation consists of looking at a problem "from the other side," so to speak. This technique assumes that the designer knows enough about the problem to step back and consider its essence and thus extract conceptual symmetries inherent in it.

Example

Allocating resources offers a good example of this phenomenon. Assume that C consumers must choose among R resources; to make it simple, consider the case where:

- all the resources are available to every consumer;
- each consumer uses at most one resource;
- each resource may be used by at most one consumer.

The apparent complexity of the problem depends on the kind of model. In practice, there are two possible models:

- The consumers choose resources.
- The resources choose consumers.

In the first case, if a constrained variable representing the chosen resource is associated with each consumer, there are C variables with a domain of size $R+1$, where R is the number of possible resources. The apparent complexity of the problem in this model is thus $(R+1)^C$.

The second case is, of course, analogous to the first: associate a constrained variable with each resource such that the variable represents the chosen consumer, so there are R variables with a domain of size $C+1$, where C is the number of possible consumers. The apparent complexity of the problem in this model is thus $(C+1)^R$.

When the difference in apparent complexity is great, it is important to consider the magnitude of C and R very carefully.

Remove symmetries

Removing symmetries is one model design principle to consider. There are multiple ways in which to remove symmetries.

Overview

Removing symmetries is one model design principle to consider. There are multiple ways in which to remove symmetries. By removing symmetries, parts of the search space can be safely ignored.

The apparent complexity of a problem can often be reduced to a much smaller practical complexity by detecting intrinsic symmetries. Parts of the search space can then be safely ignored.

Group by type

One way to remove symmetries is to group by type.

When two or more variables have identical characteristics, it is pointless to differentiate them artificially.

Principle

Rather, it is preferable to design a model that takes into account not simply the elementary entities but instead the types into which the elementary entities can be meaningfully grouped. Each such type, of course, quantitatively handles the elementary entities that belong to it.

Example

For example, consider the problem of plugging electronic cards into racks. Assume that there are five types of cards. Each rack may contain only a finite number of cards according to some constraints. The objective is to minimize the number of racks used. This problem may be formulated in two different ways:

- first model: for each card, the problem is to find the rack where it is to be plugged;
- second model: for each rack, the problem is to find the number of cards of a given type plugged into it.

In the first model, there is one variable per card. Once a solution has been found, any permutation of two cards of the same type defines a "new" solution. The search space is consequently large and contains many such symmetrical solutions that are not of interest.

In the second model, the number of cards of a given type is manipulated, rather than the cards themselves. This model suppresses the symmetries introduced by the previous one and thus reduces the search space.

Introduce order among variables

One method for removing symmetries is to introduce order among variables.

In some cases, there is really no point in examining all the possible solutions for variables and their values.

This is the case when two or more constrained variables satisfy the following conditions:

- the initial domains of these constrained variables are identical;
- these variables are subject to the same constraints;
- the variables can be permuted without changing the statement of the problem.

In fact, the permutations give rise to sets of solutions that are identical as far as the physical reality of the problem is concerned. This idea can be exploited to minimize the size of the search space.

Principle

If these domains are reduced by introducing a supplementary constraint, such as order, or by imposing a special feature on each of these variables, the size of the search space can be markedly reduced.

Example

Assume, for example, that there is the following system of equations:

$x_1 + x_2 + x_3 = 9$

$x_1 \times x_2 \times x_3 = 12$

For the ordered triple (x_1, x_2, x_3), there are six solutions:

(1, 2, 6) (1, 6, 2) (2, 1, 6) (2, 6, 1) (6, 1, 2) (6, 2, 1)

If the variables are permuted, the problem is not changed. For instance, if x_1 and x_2 are swapped, the problem becomes:

$x_2 + x_1 + x_3 = 9$

$x_2 \times x_1 \times x_3 = 12$

That problem is obviously the same as the first one. In this case, it is a good idea to introduce a supplementary constraint to enforce order among the variables. This order can be introduced in this manner:

$x1 <= x2 <= x3$

The additional constraint on the order among the variables greatly reduces the combinatorial possibilities *without removing any real solutions*.

In fact, only one solution can be returned under these conditions:

$(x_1, x_2, x_3) = (1, 2, 6)$

While removing possibilities wherever possible is a good idea, you should guard against adding a supplementary constraint that inadvertently suppresses solutions that you would like to see.

Introduce surrogate constraints

One method for removing symmetries is to introduce surrogate constraints.

Since constraint propagation decreases the size of the search space by reducing the domains of variables, it is obviously important to express all necessary constraints. In some cases, it is even a good idea to introduce implicit constraints to reduce the size of the search space by supplementary propagation.

Processing supplementary constraints inevitably slows down execution. However, this slowing down may be negligible in certain problems when it is compared with the efficiency gained from reducing the size of the search space.

Principle

A surrogate constraint makes explicit a property that satisfies a solution implicitly. Such a constraint should not change the nature of the solution, but its propagation should delimit the general shape of the solution more quickly.

Of course, there is no need to express grossly obvious redundant constraints since the highly optimized algorithms that CP Optimizer uses to ensure arc consistency already work well enough. For example, given this system of equations:

$x = y + z$

$z = a + b$

no efficiency whatsoever is gained by adding this constraint:

$x = y + a + b$

However, in any case where an implicit property makes good sense, or derives from experience, or satisfies formal computations, its explicit implementation as a surrogate constraint can be beneficial.

Example

Consider the problem of the magic sequence. Assume that there are $n+1$ unknowns, namely, $x_0, x_1, \ldots, x_n$. These x_i must respect the following constraints:

0 appears x_0 times in the solution.

1 appears x_1 times.

In general, i appears x_i times.

n appears x_n times.

The constraint of this problem can easily be written, using the specialized distribute constraint. However, the search for a solution can be greatly accelerated by introducing the following surrogate constraint that expresses the fact that $n+1$ numbers are counted.

1*x1 + 2*x2 + . . . + n*xn = n+1.

Chapter 9. Designing scheduling models

Although developing scheduling models for CP Optimizer can be straightforward, there are some principles that you should consider while working on a model.

Specifying interval bounds

Though there are various methods for restricting bounds on interval variables, specifying the bounds in the declaration of the interval variable is recommended.

If you have to specify a minimal start time, a maximal end time or a range of possible values for the size of an interval, it is recommended specify these values in the declaration of the interval itself rather than through expressions `IloStartOf`, `IloEndOf`, and `IloSizeOf`.

Specifying the values at the time of declaration avoids difficulties related to the optionality of intervals variables. For instance, the following code segments are not equivalent:

```
IloIntervalVar a(env, 10, 20, IloTrue);
```

and

```
IloIntervalVar a(env);
a.setOptional();
m.add(IloSizeOf(a) >= 10);
m.add(IloSizeOf(a) <= 20);
```

The first sample specifies a range for the size of the interval variable if the interval is present. In particular, the model will be consistent even if a is absent. The second sample will be inconsistent if a is absent because the default value of `IloStartOf(a)` will be 0 if a is absent. An equivalent model would be something like:

```
IloIntervalVar a(env);
a.setOptional();
m.add(IloSizeOf(a,10) >= 10);
m.add(IloSizeOf(a,0) <= 20);
```

Additionally, specifying the range at the declaration of the interval is more effective in the optimizer.

Specifying precedence relations between interval variables

Though there are various methods for modeling a precedence between two interval variables, using a precedence constraint is recommended.

When modeling a precedence between two intervals, it is always better to use a precedence constraint (e.g. `IloEndBeforeStart`) rather than an arithmetical constraint (<=,<,==) between end and start expressions.

Using a precedence constraint avoids difficulties related with the optionality of intervals variables. For instance, `IloEndBeforeStart(env,a,b)` is generally not equivalent to `IloEndOf(a) <= IloStartOf(b)`. Given the precedence constraint `IloEndBeforeStart(env,a,b)`, if b is absent, then the constraint will be always true and have no impact on a, which is what is usually needed. Given the constraint

IloEndOf(a) <= IloStartOf(b), if b is absent, then the constraint IloEndOf(a) <= 0 will be enforced, as 0 is the default value for IloStartOf(b) when b is absent. The form of a constraint using expressions that is equivalent to the precedence constraint would be IloEndOf(a,IloIntervalMin) <= IloStartOf(b,IloIntervalMax).

Additionally, using a precedence constraint is more effective in the optimizer.

```
IloIntervalVar a(env);
a.setOptional();
IloIntervalVar b(env);
b.setOptional();
m.add(IloEndBeforeStart(env,a,b));
```

This model is not equivalent to:

```
IloIntervalVar a(env);
a.setOptional();
IloIntervalVar b(env);
b.setOptional();
m.add(IloEndOf(a) <= IloStartOf(b));
```

Modeling resource calendars

Resource calendars can be modeled using a stepwise function that describes the intensity of "work" over time.

Many scheduling problems, especially in project scheduling, involve calendars for resources.

These calendars specify periods of time such that:
1. the resource is not available and suspends a task it is working on,
2. the resource works with a limited efficiency so that a task requiring this resource will take more time during these periods,
3. the resource cannot start or finish executing a task during such a period,
4. a task requiring the resource cannot overlap such a period or
5. a task can overlap these periods but the total overlap must be retrieved as an integer expression (typically to be minimized as part of the cost).

In CP Optimizer, this notion of a resource calendar is represented by a stepwise function that describes the intensity of "work" over time. An interval variable can be associated with an integer stepwise intensity function with possible values expressed as a percentage in [0,100]. The intensity function represents an instantaneous ratio between the size and the length of an interval variable. Cases (1) and (2) can be represented by an intensity function. Case (1) is a special case with the intensity function being equal to 0 on the time periods the resource suspends the execution of tasks. Other constraints can be modeled using the constraints IloForbidStart and IloForbidEnd (for case 3) and IloForbidExtent (for case 4).

The following sample defines an intensity function equal to 100% except for the weekends (intervals [5+7i,7+7i]) for which the intensity is 0%, that is, the resource suspends its tasks on weekends.

Task task1 will be suspended during weekends.

Task task2 will be suspended during weekends and cannot start during a weekend.

Task `task3` will be suspended during weekends and cannot start or end during a weekend.

Task `task4` must be completely executed during a week so that it cannot overlap a weekend.

Task `task5` of length 7 days can overlap weekends but the total duration of the task performed during weekends (ov5) is to be minimized.

```
IloNumToNumStepFunction we(env, 0, 364, 100);
for (IloInt w=0; w<52; ++w)
  we.setValue(5+(7*w), 7+(7*w), 0);
IloIntervalVar task1(env,10);
task1.setIntensity(we);
IloIntervalVar task2(env,10);
task2.setIntensity(we);
IloIntervalVar task3(env,10);
task3.setIntensity(we);
IloIntervalVar task4(env,4);
IloIntervalVar task5(env);
task5.setIntensity(we);
task5.setEndMax(365);
IloIntExpr ov5 = 7-IloSizeOf(task5);
m.add(IloForbidStart(env, task2, we));
m.add(IloForbidStart(env, task3, we));
m.add(IloForbidEnd(env, task3, we));
m.add(IloForbidExtent(env, task4, we));
m.add(IloLengthOf(task5)==7);
m.add(IloMinimize(env, ov5));
```

Chains of optional intervals

Though there are various methods for modeling a chains of optional interval variables, an efficient method is recommended.

Sometimes it is necessary to model a chain of n optional intervals for which, only the first k $(k<=n)$ will be present where k is an implicit decision of the problem.

For instance, this is useful for modeling a preemptive activity that can be split into at most n parts. In the following sample, there are the additional constraints that each "part" of the activity has specified a minimal (`pmin`) and a maximal duration (`pmax`) and that the total duration (size) of the parts must equal the processing time `pt` of the preemptive activity. Note that when the part i is absent, the value returned by `IloSizeOf(part[i])` will be 0 (this is the default value when no argument is passed to the expression), thus it will not be counted in the sum.

```
IloIntExpr totalSize(env);
IloIntervalVarArray part(env,n);
part[0] = IloIntervalVar(env, pmin, pmax, IloTrue);
totalSize += IloSizeOf(part[0]);
for (IloInt i=1; i < n-1; i++) {
  part[i] = IloIntervalVar(env, pmin, pmax,IloTrue);
  totalSize += IloSizeOf(part[i]);
  m.add(IloIfThen(env,IloPresenceOf(env,part[i]),IloPresenceOf(env,part[i-1])));
  m.add(IloEndBeforeStart(env,part[i-1],part[i]));
}
m.add(totalSize == pt);
```

Another example is a set of at most n flexible shifts for a worker with specific constraints on the shift duration and minimal resting time between shifts (see "Different uses of the alternative constraint" on page 98).

Different uses of the alternative constraint

The alternative constraint between interval variables should be used when an interval variable represents a set of alternative possibilities.

Alternative possibilities that can be represented by alternative constraints, include an activity that can be executed in a number of possible modes, or a discrete set of possible positions in time.

A common case is an activity a that requires a resource among a set of candidate resources `R={r_1,...,r_m}`. An optional interval a_i, 1<=i<=m can be defined that represents the possible selection of resource r_i in the alternative resource set, R. Then, an alternative constraint `IloAlternative(env,a,[a_i])` will enforce that if a is present, then one and only one of the a_i will be present. This case is illustrated here with disjunctive resources (for more details on how to model disjunctive resources, see "Modeling classical finite capacity resources" on page 101).

```
IloInt nbMachines  = 5;
IloInt nbActivities = 10;
IloIntArray2 ptime = ...; // data from file
IloIntervalVarArray activity(env,nbActivities);
IloIntervalVarArray2 actOnMach(env,nbActivities);
for (IloInt i=0; i < nbActivities; i++) {
  activity[i] = IloIntervalVar(env);
  actOnMach[i] = IloIntervalVarArray(env, actOnMach);
  for (IloInt j=0; j < nbMachines; j++) {
    actOnMach[i][j] = IloIntervalVar(env,ptime[i][j]);
    actOnMach[i][j].setOptional();
  }
  m.add(IloAlternative(env,activity[i], actOnMach[i]));
}
IloIntervalVarArray2 machHasAct(env,nbMachines);
for (IloInt j=0; j < nbMachines; j++){
  machHasAct[j] = IloIntervalVarArray(env, nbActivities);
  for (IloInt i=0; i < nbActivities; i++)
    machHasAct[j][i] = actOnMach[i][j];
  m.add(IloNoOverlap(env,machHasAct[j]));
}
```

Note:

In this sample, the processing time of the activity depends on the machine. A more complex use-case is the one of alternative resource modes where a resource mode specifies a conjunctive set of resources to be used, this case is illustrated in the delivered example on Multi-Mode Resource Constrained Project Scheduling (`sched_rcpcpmm`). In some situations, it may be useful to add redundant constraints in the model so as to increase inference on alternative constraints (see Increasing inference on alternative constraints in the engine).

Another use of the alternative constraint is to state that an activity must execute within a set of alternative time windows. Consider a worker that can perform its activities in a set of at most n flexible shifts with specific constraints on the shift duration and minimal resting time between two consecutive shifts. Each shift of the worker can be modeled as an (optional) interval variable. The set of n shifts forms a chain of n optional interval variables (see "Chains of optional intervals" on page 97) and each activity is an alternative among a set of n optional activities, one for each possible shift. Each shift spans the set of all possible activities executing in this shift. This type of model must be used with care because it multiplies the number of interval variables. It is useful when the time-windows where the

activities must be placed are not fixed. If the time-windows are fixed, it is better to use an `IloForbidExtent` constraint when it makes sense (see "Modeling resource calendars" on page 96).

```
IloInt nbActivities = 100;
IloInt nbShifts = 5;
IloInt maxShiftDuration  = 600;
IloInt minInterShiftRest = 480;
IloIntervalVarArray activity(env,nbActivities);
IloIntervalVarArray2 actInShift(env,nbActivities);
for (IloInt i=0; i<nbActivities; i++) {
  activity[i] = IloIntervalVar(env,10);
  actInShift[i] = IloIntervalVarArray(env,nbShifts);
  for (IloInt s=0; s<nbShifts; s++) {
    actInShift[i][s] = IloIntervalVar(env);
    actInShift[i][s].setOptional();
  }
  m.add(IloAlternative(env,activity[i],actInShift[i]));
}
IloIntervalVarArray shifts(env,nbShifts);
IloIntervalVarArray2 shiftActs(env,nbShifts);
for (IloInt s=0; s<nbShifts; s++) {
  shifts[s] = IloIntervalVar(env,0,maxShiftDuration,IloTrue);
  shiftActs[s] = IloIntervalVarArray(env,nbActivities);
  for (IloInt i=0; i<nbActivities; i++)
    shiftActs[s][i] = actInShift[i][s];
  m.add(IloSpan(env, shift[s], shiftActs[s]));
  if (s>0) {
    m.add(IloEndBeforeStart(env, shift[s-1], shift[s], minInterShiftRest));
    m.add(IloIfThen(env, IloPresenceOf(env,shift[s]), IloPresenceOf(env,shift[s-1])));
  }
}
```

Modeling hierarchical models and "Work Breakdown Structures"

Hierarchical models and Work Breakdown Structures can be used to model task decompositions.

Hierarchical models in which a high-level task decomposes into a set of lower-level ones are common in scheduling. For instance, it may correspond to the Work Breakdown Structure of a project in project scheduling. Many variants are possible: the depth of the hierarchy may be constant or dependent on the tasks, some tasks in the hierarchy may be optional and there may be different way to decompose a given task.

Consider a task *T* in the hierarchy. In the simplest case, task *T* decomposes into a set of subtasks *ST1*,...,*STn*. This can be modeled using a span constraint between *T* and its set of subtasks *STs = {ST1,...,STn}*; task *T* will be constrained to start at the start time of the first subtask in the set *ST1*,...,*STn* and to end at the end of the last subtask in this set.

```
IloIntervalVar T(env);
IloIntervalVar STs(env,n);
for (i=0; i<n; i++)
   STs[i] = IloIntervalVar(env);
m.add(IloSpan(env,T, STs));
```

Suppose now that some subtasks of a task *T* in the hierarchy are non-compulsory. In other words, even if *T* is executed, these subtasks may be left unperformed, perhaps incurring an additional cost. Reciprocally, if a subtask *ST* of a task *T* is compulsory, then if task *T* is executed, subtask *ST* will have to be executed as well. Of course, *T* may be non-compulsory. This notion can be modeled by using optional interval variables in the hierarchy. If subtask *ST* is compulsory, we simply

add the constraint `IloIfThen(env,IloPresenceOf(env,T),IloPresenceOf(env,ST))`.
If the top-level task of the hierarchy is compulsory, it must not be defined as
optional but as a necessarily present interval variable.

```
IloBoolArray compulsory(env,n);  //compulsory populated
IloIntervalVar T(env);
T.setOptional();
IloIntervalVar STs(env,n);
for (i=0; i<n; i++) {
   STs[i] = IloIntervalVar(env);
   STs[i].setOptional();
   if (compulsory[i])
    m.add(IloIfThen(env,IloPresenceOf(env,T),IloPresenceOf(env,STs[i])));
}
m.add(IloSpan(env,T, STs);
```

In some applications, there are several alternative ways *T1,...,Tm* to decompose a
given task *T*. A given decomposition *Ti* is described by a set of subtasks
STi1,...,STij,...,STin that will be executed if decomposition *Ti* is selected. This can be
modeled by an alternative constraint between *T* and the set of alternatives *T1,...,Tm*
(each decomposition is modeled by an optional interval variable) and by
representing each decomposition *Ti* with a span constraint between *Ti* and its
subtasks *STi1,...,STij,...,STin*.

```
IloIntArray nSubTasks(env,m);
// populated with number of subtasks of each Ts[i]
IloIntervalVar T(env);
IloIntervalVarArray Ts(env,m);
for (IloInt i=0; i<m; i++) {
  Ts[i] = IloIntervalVar(env);
  Ts[i].setOptional();
  IloIntervalArray subTasks(env,nSubTasks[i]);
  for (j=0;j<nSubTasks[i]; j++) {
    subTasks[j] = IloIntervalVar(env);
    subTasks[j].setOptional();
  }
  m.add(IloSpan(env,Ts[i], subTasks));
}
m.add(IloAlternative(env,T,Ts));
```

The final case illustrates a more complex situation that mixes the elementary cases
above: it models a complete task hierarchy with non-compulsory subtasks and
alternative decompositions. Each possible task in the hierarchy is indexed by an
integer i in `0..n-1`. The possible decompositions of a task i are described by two
arrays containing the decomposed task and the set of subtasks induced by this
decomposition. A Boolean integer array `compulsory` tells whether or not a task i is
compulsory in the hierarchy. The number of possible decompositions of a task i is
represented by `taskDecompositions[i].getSize()`. This is used to detect leaf tasks
that cannot be decomposed further (`taskDecompositions[i].getSize()=0`). The
array `taskNbParents[]` counts the number of times a given task is contained in the
subtasks of a decomposition. This is used to detect top-level tasks in the hierarchy
(`taskNbParents[i]=0`). Each task and each possible decomposition is represented as
an optional interval variable. Each compulsory top level task is constrained to be
present. If a task i is not a leaf task, then task i is constrained to be the alternative
of all its possible decompositions and each decomposition of task i is constrained
to span the decomposition subtasks. Finally, each compulsory subtask of a
decomposition is constrained to be present if the decomposition is selected.

```
// Data defining the work-breakdown structure
IloInt n=...;
IloInt totalNbDecompositions=...;
IloBoolArray compulsory(env,n); // populated
IloIntArray decomposedTask(env,totalNbDecompositions); // populated
```

```
IloIntArray2 decompositionSubtasks(env,totalNbDecompositions); // populated

IloIntervalVarArray tasks(env,n);
IloIntArray taskNbParents(env,n);
IloIntervalVarArray2 taskDecompositions(env,n);
for (IloInt i=0; i<n; i++) {
  tasks[i] = IloIntervalVar(env);
  tasks[i].setOptional();
  taskNbParents[i] = 0;
  taskDecompositions[i] = IloIntervalVarArray(env);
}

for (IloInt d=0; d<totalNbDecompositions; d++) {
  IloIntervalVar dec(env);
  dec.setOptional();
  IloInt i = decomposedTask[d];
  taskDecompositions[i].add(dec);
  IloIntervalVarArray subtasks(env);
  for (IloInt s=0; s<decompositionSubtasks[d].getSize(); s++) {
    IloInt j = decompositionSubtasks[d][s];
    taskNbParents[j]++;
    subtasks.add(tasks[j]);
    if (compulsory[j])
      m.add(IloIfThen(env,IloPresenceOf(env,dec), IloPresenceOf(env,tasks[j])));
  }
  m.add(IloSpan(env,dec, subtasks));
}

for (IloInt i=0; i<n; i++) {
  if (taskNbParents[i]==0 && compulsory[i])
    m.add(IloPresenceOf(env, tasks[i]));
  if (taskDecompositions[i].getSize()>0) {
    m.add(IloAlternative(env,tasks[i], taskDecompositions[i]));
  }
}
```

Modeling classical finite capacity resources

Classical finite capacity resources such as disjunctive/unary resources,
non-renewable/discrete resources, and renewable/producible/consumable
resources can be modeled efficiently in CP Optimizer.

Disjunctive resources are resources that can only process one activity at a time and
thus enforce a total order over the set of activities processed on the resource. There
are two ways to model disjunctive resources in CP Optimizer: using a no overlap
constraint or a cumulative function with a maximal level of 1:

• If the disjunctive resource is associated with a sequence dependent setup time,
the model using a no overlap constraint should be chosen because the cumul
function does not allow the expression of transition distances.

```
IloInt n = ...; // nActivities
IloInt m = ...; // nTypes
IloTransitionDistance setups(env, m);
for (IloInt i=0; i<m; ++i)
  for (IloInt j=0; j<m; ++j)
    setups.setValue(i, j, IloAbs(i-j));
IloIntArray type(env,n); // populate
IloIntArray ptime(env,n); // populate
IloIntervalVarArray act(env,n);
for (i=0; i<n; i++)
  act[i] = IloIntervalVar(env,ptime[i]);
IloIntervalSequenceVar res(env,act,type);
m.add(IloNoOverlap(env,res,setups));
```

- If the disjunctive resource is associated with a minimal capacity profile (specifying that the resource should be in use over some time windows, then a cumulative Function should be used specifying a constraint `IloAlwaysIn(env,s,e,1,1)` over each time-window [s,e] where the resource must be in use.

```
IloInt n = ...;
IloIntArray ptime(env,n); // populate
IloInt nMustBeInUse = ...;
IloIntArray start(env,nMustBeInUse); // populate
IloIntArray end(env,nMustBeInUse); // populate
IloIntervalVarArray act(env,n);
IloCumulFunctionExpr res(env);
for (i=0; i<n; i++)  {
   act[i] = IloIntervalVar(env,ptime[i]);
   res += IloPulse(act,1);
}
m.add(res <= 1);
for (IloInt w=0;w<nMustBeInUse; w++)
  m.add(IloAlwaysIn(env,res, start[w], end[w], 1, 1);
```

- If the resource is unavailable over some time-windows, it is preferable to use a cumul function with a set of "always in" constraints, such as `IloAlwaysIn(env,s,e,0,0)`, over each time-window [s,e] where the resource is unavailable unless the number of unavailability time windows is small enough in which case, adding fake fixed interval variables in the no overlap could make sense if there are specific reasons to use the no overlap constraint.

Non-renewable resources are resources with a finite capacity Q such that the total resource usage by all activities executing at a time point t do not exceed Q. Each activity executing on the resource requires a given amount of resource from its start time to its end time. In CP Optimizer, non-renewable resources are modeled using a cumul function defined as a sum of elementary pulse functions. Constraints <=, >= or `IloAlwaysIn` are used to constrain the maximal or minimal amount of resource used over time.

```
IloInt nSteps = ...;
IloIntArray start(env,nSteps); // populate
IloIntArray end(env,nSteps); // populate
IloIntArray capmax(env,nSteps); // populate
IloInt n = ...;
IloIntArray ptime(env,n); // populate
IloIntArray qty(env,n); // populate
IloIntervalVarArray act(env,n);
IloCumulFunctionExpr res(env);
for (IloInt i=0; i<n; i++)  {
   act[i] = IloIntervalVar(env,ptime[i]);
   res += IloPulse(act,qty[i]);
}
 for (IloInt s=0; s<nSteps; s++)
  m.add(IloAlwaysIn(env,res, start[s], end[s], 0, capmax[s]));
```

Renewable resources are resources that can be produced and consumed by activities. In CP Optimizer a renewable resource can be modeled as a cumul function defined as a sum of elementary step functions or their opposite. Resource production can be modeled as a +IloStepAtStart (if the activity produces at its start time) or a +IloStepAtEnd function (if the activity produces at its end time). Resource consumption can be modeled as a -IloStepAtStart (if the activity consumes at its start time) or a -IloStepAtEnd function (if the activity consumes at its end time). Constraints <=, >= or IloAlwaysIn are used to constrain the maximal or minimal amount of resource over time. The following sample illustrates a renewable resource representing a tank with a maximal capacity and a safety level with activities producing or consuming material stored in the tank.

```
IloInt Horizon = ...;
IloInt Capacity = ...;
IloInt SafetyLevel = ...;
IloInt StartLevel = ...;
IloInt NProd = ...;
IloInt NCons = ...;
IloIntArray QProd(env,NProd); // populate
IloIntArray QCons(env,NCons); // populate
IloCumulFunctionExpr level(env);
level += IloStep(env,0,StartLevel);
IloIntervalVarArray AProd(env,NProd);
for (IloInt i=0; i<NProd; i++) {
  AProd[i]=IloIntervalArray(env,1);
  level += IloStepAtEnd(AProd[i], QProd[i])
}
for (i=0; i<NCons; i++) {
  ACons[i]=IloIntervalArray(env,1);
  level -= IloStepAtStart(ACons[i], QCons[i]);
}
m.add(IloAlwaysIn(env,level, 0, Horizon, SafetyLevel, Capacity));
```

Modeling classical scheduling costs

Classical scheduling costs such as makespan, earliness/tardiness, resource
allocation costs and activity execution costs can be modeled efficiently in CP
Optimizer.

Non-execution costs

An activity non-execution cost *K* is modeled by an expression *K * presenceOf(a)* if *a*
is the optional interval variable representing the activity.

Makespan

A makespan cost is modeled as the maximal value of the end of a set of interval
variables. In the case of an optional interval variable, the value of expression
IloEndOf(a) is 0 when interval a is absent.

```
dvar interval act[i in 1..n];
dexpr int makespan = max(i in 1..n) endOf(act[i]);
 minimize makespan;
subject to {      // ...    }

IloIntervalVarArray act(env,n);
IloIntExprArray end(env);
for (IloInt i=0; i<n; i++ ) {
  act[i] = IloIntervalVar(env);
  ends.add(IloEndOf(act[i]));
}
m.add(IloMinimize(env,IloMax(ends)));
```

Earliness/tardiness costs

An earliness/tardiness cost can be modeled by a set of piecewise linear functions *f*
that represent the cost *f(t)* of finishing (or starting) an activity at a date *t*. Integer
expressions IloStartEval and IloEndEval are used to evaluate the function on the
start or end point of an interval variable. An example of earliness/tardiness cost
where the cost of an activity is expressed as a V-shaped function evaluated at the
activity end time follows. This sample combines the earliness/tardiness cost with a
non-execution cost: in the example, activities are supposed to be optional and
leaving the activity unperformed incurs a cost.

```
pwlFunction etcost[i in 1..n] = piecewise{-earliW[i]->targetEnd[i];
tardiW[i]}(targetEnd[i],0);
dexpr float cost = sum(i in 1..n) endEval(act[i], etcost[i], nonExecCost[i]);

IloIntArray targetEnd(env,n); // populate
IloNumArray earliW(env,n); // populate
IloNumArray tardiW(env,n); // populate
IloNumArray nonExecCost(env,n); // populate
IloNumExpr cost(env);
IloIntervalArray act(env,n);
IlNumExpr cost(env);
for (IloInt i=0; i<n; i++ ) {
  act[i] = IloIntervalVar(env);
  act[i].setOptional();
  IloNumToNumSegmentFunction etcost = IloPiecewiseLinearFunction(env,
                                       IloNumArray(env,1, targetEnd[i]),
                                       IloNumArray(env,2,-earliW[i],tardiW)
                                       targetEnd[i],0);
  cost += IloEndEval(act[i],etcost,nonExecCost[i]);
}
m.add(IloMinimize(env,cost));
```

When the function is very simple such as a pure tardiness cost and the activity is not optional, it is slightly more efficient to use a `IloMax` expression rather than a piecewise linear function as illustrated.

```
IloIntArray dueDate(env,n); // populate
IloNumArray tardiW(env,n); // populate
IloIntervalVarArray act(env,n);
IloNumExpr cost(env);
for (IloInt i=0; i<n; i++ ) {
  act[i] = IloIntervalVar(env);
  cost += tardiw[i]*IloMax(0,IloEndOf(act[i])-dueDate[i]);
}
m.add(IloMinimize(env,cost));
```

Resource allocation costs

A resource or mode allocation cost specifying a cost K incurred by an activity executing on a particular resource or in a particular mode is modeled by an expression K * IloPresenceOf(a) if a is the optional interval variable representing the execution of the activity on the resource or in the specified mode (see "Different uses of the alternative constraint" on page 98). The following sample illustrates a simple resource allocation cost for an activity executing on a set of alternative resources.

```
IloInt nbMachines   = 5;
IloInt nbActivities = 10;
IloIntArray ptime(env,nbActivities); // populate
IloIntArray2 allocCost(env,nbActivities);
IloIntervalVar activity(env,nbActivities);
IloIntervalVarArray2 actonMach(env,nbActivities);
for (IloInt i=0; i<nbActivities;i++) {
  activity[i] = IloIntervalVar(env,ptime[i]);
  allcost[i] = IloIntArray(env,nbMachines); // ppulate
  actOnMach[i] = IloIntervalVarArray(env,nbMachines);
  for (IloInt j=0; j< nbMachines; j++) {
    actOnMach[i][j] = IloIntervalVar(env);
    actOnMach[i][j].setOptional();
  }
  m.add(IloAlternative(env,activity[i],actOnMach[i]));
}

IloIntervalVarArray2 resOnAct(env,nbMachines);
for (IloInt j=0; j<nbMachines; j++) {
  resHasAct[j] = IloIntervalVarArray(env,nbActivities);
```

```
      for (IloInt i=0; i<nbActivities;i++) {
        resHasAct[j][i] = actOnMach[i][j];
      }
      m.add(IloNoOverlap(env, resOnAct[j]));
    }

    IloIntExpr cost(env);
    for (i=0; i<n; i++ )
      for (IloInt j=0; j<m; j++)
        cost += allocCost[i][m]*IloPresenceOf(env,actOnMach[i][m]);
    m.add(IloMinimize(env,cost));
```

Sequence-dependent setup costs

A sequence-dependent setup cost on a machine is usually expressed as a cost
matrix M[ti][tj] that gives the setup cost for the machine to switch from an
activity of type ti to an activity of type tj. Such a setup cost can be modeled by
using expressions *typeOfNext* and *typeOfPrev* on the sequence variable representing
the machine as illustrated in the following sample. A complete example modeling a
total setup cost is available in the delivered example <Install_dir>/cpoptimizer/
examples/src/cpp/sched_tcost.cpp.

```
IloEnv env;
IloModel model(env);
IloInt m    = ...; // Number of types { 0, 1, ..., m-1 }
IloInt last = m;   // Type of next for the last activity on the machine
IloInt abs  = m+1; // Type of next for a non-executed activity on the machine

IloIntArray2 M(env, m);
for (IloInt ti=0; ti<m; ++ti) {
  M[ti]= IloIntArray(env, m+2);
  for (IloInt tj=0; tj<m; ++tj) {
    M[ti][tj] = ...; // Setup cost between types ti and tj
  }
  M[ti][last] = ...; // Cost if an activity of type ti is last
  M[ti][abs]  = ...; // Cost if an activity of type ti is not executed
}

IloInt n = ...;    // Number of activities on the machine
IloIntervalVarArray act(env, n); // Activities on the machine
IloIntArray        type(env, n); // Activity types

// ...

IloIntervalSequenceVar machine(env, act, type);
model.add(IloNoOverlap(env, machine));

IloIntExpr totalCost(env);
for (IloInt i=0; i<n; ++i)
  totalCost += M[type[i]][IloTypeOfNext(machine, act[i], last, abs)];
model.add(IloMinimize(env, totalCost));

IloCP cp(model);
cp.solve();
env.end();
```

Modeling sequence-dependent setup times

Setup times on disjunctive resources can be modeled by a no-overlap constraint
with a transition distance.

It is quite common that a certain minimal amount of time must elapse between the
execution of two successive operations on a resource (e.g. a machine), and, often,
this amount of time depends on the types of the two successive operations. This is

the notion of sequence-dependent setup time and can be captured in CP Optimizer by a *no-overlap* constraint with *transition distance*.

Operations on the machine are represented by interval variables a[i] and a sequence variable seq is created on these interval variables to model the sequence of operations on the machine. The *type* of the different operations (used to compute the setup time) are specified when building the sequence variable.

A transition distance is represented by a transition distance class and stored as a matrix of setup times. A no-overlap constraint must be posted with the sequence variable and the transition distance to state that interval variables of the sequence do not overlap and that the sequence-dependent setup time of the transition distance applies to the intervals of the sequence.

CP Optimizer distinguishes two kinds of behavior of the no-overlap constraint with respect to transition distances: (1) transition distance between immediate successors and (2) transition distances between all successors.

A transition distance between immediate successors is generally useful for modeling the duration of a setup activity to switch the state of the resource from one interval to the next interval in the sequence. Here, the state of the resource is represented by the type of the interval in the sequence. This is illustrated in Sample 1; the Boolean flag passed at the construction of the no-overlap constraint specifies if the transition distance must only be applied between **immediate** successors on the sequence variable. A complete example of transition distance between immediate successors is available in the delivered example `<Install_dir>/cpoptimizer/examples/src/cpp/sched_setup.cpp`. In some more complex cases, the setup activity will have to be explicitly modeled as an interval variable because, for instance, it requires some additional resource. In this case, you can use the expressions *typeOfNext* and *typeOfPrev* on the sequence variable to constrain the length of the setup activity as illustrated in Sample 2. See "Modeling classical scheduling costs" on page 103 for modeling sequence-dependent setup costs.

In some specific cases, the transition distance must be applied between **all** pairs of intervals succeeding each other on the sequence, no matter if there are other intervals in between. For example, consider a set of movies to be scheduled on a TV channel. If a movie of type *ti* is scheduled after a movie of type *tj* (no matter which other movies are shown in between), depending on the types *ti,tj*, one would like a minimal amount of time to elapse between the two occurrences to avoid showing movies of similar types too close to each other. Sample 3 illustrates such a use-case using a minimal delay `separationTime[ti]` between movies of type ti; the Boolean flag passed at the construction of the no-overlap constraint specifies that the transition distance must be applied between **all** successors on the sequence variable.

SAMPLE 1: Sequence-dependent setup time on immediate successors

```
IloEnv env;
IloModel model(env);
IloInt m = ...; // Number of types { 0, 1, ..., m-1 }

IloTransitionDistance setupTimes(env, m);
for (IloInt ti=0; ti<m; ++ti)
  for (IloInt tj=0; tj<m; ++tj)
    setupTimes.setValue(ti,tj, ...); // Setup time between types ti and tj

IloInt n = ...;    // Number of activities on the machine
```

```
IloIntervalVarArray act(env, n); // Activities on the machine
IloIntArray         type(env, n); // Activity types
// ...

IloIntervalSequenceVar machine(env, act, type);
// Transition distance applies between immediate successors
model.add(IloNoOverlap(env, machine, setupTimes, IloTrue));
```

SAMPLE 2: Sequence-dependent setup activities

```
IloEnv env;
IloModel model(env);
IloInt m = ...; // Number of types { 0, 1, ..., m-1 }

IloInt n = ...;      // Number of activities on the machine
IloIntervalVarArray act  (env, n); // Activities on the machine
IloIntervalVarArray setup(env, n); // Setup activities on the machine
IloIntervalVarArray cover(env, n); // Covering activities on the machine
IloIntArray         type (env, n); // Activity types

// ...
for (IloInt i=0; i<n; ++i) {
  act  [i] = ...;
  type [i] = ...;
  cover[i] = IloIntervalVar(env);
  setup[i] = IloIntervalVar(env);
  IloIntervalVarArray dec(env); dec.add(act[i]); dec.add(setup[i]);
  model.add(IloSpan(env, cover[i], dec));
  model.add(IloEndBeforeStart(env, act[i], setup[i]));
}

IloIntervalSequenceVar machine(env, cover, type);
model.add(IloNoOverlap(env, machine));  // Setup activities

IloIntArray2 setupDuration(env, m);
IloInt last = m;  // Type of next for the last activity on the machine
for (IloInt ti=0; ti<m; ++ti) {
  setupDuration[ti]= IloIntArray(env, m+1);
  for (IloInt tj=0; tj<m; ++tj)
    setupDuration[ti][tj] = ...; // Length of setup activity between types ti and tj
  setupDuration[ti][last] = 0;   // Length of last setup activity
 }
for (IloInt i=0; i<n; ++i)
 model.add(IloLengthOf(setup[i])==
                     setupDuration[type[i]][IloTypeOfNext(machine,cover[i],last)]);
```

SAMPLE 3: Sequence-dependent setup time on all successors

```
IloEnv env;
IloModel model(env);
IloInt m = ...; // Number of types { 0, 1, ..., m-1 }

IloTransitionDistance separationTimes(env, m);
for (IloInt ti=0; ti<m; ++ti)
  separationTimes.setValue(ti,ti, ...); // Separation time between two movies of type ti

IloInt n = ...;      // Number of movies
IloIntervalVarArray movie(env, n); // Movies
IloIntArray         type(env, n); // Types
// ...

// Transition distance applies between all successors
IloIntervalSequenceVar movieSequence(env, movie, type);
model.add(IloNoOverlap(env, movieSequence, separationTimes, IloFalse));
```

Modeling stochastic scheduling problems on disjunctive resources

Stochastic programming problems on disjunctive resources can be modeled using same sequence constraints.

Scheduling problems often involve some uncertainty in the data and require computing a schedule that is, in a certain sense, robust with respect to these uncertainties. Data can be uncertain for several potential reasons such as activities taking more or less time than originally estimated, resource availability that is subject to variability (machine breakdowns, manpower absenteeism) or different unexpected delays (materials arriving later than planned, weather conditions), etc.

When facing a scheduling problem with uncertainties two important questions should be answered first:

1. Can I solve the deterministic version of the problem? Modeling and solving a stochastic scheduling problem is more complex than modeling and solving the deterministic variant of the problem where all uncertainties have been ruled out. Being capable of efficiently solving the deterministic version of the problem is always a *prerequisite*.

2. What is a solution to the robust scheduling problem? This question is often harder than it seems. The traditional solution to a scheduling problem, in the form of fixed start and end dates for activities is in general too rigid and too brittle to be considered as robust. One generally looks for a more flexible solution that implicitly represents a *set* of fully fixed schedules. The decision of which fixed schedule to use, that is, the selection of actual start and end dates for activities, is usually performed at schedule execution time. Stated otherwise, a solution to a robust scheduling problem is in general a particular *execution policy* computed off-line that will be used for on-line execution.

When the scheduling problem is dominated by disjunctive resources, like machines, tools or workers that can only perform one activity at a time, a natural representation of a flexible schedule is the sequence of activities on the resource. In this context, finding a robust schedule boils down to finding a robust allocation and sequencing of activities on resources.

The objective function of stochastic scheduling problems is often related with the expectation of a certain performance criterion. Typical examples include minimizing the expected project makespan or project cost or maximizing the probability that a given level of service can be satisfied.

Most of the existing approaches for stochastic scheduling implement a scenario-based framework, in which the uncertainty is modeled through the use of a number of n scenarios. Each scenario is a deterministic sub-problem that represents a particular realization of the uncertainties. In practical applications, these scenarios may be given by a selection of representative past execution of the system or they may be computed by sampling the probability distributions of the random variables related with uncertainties.

A global model is built by aggregating these n scenarios. Some constraints are added to ensure that the decisions taken off-line in order to build a robust schedule (for instance allocation and sequencing decisions) are consistent across all scenarios. The global objective function is a combination of the individual objective function of each scenario.

CP Optimizer provides a specific constraint `sameSequence` to state that two sequence variables should be ordered in the same way. This constraint is particularly useful in stochastic scheduling problems involving the creation of robust sequences on disjunctive resources. It will be used to state that, for a given disjunctive resource, the same sequencing of activities should be used across all n scenarios. This is illustrated in Sample 1.

SAMPLE 1: Scenario-based model for a single disjunctive resource

```
IloEnv env;
IloModel model(env);
IloInt n = ...; // Number of scenarios
IloInt m = ...; // Number of activities on the machine
IloIntArray2 dur = ...; // dur[i][j]: duration of activity j in scenario i

IloIntervalVarArray2 act(env, n);          // act[i][j]: activity j in scenario i
IloIntervalSequenceVarArray seq(env, n); // seq[i]: sequence of activities for scenario i
for (IloInt i=0; i<n; ++i) {
   // Scenario i
   act[i] = IloIntervalVarArray(env, m);
   for (IloInt j=0; j<m; ++j) {
     act[i][j] = IloIntervalVar(env, dur[i][j]);
   }
   seq[i] = IloIntervalSequenceVar(env, act[i]);
   model.add(IloNoOverlap(env, seq[i]));
   if (0<i) {
     // Same sequence across all scenarios
     model.add(IloSameSequence(env, seq[0], seq[i]));
   }
}
```

A complete example of stochastic scheduling problem is available in the delivered example `<Install_dir>/cpoptimizer/examples/src/cpp/ sched_stochastic_jobshop.cpp`. This problem is a stochastic version of the classical job-shop scheduling problem where the duration of activities is uncertain and the objective is to minimize the expected makespan.

Increasing inference on alternative constraints in the engine

Adding redundant cumul functions can increase the inference of the optimizer on alternative constraints.

There may be situations where stronger inference on alternative constraints will help the optimizer to find a solution or to converge quicker. Given a pool of m alternative resources, some activities `act[i]`, i in 0..n−1 in the schedule need to select one resource from this pool. As described in the section "Different uses of the alternative constraint" on page 98, this can be modeled by a set of m optional interval variables `actOnRes[i][k]` for each possible resource k in 0..m−1 and an alternative constraint between `act[i]` and these `actOnRes[i][k]`. The resource pool can also be seen, globally, as a renewable resource of capacity m, each activity `act[i]` requiring one unit of this resource. This additional redundant constraint will provide the engine a more global view on the number of resources simultaneously used at each time point, independently from the actual resource that is being allocated to each activity. This redundant constraint can be modeled as a cumul function with maximal level as illustrated on the following sample.

Note:

Sometimes, although the description of the problem in natural language mentions the allocation of individual resources to an activity, it is not necessary to use a fine

grained model and the decomposition of the activities as an alternative of optional interval variable on each resource will not be necessary. A cumul function with maximal level will be sufficient for these cases where resources from a resource pool do not have individual characteristics. When possible, this will result in a much lighter and efficient model.

```
IloIntervalVarArray act(env,n);
IloIntervalVarArray2 actOnRes(env,n);
for (IloInt i=0; i<n; i++ ) {
  act[i] = IloIntervalVar(env,pt[i]);
  actOnRes[i] = IloIntervalVarArray(env,m);
  for (IloInt j=0; j<m; j++) {
    actOnRes[i][j] = IloIntervalVar(env);
    actOnRes.setOptional();
  }
  m.add(IloAlternative(env,act[i], actOnRes[i]));
}

IloIntervalVarArray2 ResHasAct(env,m);
for (IloInt j=0; j<m; j++) {
  resHasAct[j] = IloIntervalVarArray(env,n);
  for (IloInt i=0; i<n; i++ )
    resHasAct[j][i] = actOnRes[i][j];
  m.add(IloNoOverlap(env,resHasAct[j]));
}

cumulFunction nbUsed(env);
nbused += IloPulse(act[i], 1);

// Redundant constraint
m.add(nbUsed <= m);
```

Chapter 10. Debugging applications

CP Optimizer offers debugging features including exception handling, logging and tracing.

Overview

CP Optimizer offers debugging features including exception handling, logging and conflict refining.

When you are debugging an application that employs CP Optimizer, this library offers several features, such as logging and conflict refining, to make the task easier in whatever development environment you prefer. "The search log" on page 65 introduces the search log of CP Optimizer. In the search log, you can see the effectiveness of propagation, the progress of the search and the status of a solution. This section documents other features of CP Optimizer that facilitate debugging.

Catching exceptions

Exception handling for managing anomalies is a good programming practice when using CP Optimizer.

When programming an application, it is a good programming practice to enclose parts of your application in try-catch statements. In that way, when anomalies arise during execution, they can be managed as exceptions, so that your application can recover as cleanly as possible. You will find samples of try-catch statements in the examples distributed with CP Optimizer.

If you use try-catch statements, it is possible to distinguish exceptions raised by CP Optimizer from exceptions raised by the rest of the application. Specifically, when an error condition is encountered, CP Optimizer raises an exception of type `IloException.` Your application can catch these exceptions within try-catch statements, and you can thus determine directly whether the anomaly arises within the CP Optimizer part of your application or in another part of your application. Here's the conventional way to catch an error exception from CP Optimizer in the C++ API:

```
catch (IloException& e) {
  ...
  e.end();
}
```

The reference manual documents exceptions specific to IBM ILOG Concert Technology and CP Optimizer.

Note:

Catch exceptions by reference

Catch exceptions by reference, not by value, to avoid losing information and to prevent leaks from expressions or arrays.

In the C++ API, exception classes are not handle classes. Thus, the correct type of an exception is lost if it is caught by value rather than by reference (that is, using

`catch(IloException& e) {...}`). This is one reason that catching `IloException` objects by reference is a good idea, as demonstrated in all the examples distributed with CP Optimizer. Some derived exceptions may carry information that would be lost if caught by value. So if you output an exception caught by reference, you may get a more precise message than when outputting the same exception caught by value.

There is a second reason for catching exceptions by reference. Some exceptions contain arrays to communicate the reason for the failure to the calling function. If this information were lost by calling the exception by value, the method `end` could not be called for such arrays, and their memory would be leaked (until `env.end` is called). After catching an exception by reference, calling the `end` method of the exception will free all the memory that may be used by arrays (or expressions) of the actual exception that was thrown.

You can also control where warnings and error messages of the CP Optimizer part of your application are displayed. For example, during debugging, you might want all warning or error messages directed to your monitor, whereas when your application goes into production for use by customers, for example, you might want to direct warnings and error messages to a log file or some other discreet channel.

In the C++ API of Concert Technology, the class `IloEnv` initializes output streams for general information, for error messages and for warnings. The class `IloAlgorithm` supports these communication streams and the class `IloCP` inherits its methods. For general output, there is the method `IloAlgorithm::out`. For warnings and nonfatal conditions, there is the method `IloAlgorithm::warning`. For errors, there is the method `IloAlgorithm::error`.

In the C++ API, an instance of `IloEnv` defines the output stream referenced by the method out as the system `cout` in the C++ API, but you can use the method `setOut` to redefine it as you prefer. For example, to suppress output to the screen in a C++ application, use this method with this argument:

```
setOut(IloEnv::getNullStream())
```

Likewise, you can use the methods `IloAlgorithm::setWarning` and `setError` to redefine those channels as you prefer.

In the Microsoft .NET Framework languages and Java APIs, the native streams are used directly.

Using the conflict refiner

The conflict refiner is used to find an explanation for the infeasibility of a model.

In the process of developing a model it often happens that, by nature of the model or because some errors slipped in the model or in the data, the problem is infeasible. In this case, you can use the conflict refiner functionality to identify a minimal infeasible subset of constraints.

Consider for instance the following map-coloring problem in which there are two available colors (blue, white) and six contiguous countries (Belgium, Denmark, France, Germany, Luxembourg, Netherlands).

```
#include <ilcp/cp.h>

const char* Names[] = {"blue", "white"};
```

```
int main(){
  IloEnv env;
  try {
    IloModel model(env);
    IloIntVar Belgium    (env, 0, 1, "Belgium    "),
              Denmark    (env, 0, 1, "Denmark    "),
              France     (env, 0, 1, "France     "),
              Germany    (env, 0, 1, "Germany    "),
              Netherlands(env, 0, 1, "Netherlands"),
              Luxembourg (env, 0, 1, "Luxembourg ");
    // Constraints
    model.add(Belgium    != France);
    model.add(France     != Luxembourg);
    model.add(Luxembourg != Belgium);
    model.add(Belgium    != Netherlands);
    model.add(Luxembourg != Germany);
    model.add(Belgium    != Germany);
    model.add(France     != Germany);
    model.add(Germany    != Netherlands);
    model.add(Denmark    != Germany);
    IloCP cp(model);
    // Search for a solution
    if (cp.solve()) {
      cp.out() << cp.getStatus() << " Solution" << std::endl;
      cp.out() << "Belgium    :" << Names[(IloInt)cp.getValue(Belgium)] << std::endl;
      cp.out() << "Denmark    :" << Names[(IloInt)cp.getValue(Denmark)] << std::endl;
      cp.out() << "France     :" << Names[(IloInt)cp.getValue(France)] << std::endl;
      cp.out() << "Germany    :" << Names[(IloInt)cp.getValue(Germany)] << std::endl;
      cp.out() << "Netherlands:" << Names[(IloInt)cp.getValue(Netherlands)] << std::endl;
      cp.out() << "Luxembourg :" << Names[(IloInt)cp.getValue(Luxembourg)]  << std::endl;
    } else {
      // If the problem was shown to be infeasible,
  //    find a minimal explanation for infeasibility
      if (cp.refineConflict()) {
        cp.writeConflict(cp.out());
      }
    }
  } catch (IloException& ex) {
    env.error() << "Error: " << ex << std::endl;
  }
  env.end();
  return 0;
}
```

This problem is infeasible. Running the conflict refiner on this model provides the following explanation:

```
 // ------ Conflict members: -------------------------------------------------
France != Luxembourg;
Luxembourg != Germany;
France != Germany;
```

The explanation is that it is not possible to color Germany, France and Luxembourg with only two colors. At least three different colors are required as these countries are contiguous. This explanation is minimal in the sense that if one removes any of these three constraints, the two remaining ones are feasible.

For more information on the conflict refiner, please refer to the concept "Conflict Refiner in CP Optimizer" in the CP Optimizer C++ API Reference Manual.

Another example of the use of the conflict refiner for a scheduling problem is available in the delivered example <Install_dir>/cpoptimizer/examples/src/cpp/ sched_conflict.cpp.

In this example you will learn how to use constraint weights to select a preferred explanation and how the conflict refiner can be used to iterate over several explanations.

Testing with a known solution

Using an instance of the problem with a known solution is a useful method for debugging.

Chapter 8, "Designing models," on page 89 explains how to design good models for various problems. However, if errors slip in when you implement the constraints of the problem, it can be very difficult to understand why the application finds wrong answers, even though the model is very good.

There is a simple way to address such cases: use an instance of the problem with a known solution to test the constraints. For testing a solution, you assign the values corresponding to a solution before adding the constraints to the model. Then no failure should happen within this function. If a failure occurs with the known solution, you know there is a problem among the constraints.

To see how to test with a known solution, consider a map-coloring problem in which there are four available colors (blue, white, red, green) and six contiguous countries (Belgium, Denmark, France, Germany, Luxembourg, Netherlands). One known solution is to color Belgium red; Denmark, France and the Netherlands blue; Germany green; and Luxembourg white.

This C++ sample tests a known solution of the map-coloring problem.

```
#include <ilcp/cp.h>
const char* Names[] = {"blue", "white", "red", "green"};
int main(){
  IloEnv env;
  try {
    IloModel model(env);
    IloIntVar Belgium(env, 0, 3), Denmark(env, 0, 3),
      France(env, 0, 3), Germany(env, 0, 3),
      Netherlands(env, 0, 3), Luxembourg(env, 0, 3);
    // Test a known solution
    model.add(Belgium == 2);
    model.add(Denmark == 0);
    model.add(France == 0);
    model.add(Germany == 3);
    model.add(Netherlands == 0);
    model.add(Luxembourg == 1);
    // Constraints
    model.add(Belgium != France);
    model.add(Denmark != Germany );
    model.add(France != Germany);
    model.add(Belgium != Netherlands);
    model.add(Germany != Netherlands);
    model.add(France != Luxembourg);
    model.add(Luxembourg != Germany);
    model.add(Luxembourg != Belgium);
    model.add(Belgium != Germany);
    IloCP cp(model);
    // Search for a solution
    if (cp.solve()) {
      cp.out() << cp.getStatus() << " Solution" << std::endl;
      cp.out() << "Belgium:" << Names[(IloInt)cp.getValue(Belgium)] << std::endl;
      cp.out() << "Denmark:" << Names[(IloInt)cp.getValue(Denmark)] << std::endl;
      cp.out()  << "France:" << Names[(IloInt)cp.getValue(France)] << std::endl;
      cp.out() << "Germany:" << Names[(IloInt)cp.getValue(Germany)] << std::endl;
```

```
        cp.out() << "Netherlands:" << Names[(IloInt)cp.getValue(Netherlands)] << std::endl;
        cp.out() << "Luxembourg:" << Names[(IloInt)cp.getValue(Luxembourg)]  << std::endl;
                      }
  }
  catch (IloException& ex) {
    env.error() << "Error: " << ex << std::endl;
  }
  env.end();
  return 0;
}
```

If the solution turns out not to be feasible for the constraints of the model, you can use the conflict refiner to identify which constraints are violated in the solution as described in "Using the conflict refiner" on page 112.

Chapter 11. Developing an application with CP Optimizer

Developing an application with CP Optimizer involves data preparation, modeling and solving.

Overview

Developing an application with CP Optimizer involves data preparation, modeling and solving.

This section offers a few guidelines for creating an application that exploits CP Optimizer. Recognizing that the translation from a specification to an application can be challenging, this section covers topics regarding data preparation, modeling and solving.

Preparing data

Preparing data for use in a CP Optimizer application involves ensuring that you have realistic, representative data and that the format is specified.

Start with clean, **realistic data**. If you don't have access to real data, you should consider fabricating realistic data for this purpose.

For example, imagine that you are developing a rostering application for nurses in a hospital, where the roster covers six months for 30 nurses with different levels of skill; six nurses are highly qualified, 20 have standard qualifications, four are beginners. The data set does not need to represent the nurses individually in detail, but it needs to satisfy the number of nurses per day for each service and the level of qualification of the nurses. Realistic data for this rostering application must involve the same **proportion** of qualified nurses, the same **type** of service requests and so on.

Realistic data must also be **representative** even when you are testing a reduced data set on a smaller version of the problem. This principle means that some data in a smaller version of a problem can simply be reduced, but other data must be reduced only in ways that respect the proportions of the original problem because changing the proportions among those data would effectively change the problem to solve. To understand this difference between data that can be reduced arbitrarily and representative data that must respect proportions when it is reduced, consider the nurse rostering example again. One way to reduce the size of the problem is simply to consider a shorter period of time, for example, one month instead of six months. In other words, from a constraint programming point of view, the period of time can be reduced almost arbitrarily. In contrast, if you reduce the number of nurses in your test data in order to work with a smaller problem, your test data must still respect the **proportions** among their levels of skill. For example, if you decide to test your application on half the number of nurses (15 instead of 30), then a **representative data set** must still include three highly qualified, ten with standard qualifications and two beginners in order to respect the proportions of the original problem.

The solution of a combinatoric problem is quite sensitive to variations in data, so you need to run, test and optimize an CP Optimizer application with respect to **multiple sets of data** to have a reliable effect. In fact, the robustness of your

application will depend heavily on tests run over several sets of data. This point about using multiple sets of data to test your program is so important that if the client for your application cannot supply multiple sets of real data, then you should consider generating multiple sets of realistic data, for example, by random variation.

Early in development, you should also settle the **format of data**. If, for example, it is straightforward and quick to sort an array by posing a few constraints, it will be even quicker to use a conventional sorting technique instead. This guideline can be generalized: most ordinary preprocessing of data (unrelated to constraint programming) can be handled more efficiently in your chosen programming language rather than in CP Optimizer.

Use multiple sets of clean, valid, realistic data to validate the model that you design. After you have validated it, your first model itself will play the role of a reference. It will enable you to test new solutions that you get from the implementations you develop.

Later, multiple sets of data may also help you tune performance, as variations between data sets can highlight different aspects of your application that may allow improvement.

Clarifying the cost function

A good cost function accurately represents which solutions are preferable.

Overview

A good cost function accurately represents which solutions are preferable.

Clarifying the cost function may be a truly difficult step in developing your application if the cost function provided to you is not really representative of a good solution for the client.

The client may be mainly interested in a good solution, but as a programmer, you need a good cost function accurately representing what a good solution is. Even if your client supplies a cost function as part of the specification, you may need to look more closely at it or modify it in some respect.

Dissect the cost function

It is important to separate distinct types of costs so that each aspect of the cost function can be appropriately studied while the model is being built.

It is important to have a clear definition of the cost function (or objective) to optimize. Often this function will contain diverse components or terms to optimize, each relating to a different part of the business.

Only part of this cost function may be a measure of monetary profit or cost. Other parts may be related to subjective criteria or perceived inconvenience to the business. For example, some parts of the cost function may depend on work practices, labor conventions, customary habits of employees. Parts of the cost function corresponding to these optimizations can often change quite dramatically during the course of a project.

It is important to separate distinct types of costs so that each aspect of the cost function can be appropriately studied while the model is being built. Such a

separation will also help you represent the problem more accurately in your model. A good representation of the cost function will also help you communicate effectively with your client.

Approximate the cost function

In some cases, it may be best to approximate the cost function.

The client for your application may supply, as part of the specification, a cost function that takes into account various complex measures of solution quality from a business point of view. Since it is familiar and business-like, your client may expect to see results in terms of this cost function, as he or she already understands this measure.

However, cost functions coming from clients may not necessarily be the best way to get good solutions to an optimization problem. In such situations, a *proxy* for the cost function supplied by the client may be a better way to proceed.

The idea behind this advice is that the proxy approximates the cost function supplied by the client. The proxy should correspond closely to the client-supplied cost function, but it should be significantly more amenable to optimization with constraint programming. CP Optimizer will optimize the proxy to produce a good solution. However, the client need never see this proxy, as you can write code in your application to calculate his or her preferred cost function from the solution produced by the proxy.

As an example of a good place to use a proxy instead of a client-supplied cost function, consider the nursing roster again. In order to limit the number of times that a given nurse must work at night, it is conventional in some rostering applications to use a Big-M formulation, where a counting variable with a very large maximal value is introduced as part of the cost function. A better approach is to add a constraint that limits the maximum value of this counting variable, instead of adding this parameter in the cost function.

Defining a solution

A solution satisfies the constraints of the problem.

In CP Optimizer, a solution can never violate a constraint. By definition, a solution satisfies the constraints of the problem, and any problem in which a constraint is violated has no solution in constraint programming terms.

In contrast, when a client specifies the requirements of a problem, the client may have loose or vague ideas about ways in which an acceptable solution might violate a constraint *just a little bit*; or a solution might ignore certain restrictions *sometimes*; or a solution could overlook a requirement *if necessary*.

To cope with this difference between how constraint programming construes a solution and what your client expects as a solution, you must first understand what the client regards as a solution. One approach to this task is to ask your client for a sample solution.

If the known solution fails in your application, then you know that your model does not correspond to the client's idea of an acceptable solution. Your next task, then, is to agree with the client about what really constitutes a solution to the problem. A first step toward this agreement is to identify which constraints can

never be violated in the client's mind. For example, security regulations might be such: no solution acceptable to the client should violate a security regulation. In such a problem, security regulations will consequently be expressed as constraints.

Identifying the constraints

It is important to identify all of the constraints, to understand which ones are real ones and which have been used in the past for tuning, and to know which ones may be relaxed.

Overview

It is important to identify all of the constraints, to understand which ones are real ones and which have been used in the past for tuning, and to know which ones may be relaxed.

When you consider the aspects of a specification that involve constraints, there are a number of important points to keep in mind.

- You need a complete and exhaustive statement of all the constraints, as a minimal part of the problem description. "Define all constraints" offers an example of this point.
- You need a way to identify those conditions that your client would like to satisfy but that you may relax or drop altogether if the problem proves intractable. Your client helps you identify those conditions. "Distinguish constraints from preferable conditions" on page 121 suggests ways to do so.
- You must clearly separate the real constraints from any ad hoc rules of thumb used to solve the problem in the past. Make the real constraints part of the model, but consider whether the ad hoc rules of thumb are still necessary. Perhaps they can be replaced by constraint programming techniques.

Define all constraints

A model must contain all of the constraints.

Technically satisfying the constraints in a specification may not always meet the client's expectation.

For example, in the design of a telecommunication network, the problem description might demand that every pair of nodes in the network must have at least two paths between those nodes for robustness and reliability.

To meet that demand, one might design a model, taking into account that constraint on every pair of nodes. In a first-cut of that model, CP Optimizer could find a solution that is quite logical: the solution consists of a huge loop, passing through every node. Since a such a loop can be traversed in both directions, this first-cut solution technically contains two paths to between every pair of nodes, though the paths between some pairs may be rather long. When the solution is shown to the client, however, the designer learns, alas, that solution is not acceptable in practice, so the designer goes back to the phase of describing the problem again and adds constraints on the length of paths between pairs of nodes. In consequence of this change in the problem *description*, the *model* and the *implementation* must change as well.

A better problem description in the first place (one that included all the constraints, especially those on the lengths of paths between pairs of nodes) might

have avoided this repetition. Certainly a sound initial problem description is the surest safeguard against the infamous encroachments of an incremental specification.

Distinguish constraints from preferable conditions

A model must distinguish constraints from preferable conditions.

Sometimes a specification includes conditions that the client would like to satisfy, but that might be relaxed or dropped if the problem proves intractable. These are the conditions that might lead to the elimination of potential solutions if the application insists on imposing them. At the same time, if the application fails to take these conditions into account, the application will come up with a "solution" that is not satisfactory in reality because these conditions embody some fundamental and practical aspect of the problem.

Separating these conditions from the hard constraints of the problem accomplishes two tasks: to qualify acceptable solutions; and to measure the quality of proposed solutions.

This work (separating constraints from other conditions so that the adequacy of various solutions can be compared) is often very difficult, especially when the conditions represent social rules, accumulated and elaborated over time, such social rules as customary working hours, conventional combinations of tasks, habitual job assignments and so forth. If the model ignores such conditions, the application is likely to find unacceptable solutions; yet if the application enforces such conditions as constraints, there may be no solution at all. In any case, the nature of such conditions has an impact on the model and implementation since they influence the cost function.

These kinds of conditions (social rules, customs, conventions, habits) depend strongly on the type of problem you are solving. Consequently, there are few or no general guidelines that always apply. As you consider how to manage such conditions in your application, keep in mind the fact that CP Optimizer performs well when constraints are tight and accurately represent conditions that cannot be violated. (Tight constraints are those that limit the possible combinations of values of variables.) Consequently, an approach that consists of adding quantities of terms in the objective function is generally not a good idea.

There are, however, various effective ways of managing these conditions in an application of CP Optimizer.

- In some applications, it is appropriate to represent these conditions as *slack* which can be pushed into the cost function.
- In some models, it is practical to start with all possible constraints, including these conditions. If no solution can be found, then remove the conditions one by one, until a solution is possible.
- In other models, it is more practical to start with a minimal set of hard constraints. Find a solution to this minimal set of constraints. Then add conditions as constraints, one by one, until failure occurs. Remove that last condition, the constraint that led to failure.

Abstracting a miniproblem

A representative miniproblem I useful for validation and testing.

Before you actually embark on the total application with a final model, it's a good idea to experiment with small-scale models. In that way, you'll acquire useful knowledge about the problem, and you'll get this information early enough in development to be useful later.

In the model-prototype, the first decision to make is how to identify a miniproblem from the overall problem. The miniproblem provides the basis for a first-cut model of the whole problem. In practice, it's important to get a first-draft model as early as possible. That model serves two purposes: as a validation field for the specifications and as an experimental domain for the feasibility of the project.

It's thus critical to choose a **representative** miniproblem along the same lines as the choice of representative data discussed in "Preparing data" on page 117. If the miniproblem that you identify is too simple, it won't show you how the application will perform. If it's too complicated, it will demand too much work and very possibly wasted effort if the model reveals a flaw in the design. In other words, the choice of a miniproblem has serious consequences. Your general experience and your knowledge of similar applications will greatly facilitate your choice. One fallback choice (sometimes the only choice) is to use a small-scale instance of the entire problem.

This small-scale instance of the entire problem can help you identify any difficult subproblems. After you identify a difficult subproblem, you can focus on **resolution** of this subproblem and on **communication** between this subproblem and the rest of the problem. In this context, communication between a subproblem and the rest of a greater problem can be expressed through constraints between the variables in the subproblem and variables in the rest of the problem as a whole.

Designing the model

As you design a model for your application, there are several principles to guide you.

Overview

As you design a model for your application, there are several principles to guide you.

There are practical ways to make designing your CP Optimizer application easier, as suggested in the following sections.

Decompose the model

Decomposing the model can lead to faster solution times.

If the model can be logically decomposed, you should do so. You can decompose the model in the same way that you partition the variables into subsets with few constraints linking the variables of different subsets. It is really important to find such a decomposition and to carry it out because the decomposition may save an exponential factor in the solution time.

In fact, if you identify subproblems that are independent of one another, CP Optimizer will take into account their independence. For example, consider a problem consisting of subproblems A, B and C, where B and C are independent of each other. If you inform CP Optimizer about the independence of subproblem B

from subproblem C, and in addition, CP Optimizer finds a solution of subproblem A that renders subproblem C unsatisfiable, then CP Optimizer will not waste time attempting to solve subproblem B.

Determine precision

Determining the degree of precision that represents a reasonable trade-off between numeric accuracy in the solution and acceptable performance in your application is important.

Clearly, the degree of precision (that is, the granularity of the definition and solution) needs to be adapted to the aims of the application. There is no point in counting seconds if your problem works in terms of years. In short, choose an appropriate unit, both for the problem representation and for the solution.

High precision may mean slower performance. You may need to experiment to determine the degree of precision that represents a reasonable trade-off between numeric accuracy in the solution and acceptable performance in your application.

Validate the model with data and constraints

Validating the model with data and constraints helps you to check the constraints.

As mentioned in "Preparing data" on page 117, you need to validate your model with clean, realistic data. In fact, it is a good idea to validate your model with multiple sets of data, as combinatoric problems are notoriously sensitive to slight variations in data.

Validation is greatly simplified if a solution is already available to you. With a known solution, you can check the way that your constraints have been written so that you will be alerted to any errors in interpretation or any bug in translation between the specification and the implementation. You can also check the global coherence of your initial constraints, their coherence as a set. That is, when you have a known solution already, you can use it as a reference point in testing and verification.

Experiment with variant models

experimenting with variant models may help you in finding a good fit.

In spite of these common sense rules, there are many ways to design the model and create the prototype for any given problem. You'll need to try more than one model to find the best fit for your problem, but don't think of these trials as wasted effort. The time that you spend in designing the model, even time spent "playing around" with it, is time well spent since that time will be saved later.

"Consider symmetry" on page 125 offers an example where experimentation is a good idea.

Tuning Performance

Tuning the performance of the optimizer is useful. There are practical ways to make your CP Optimizer application more efficient, as suggested in the following sections.

Use multiple data sets for tuning

Using multiple data sets for tuning is important.

As suggested in "Preparing data" on page 117, you will need multiple data sets. Not only are these multiple data sets useful for validation of your model, but they are also practical for tuning aspects of your application. Different data sets highlight different aspects of performance.

Optimize propagation and search

CP Optimizer offers features for optimizing propagation and search.

With respect to CP Optimizer, the areas to optimize in your application are the constraint propagation and the solution search. Improvements in those two areas are interdependent, and the greatest gains in efficiency come about when you work simultaneously on both aspects.

To optimize propagation and search, CP Optimizer provides several features for you to consider:
• Inference levels, introduced in Chapter 5, "Constraint propagation in CP Optimizer," on page 45
• Search phases, documented in "Ordering variables and values" on page 79

Look at changes

The search and propagation log may point out weaknesses in propagation and search.

To help you see what you're doing, it's a good idea to look hard at the search log and at the trace.

"The search log" on page 65 introduces this log and explains how to read it.

The trace lets you rapidly distinguish which variables to take into account early in the solution search. It also shows you the areas where propagation is weakest, areas that you should re-enforce with surrogate constraints.

Use expertise from the problem domain

Using expertise from the problem domain may help you tune the search.

At some point in developing your application with CP Optimizer, you'll need to think about the criteria for choosing variables and values. These choices are so specific to the problem under consideration, sometimes even specific to the end-user, that there are no "canned" rules ready to use "off-the-shelf." You really must exploit professional expertise from the problem domain.

Professional knowledge about the problem itself proves useful in tuning the search. In real-world problems (unlike purely theoretical problems) the hardest parts are actually often confined to one sole part of the model. In consequence, you have to start searching for a partial solution in that region in order to minimize the exponential factor.

For example, when you're building a schedule, you have to assign the most difficult or the most critical items first. The easier parts fall into place around them. A domain expert helps you identify the difficult or critical items of a scheduling application.

This kind of solution clearly shows that the procedure for solving a problem depends very much on the instance of the problem under consideration, not strictly on its formal statement.

Reconsider what makes a solution optimal

Examining the tolerance of the objective function is important in tuning the performance.

CP Optimizer treats optimization problems by enforcing the idea that the current solution must be better than the preceding solution at each iteration in terms of a cost. Parameters, such as `OptimalityTolerance` and `RelativeOptimalityTolerance`, enable you to control how much better the following solution must be.

The essential factor is in the propagation of the cost variable. At every moment, its boundaries must be the best evaluation possible of the cost of the solution partially examined.

"Optimality at all costs" is frequently unrealistic or even impossible. The behavior that you, as a developer, want from an optimization application is to achieve significantly better solutions very fast. With such an aim, it may be relevant to impose a rule such as, "the next solution must be 10% better than the current one." Or, for example, it may be a good idea to eliminate once and for all the solutions close to the one already obtained by removing the values near the value belonging to the most recent solution.

Consider symmetry

Considering whether removing symmetry helps in the search for a solution is important in tuning the performance.

Very often, real-world problems have symmetries. Consequently, a direct model of the real-world problem will also have symmetries. Examples of such symmetries are a fleet of identical trucks, a set of identical containers, engineers with identical skills and so on.

Eliminating symmetries in your models by adding constraints can be beneficial, especially if the aim is to prove the optimality or non-existence of a solution. For example, given a fleet of identical trucks or containers, it is conventional to impose an arbitrary order among them simply by numbering them. In that way, they are no longer technically identical, and symmetries arising from their original interchangeability are thus eliminated.

However, eliminating symmetries does not always accelerate the speed at which good solutions are found with CP Optimizer. There are problems in which eliminating symmetry can force the search to look in a less productive area of the solution space than it might have pursued had all the possible paths remained available.

In other words, while it may be a good idea to eliminate symmetry in some problems, it is also necessary to check whether symmetry serves a useful purpose in others. Consequently, you may need to experiment, to test whether eliminating symmetry is a good idea in your model. The best way to do so is to group all symmetry constraints in one easily identifiable place in your model so that each one can be activated or deactivated as you carry out these experiments.

Notices

This information was developed for products and services offered in the US. This material might be available from IBM in other languages. However, you may be required to own a copy of the product or product version in that language in order to access it.

IBM may not offer the products, services, or features discussed in this document in other countries. Consult your local IBM representative for information on the products and services currently available in your area. Any reference to an IBM product, program, or service is not intended to state or imply that only that IBM product, program, or service may be used. Any functionally equivalent product, program, or service that does not infringe any IBM intellectual property right may be used instead. However, it is the user's responsibility to evaluate and verify the operation of any non-IBM product, program, or service.

IBM may have patents or pending patent applications covering subject matter described in this document. The furnishing of this document does not grant you any license to these patents. You can send license inquiries, in writing, to:

IBM Director of Licensing
IBM Corporation
North Castle Drive, MD-NC119
Armonk, NY 10504-1785
US

For license inquiries regarding double-byte character set (DBCS) information, contact the IBM Intellectual Property Department in your country or send inquiries, in writing, to:

Intellectual Property Licensing
Legal and Intellectual Property Law
IBM Japan Ltd.
19-21, Nihonbashi-Hakozakicho, Chuo-ku
Tokyo 103-8510, Japan

INTERNATIONAL BUSINESS MACHINES CORPORATION PROVIDES THIS PUBLICATION "AS IS" WITHOUT WARRANTY OF ANY KIND, EITHER EXPRESS OR IMPLIED, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF NON-INFRINGEMENT, MERCHANTABILITY OR FITNESS FOR A PARTICULAR PURPOSE. Some jurisdictions do not allow disclaimer of express or implied warranties in certain transactions, therefore, this statement may not apply to you.

This information could include technical inaccuracies or typographical errors. Changes are periodically made to the information herein; these changes will be incorporated in new editions of the publication. IBM may make improvements and/or changes in the product(s) and/or the program(s) described in this publication at any time without notice.

Any references in this information to non-IBM websites are provided for convenience only and do not in any manner serve as an endorsement of those

websites. The materials at those websites are not part of the materials for this IBM product and use of those websites is at your own risk.

IBM may use or distribute any of the information you provide in any way it believes appropriate without incurring any obligation to you.

Licensees of this program who wish to have information about it for the purpose of enabling: (i) the exchange of information between independently created programs and other programs (including this one) and (ii) the mutual use of the information which has been exchanged, should contact:

IBM Director of Licensing
IBM Corporation
North Castle Drive, MD-NC119
Armonk, NY 10504-1785
US

Such information may be available, subject to appropriate terms and conditions, including in some cases, payment of a fee.

The licensed program described in this document and all licensed material available for it are provided by IBM under terms of the IBM Customer Agreement, IBM International Program License Agreement or any equivalent agreement between us.

Information concerning non-IBM products was obtained from the suppliers of those products, their published announcements or other publicly available sources. IBM has not tested those products and cannot confirm the accuracy of performance, compatibility or any other claims related to non-IBM products. Questions on the capabilities of non-IBM products should be addressed to the suppliers of those products.

This information is for planning purposes only. The information herein is subject to change before the products described become available.

This information contains examples of data and reports used in daily business operations. To illustrate them as completely as possible, the examples include the names of individuals, companies, brands, and products. All of these names are fictitious and any similarity to actual people or business enterprises is entirely coincidental.

COPYRIGHT LICENSE:

This information contains sample application programs in source language, which illustrate programming techniques on various operating platforms. You may copy, modify, and distribute these sample programs in any form without payment to IBM, for the purposes of developing, using, marketing or distributing application programs conforming to the application programming interface for the operating platform for which the sample programs are written. These examples have not been thoroughly tested under all conditions. IBM, therefore, cannot guarantee or imply reliability, serviceability, or function of these programs. The sample programs are provided "AS IS", without warranty of any kind. IBM shall not be liable for any damages arising out of your use of the sample programs.

Each copy or any portion of these sample programs or any derivative work must include a copyright notice as follows:

© (your company name) (year).
Portions of this code are derived from IBM Corp. Sample Programs.
© Copyright IBM Corp. _enter the year or years_.

Terms and conditions for product documentation

Permissions for the use of these publications are granted subject to the following terms and conditions.

Applicability

These terms and conditions are in addition to any terms of use for the IBM website.

Personal use

You may reproduce these publications for your personal, noncommercial use provided that all proprietary notices are preserved. You may not distribute, display or make derivative work of these publications, or any portion thereof, without the express consent of IBM.

Commercial use

You may reproduce, distribute and display these publications solely within your enterprise provided that all proprietary notices are preserved. You may not make derivative works of these publications, or reproduce, distribute or display these publications or any portion thereof outside your enterprise, without the express consent of IBM.

Rights

Except as expressly granted in this permission, no other permissions, licenses or rights are granted, either express or implied, to the publications or any information, data, software or other intellectual property contained therein.

IBM reserves the right to withdraw the permissions granted herein whenever, in its discretion, the use of the publications is detrimental to its interest or, as determined by IBM, the above instructions are not being properly followed.

You may not download, export or re-export this information except in full compliance with all applicable laws and regulations, including all United States export laws and regulations.

IBM MAKES NO GUARANTEE ABOUT THE CONTENT OF THESE PUBLICATIONS. THE PUBLICATIONS ARE PROVIDED "AS-IS" AND WITHOUT WARRANTY OF ANY KIND, EITHER EXPRESSED OR IMPLIED, INCLUDING BUT NOT LIMITED TO IMPLIED WARRANTIES OF MERCHANTABILITY, NON-INFRINGEMENT, AND FITNESS FOR A PARTICULAR PURPOSE.

Copyright notice

Describes general use restrictions and trademarks related to this document and the software described in this document.

Index

Special characters

.NET languages   2

A

all different constraint
    definition   34
    example   34
alternative constraint   40
application
    debugging   111
    designing   89
    developing   117
arithmetic constraint   24
    definition   29
arithmetic expression
    definition   29
    example   29
arity   16
array   14

B

backtrack   61
bound reduction   47
branch   61, 66, 72
branch decision   66

C

C++   2
channel   13
choice point   66
Concert Technology   5
conjunction   32
constrained variables   6
constraint   24
    aggregation   56, 59
    all different   34
    allowed assignment   33, 57
    alternative   40
    arithmetic   24, 29, 47
    cardinality   49
    compatibility   33, 57
    conjunction   32
    cumul function expression   43
    definition   5, 7, 24
    disjunction   32
    distribution   36, 56
    implication   32
    interval sequence variable   41
    interval variable   38, 39, 40
    inverse   35
    isomorphism   40
    lexicographic   36
    logical   26, 32, 49
    minimum distance   34
    negation   32
    no overlap   41

constraint *(continued)*
    number   72
    packing   35
    precedence   38
    presence   39
    sequence   41
    spanning   40
    specialized   25, 34, 51
    state function   44
    surrogate   93, 125
    symmetry reducing   91
    synchronize   40
    table   33, 57
constraint propagation
    during search   9
    initial   7
    types   7
constructive disjunction   49
constructive search   61, 76
cost function
    proxy for   119
cumul function expression   42, 43

D

data
    multiple sets   124
    preparing   117
debugging   111, 112
decision variable
    definition   5, 6, 19
    domain   23
    evaluating   80
    evaluator   80
    fixed   46
    introducing order   92
    number   72
    selecting   80
design principle   89
detecting symmetry   91
disjunction   32
distribution constraint   36
div   32
division
    integer   32
domain
    definition   45
    example   6
    expression   23
    holes   47
    reduction   45
    variable   5, 23
domain reduction   7
    absolute value   47
    arithmetic constraint   47
    bound   47
    counting expression   55
    definition   45
    element expression   54
    example   7, 45
    inference level   52

domain reduction *(continued)*
    introduction   7
    specialized constraint   51
dual representation   90

E

element expression
    definition   31
    domain reduction   54
elementary cumul function
  expression   42
environment   13
evaluator   80, 81, 83
    custom   83
    value   81
example
    proxy for cost function   119
exception   111
expression   21
    arithmetic   29, 47
    counting   55
    cumul function   42
    cumul function height   43
    domain   23
    element   31
    elementary cumul function   42
    integer division   32
    objective   118
extraction   66

F

fail   72
function
    state   44

I

IConstraint interface   24
IIntervalSequenceVar interface   21
IIntervalVar interface   19
IIntExpr interface   21
IIntVar interface   19
IloAllowedAssignmens function   57
IloAllowedAssignments function   33
IloConstraint class   24
IloCP class
    AllDiffInferenceLevel parameter   51
    constructor   62
    DepthFirst parameter value   76
    domain method   46
    getObjValue method   62
    getValue method   62
    printInformation method   72
    propagate method   46
    SearchType parameter   76
    setParameter method   11
    solve method   62
    startNewSearch method   63

IBM®

Printed in USA